

EECE 5639 Computer Vision I

Lecture 11

Generalized Hough Transform, RANSAC, Snakes, Region Segmentation

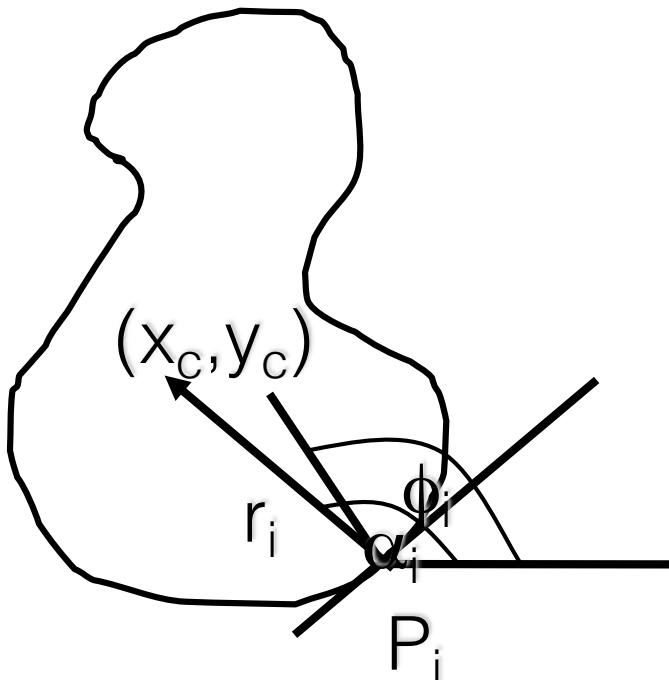
Project 2 is out ...

Next Class

Region Segmentation, Homographies

Generalizing the H.T.

The H.T. can be used even if the curve has not a simple analytic form!



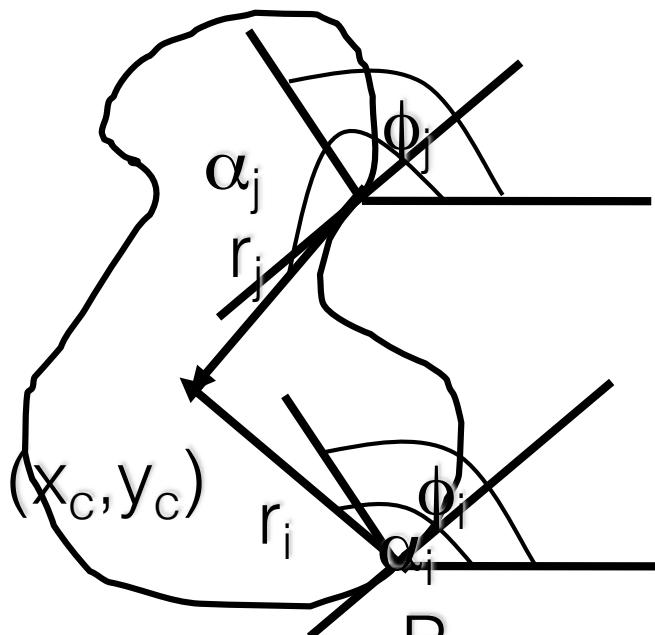
$$x_c = x_i + r_i \cos(\alpha_i)$$

$$y_c = y_i + r_i \sin(\alpha_i)$$

1. Pick a reference point (x_c, y_c)
2. For $i = 1, \dots, n$:
 - a. Draw segment to P_i on the boundary.
 - b. Measure its length r_i , and its orientation α_i .
 - c. Write the coordinates of (x_c, y_c) as a function of r_i and α_i
 - d. Record the gradient orientation ϕ_i at P_i .
5. Build a table with the data, indexed by ϕ_i .

Generalizing the H.T.

Suppose, there were m **different** gradient orientations: ($m \leq n$)



$$x_c = x_i + r_i \cos(\alpha_i)$$

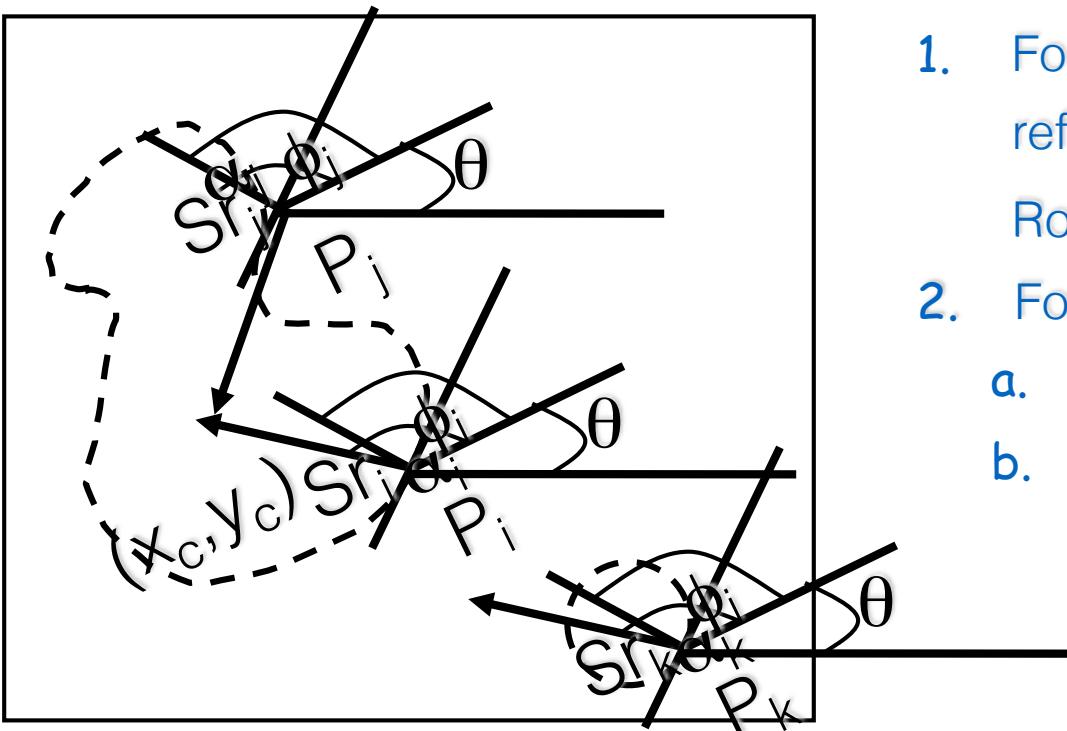
$$y_c = y_i + r_i \sin(\alpha_i)$$

ϕ_1	$(r^1_1, \alpha^1_1), (r^1_2, \alpha^1_2), \dots, (r^1_{n1}, \alpha^1_{n1})$
ϕ_2	$(r^2_1, \alpha^2_1), (r^2_2, \alpha^2_2), \dots, (r^2_{n2}, \alpha^2_{n2})$
.	.
.	.
.	.
ϕ_m	$(r^m_1, \alpha^m_1), (r^m_2, \alpha^m_2), \dots, (r^m_{nm}, \alpha^m_{nm})$

H.T. table

Generalized H.T. Algorithm:

Finds a rotated, scaled, and translated version of the curve:



$$x_c = x_i + r_i \cos(\alpha_i)$$

$$y_c = y_i + r_i \sin(\alpha_i)$$

1. Form an A accumulator array of possible reference points (x_c, y_c) , scaling factor S and Rotation angle θ .
2. For each edge (x, y) in the image:
 - a. Compute $\phi(x, y)$
 - b. For each (r, α) corresponding to $\phi(x, y)$ do:
 1. For each S and θ :
 - a. $x_c = x_i + r(\phi) S \cos[\alpha(\phi) + \theta]$
 - b. $y_c = y_i + r(\phi) S \sin[\alpha(\phi) + \theta]$
 - c. $A(x_c, y_c, S, \theta) ++$
 3. Find maxima of A .

H.T. Summary

H.T. is a “voting” scheme

points vote for a set of parameters describing a line or curve.

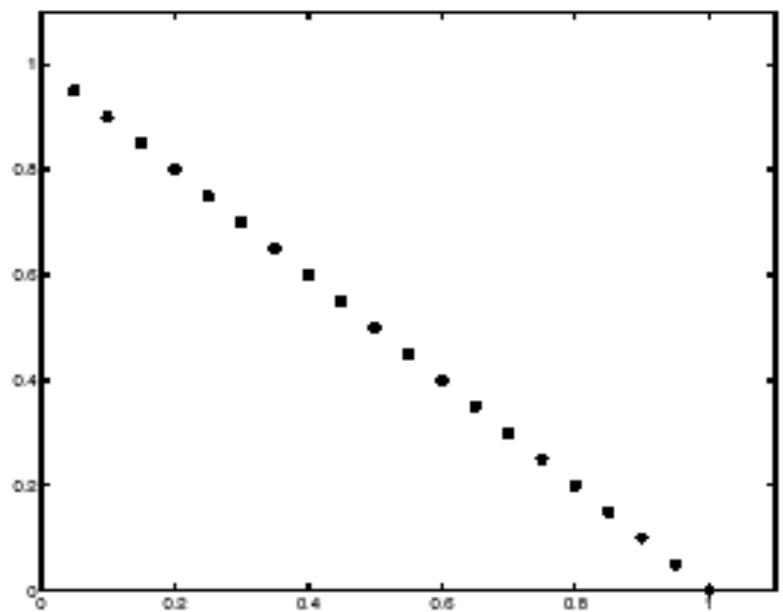
The more votes for a particular set

the more evidence that the corresponding curve is present in the image.

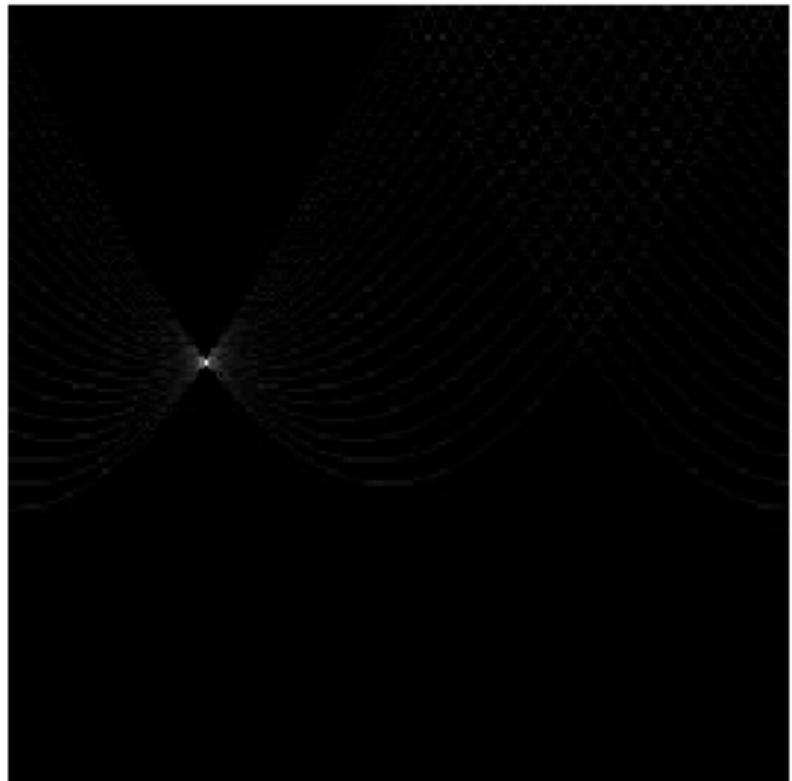
Can detect MULTIPLE curves in one shot.

Computational cost increases with the number of parameters describing the curve.

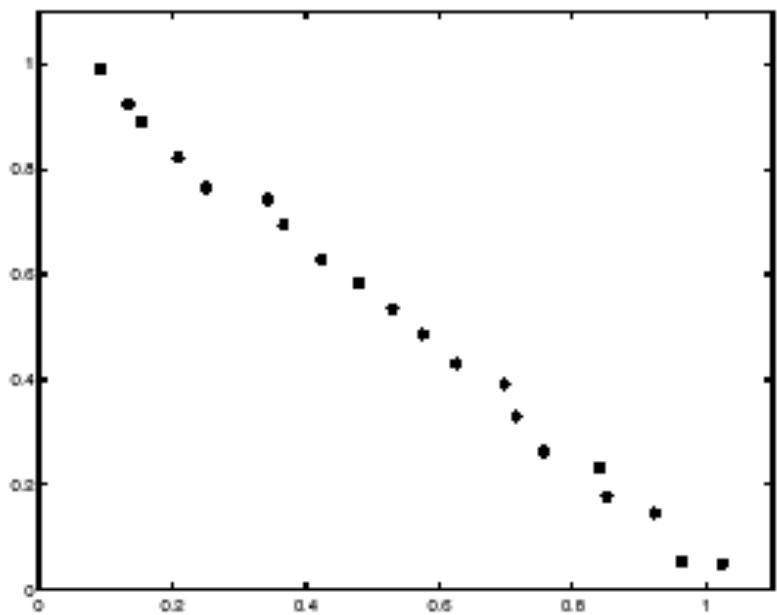
Hough Transf. & Noise



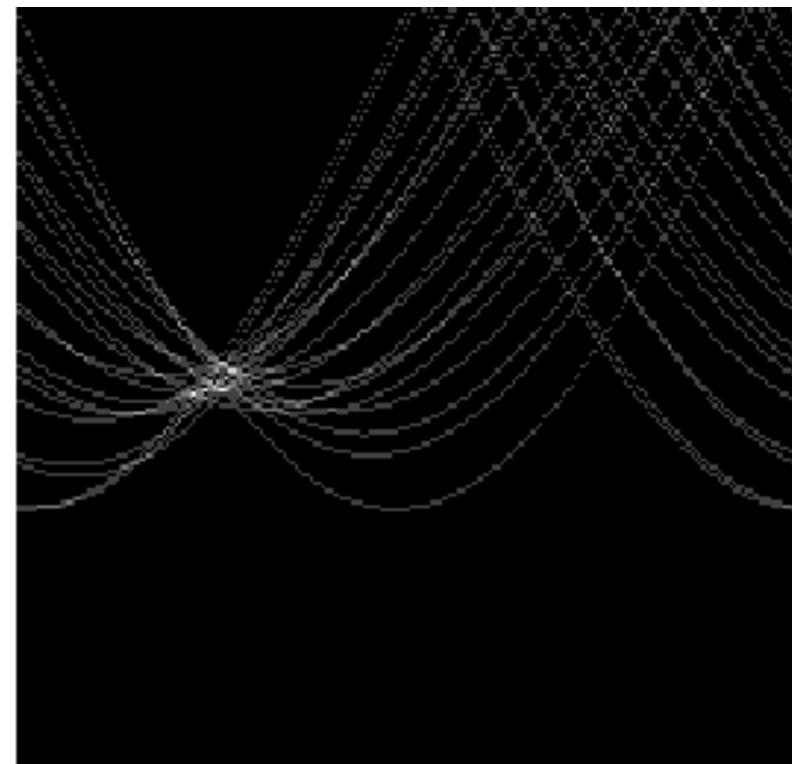
tokens



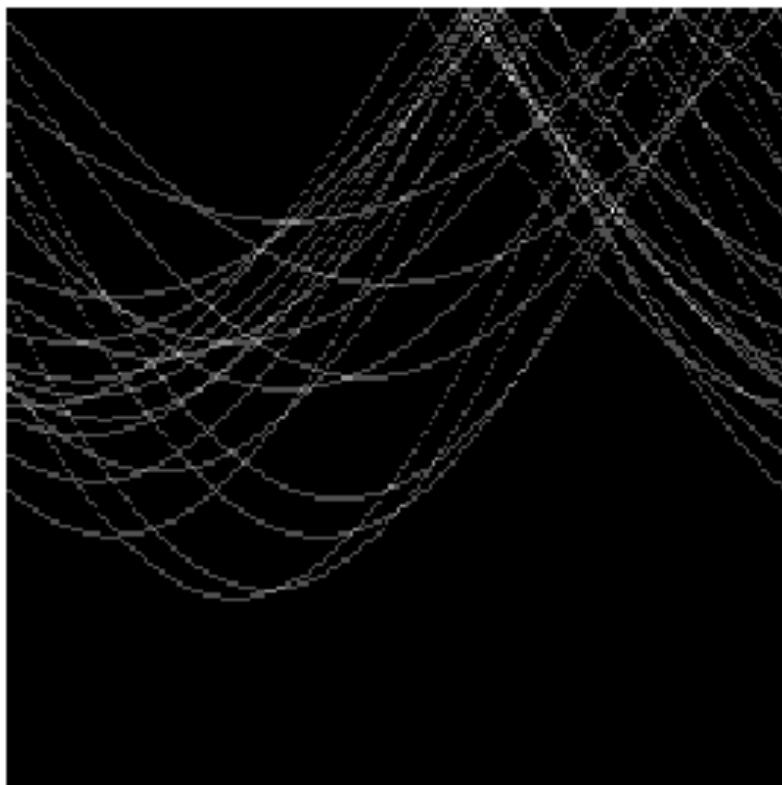
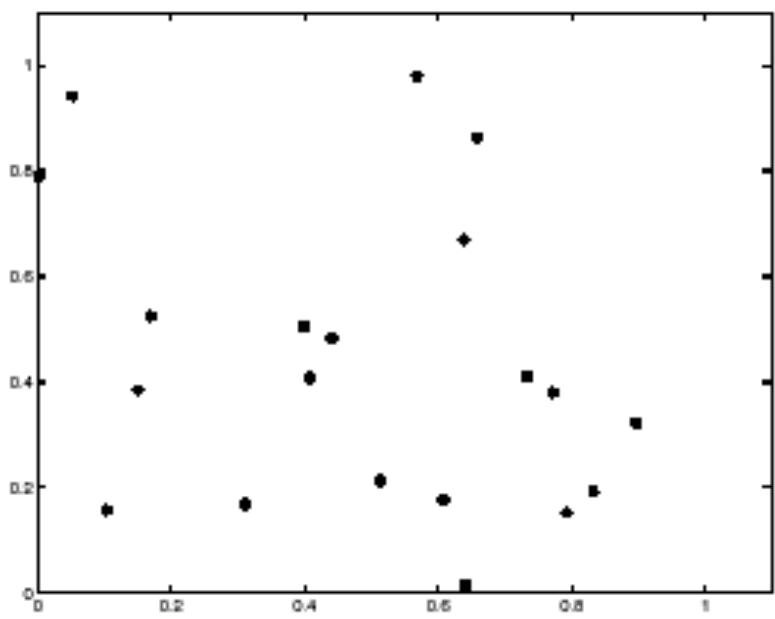
votes

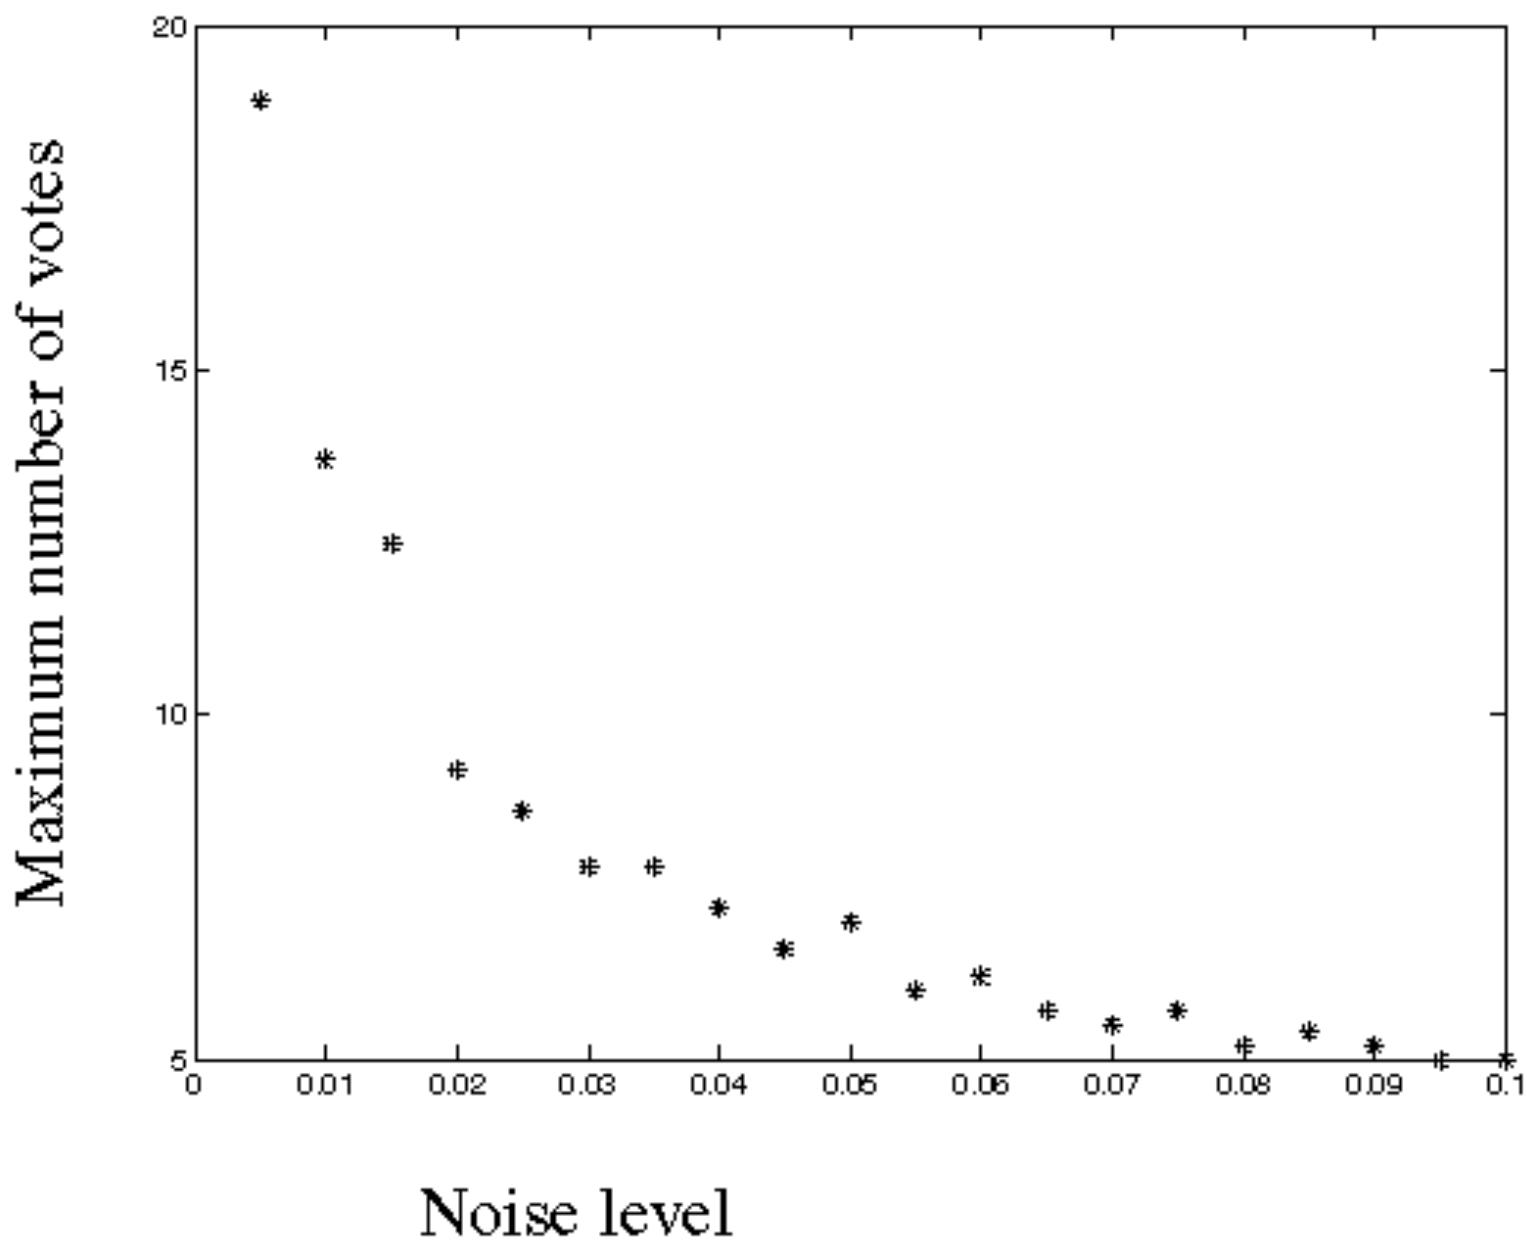


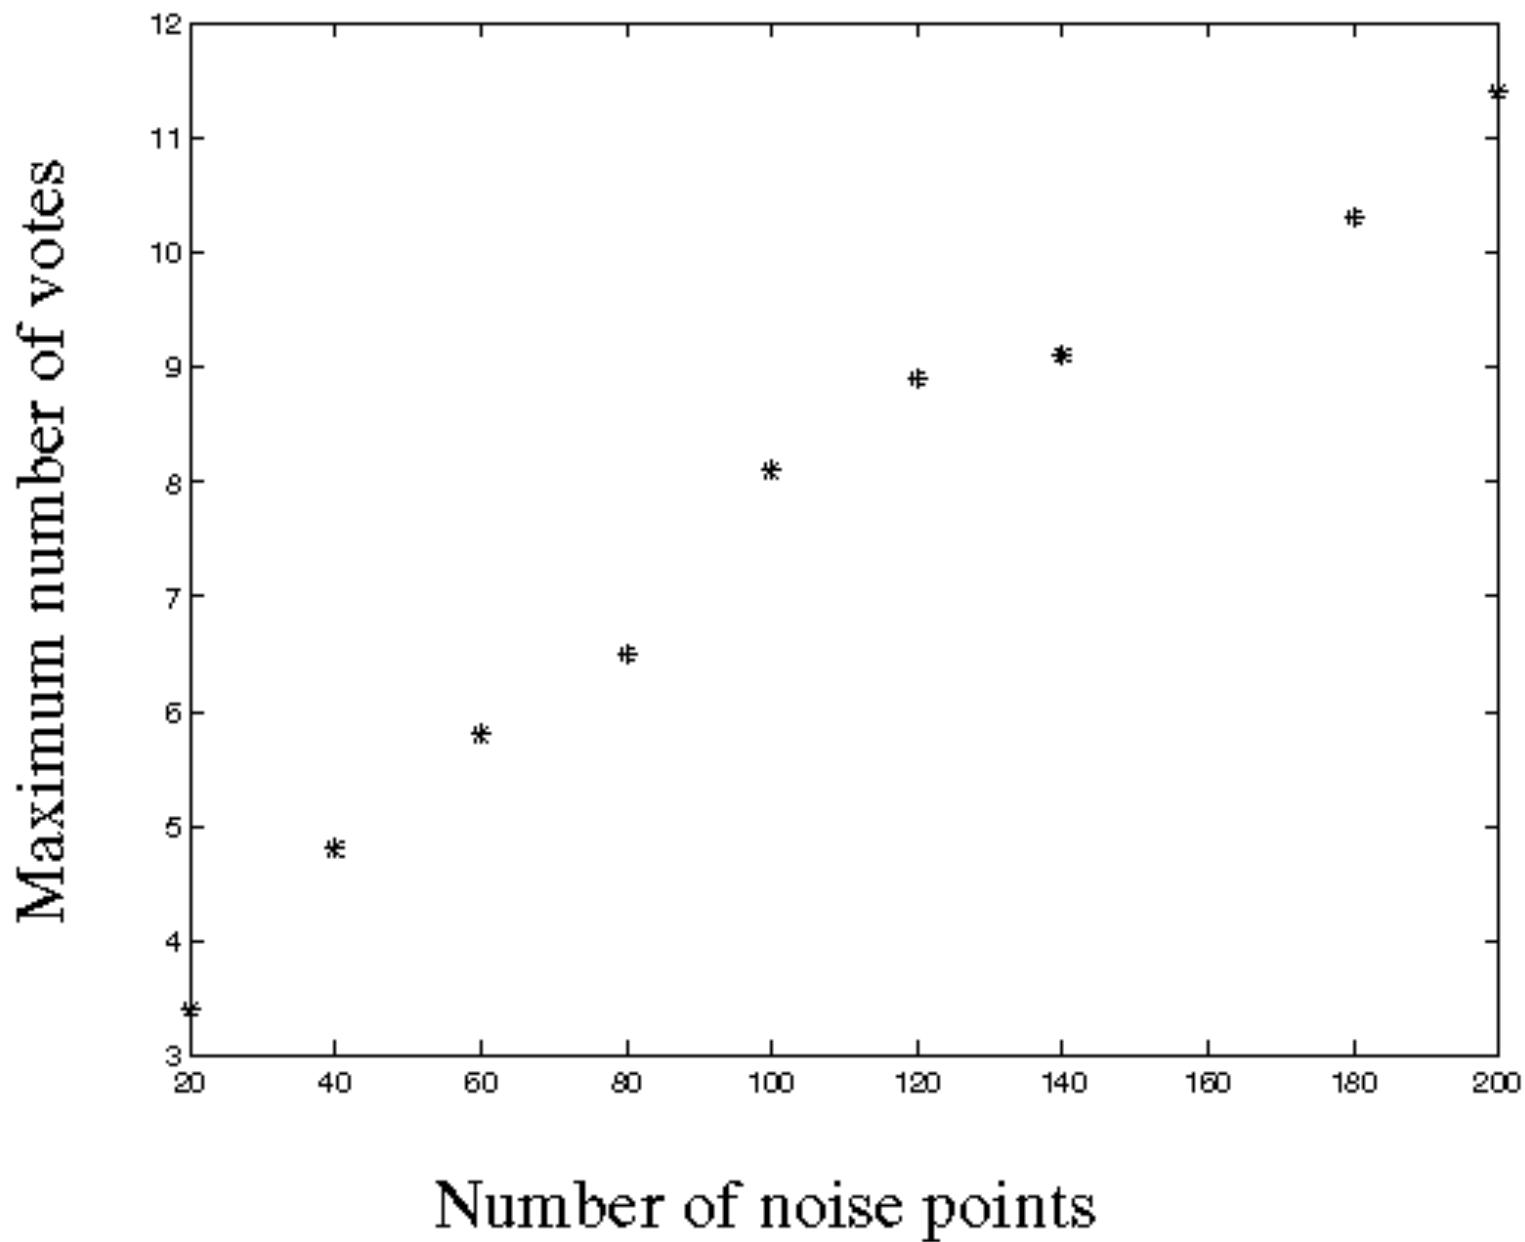
tokens



votes



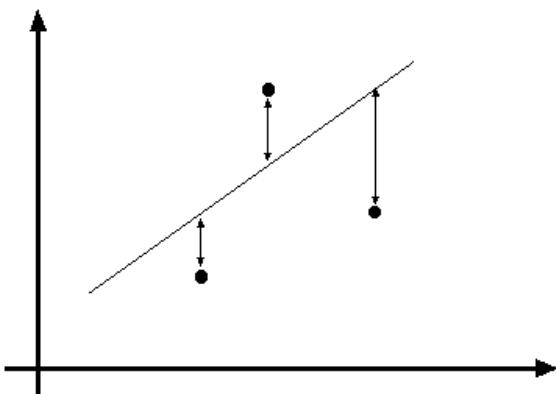




Segmentation by Fitting

Fitting Lines

Using Least Squares Fitting Error:



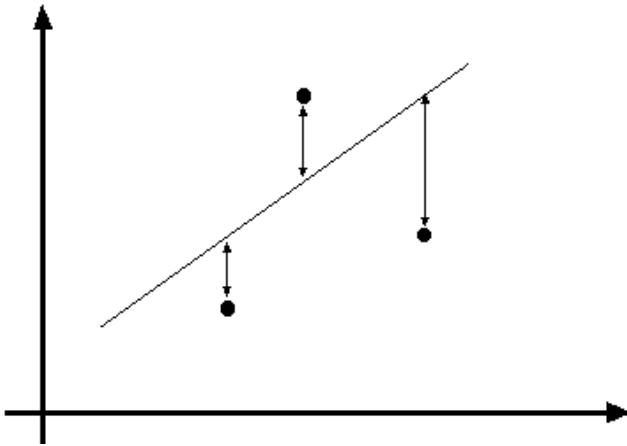
$$\min \Phi = \sum_i (y_i - ax_i - b)^2$$

$$\frac{\partial \Phi}{\partial a} = 2a \sum_i x_i^2 - 2 \sum_i x_i(y_i - b) = 0$$

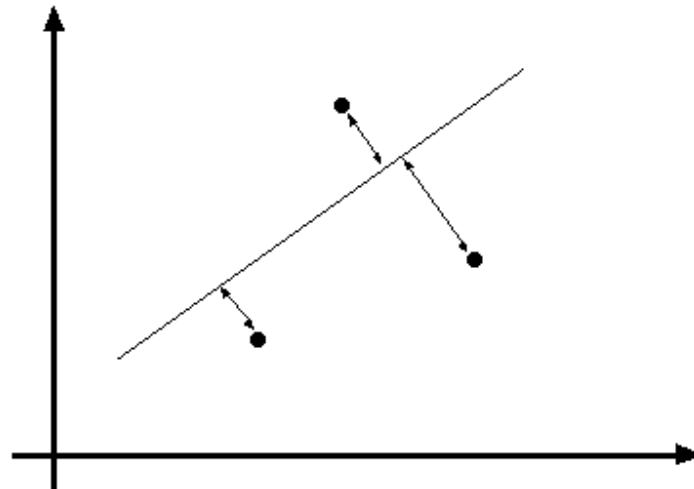
$$\frac{\partial \Phi}{\partial b} = -2 \sum_i (y_i - ax_i - b) = 0$$

$$\begin{aligned}\bar{xy} &= \bar{x}^2 a + \bar{x} b \\ \bar{y} &= \bar{x} a + b\end{aligned}$$

$$\begin{bmatrix} \bar{xy} \\ \bar{y} \end{bmatrix} = \begin{bmatrix} \bar{x}^2 & \bar{x} \\ \bar{x} & 1 \end{bmatrix} \begin{bmatrix} a \\ b \end{bmatrix}$$



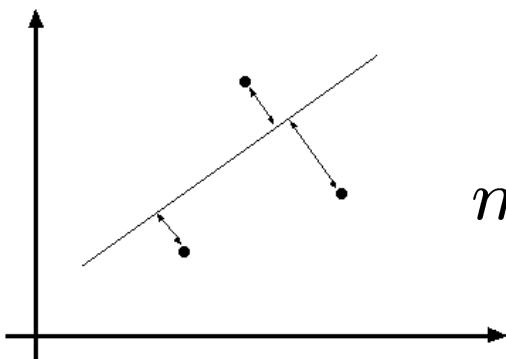
Line fitting can be max.
likelihood - but choice of
model is important



Fitting Lines

Using Total Least Squares Fitting Error:

$$\min \Phi = \sum (ax_i + by_i + c)^2$$
$$\text{s.t. } a^2 + b^2 = 1$$



$$\min \Lambda = \sum_i (ax_i + by_i + c)^2 + \lambda(a^2 + b^2 - 1)$$

$$\frac{\partial \Lambda}{\partial a} = 0 \quad \frac{\partial \Lambda}{\partial b} = 0 \quad \frac{\partial \Lambda}{\partial c} = 0 \quad \frac{\partial \Lambda}{\partial \lambda} = 0$$

$$\begin{bmatrix} \bar{x^2} - \bar{x}\bar{x} & \bar{xy} - \bar{x}\bar{y} \\ \bar{xy} - \bar{x}\bar{y} & \bar{y^2} - \bar{y}\bar{y} \end{bmatrix} \begin{bmatrix} a \\ b \end{bmatrix} = \mu \begin{bmatrix} a \\ b \end{bmatrix}$$

Eigenvalue
Problem!

Who came from which line?

Assume we know how many lines there are - but which lines are they?

easy, if we know who came from which line

Possible strategies

Hough transform

Incremental line fitting

K-means

Algorithm 15.1: Incremental line fitting by walking along a curve, fitting a line to runs of pixels along the curve, and breaking the curve when the residual is too large

Put all points on curve list, in order along the curve

Empty the line point list

Empty the line list

Until there are too few points on the curve

 Transfer first few points on the curve to the line point list

 Fit line to line point list

 While fitted line is good enough

 Transfer the next point on the curve

 to the line point list and refit the line

 end

 Transfer last point(s) back to curve

 Refit line

 Attach line to line list

end

Algorithm 15.2: K-means line fitting by allocating points to the closest line and then refitting.

Hypothesize k lines (perhaps uniformly at random)

or

Hypothesize an assignment of lines to points
and then fit lines using this assignment

Until convergence

 Allocate each point to the closest line

 Refit lines

end

Robustness

As we have seen, squared error can be a source of bias in the presence of noise points

One fix is EM - we'll not do this in this class

Another is an M-estimator

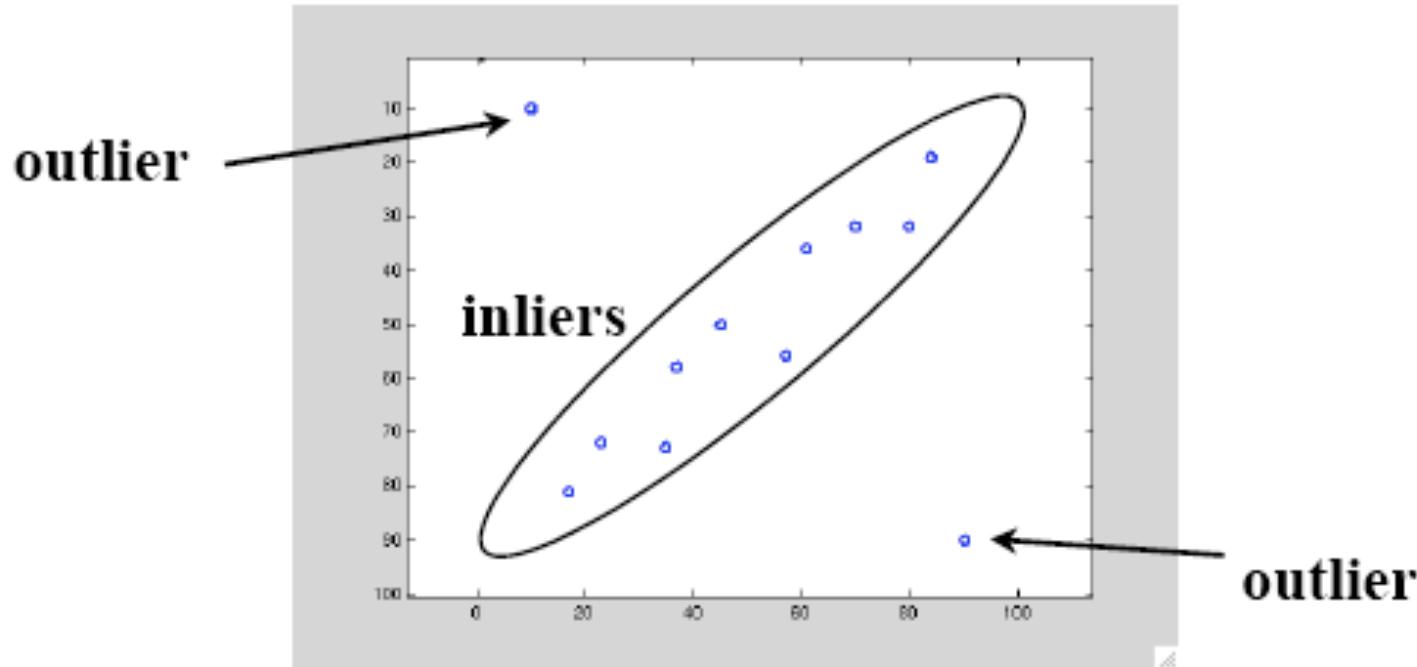
 Square nearby, threshold far away

A third is RANSAC

 Search for good points

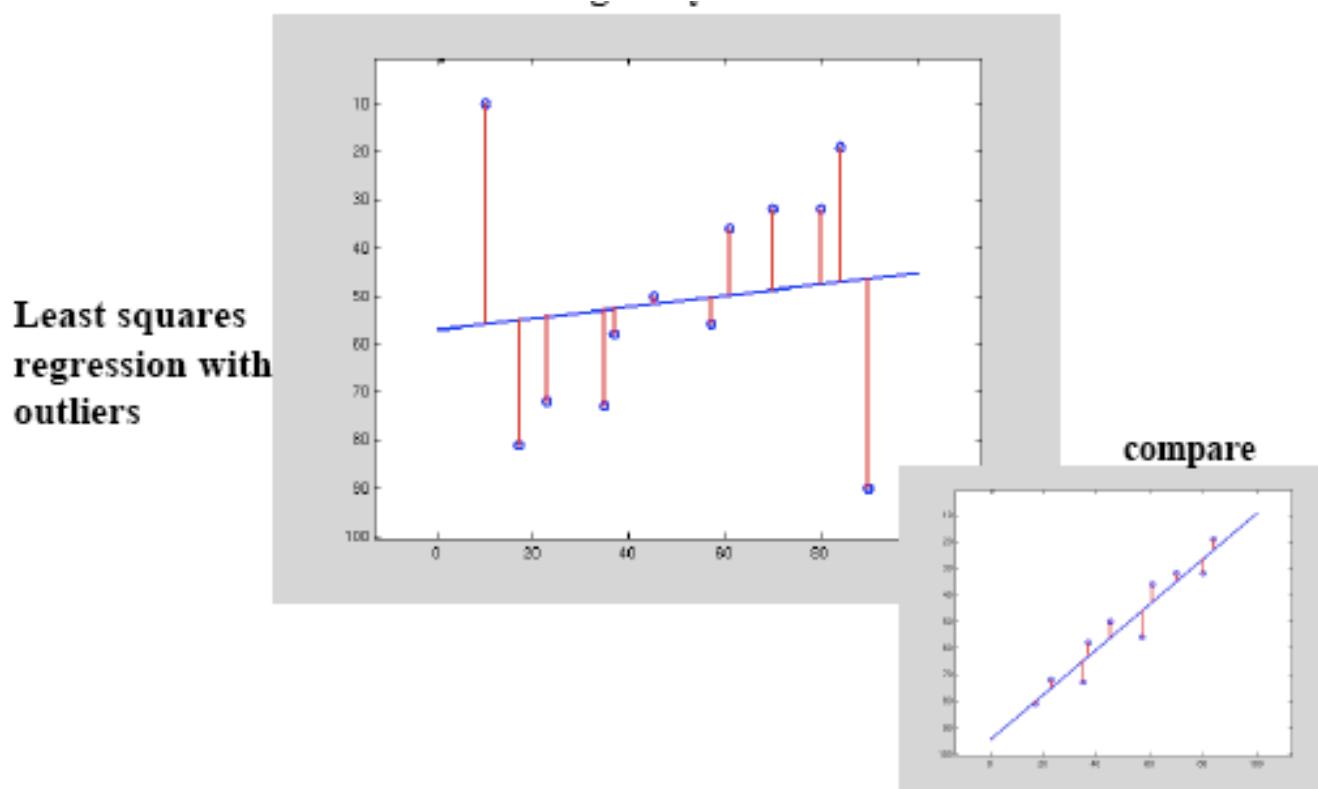
Inliers-Outliers

Loosely speaking, **outliers** are “bad data” points that do not fit the model. Points that fit the model are **inliers**.



Problems with Outliers

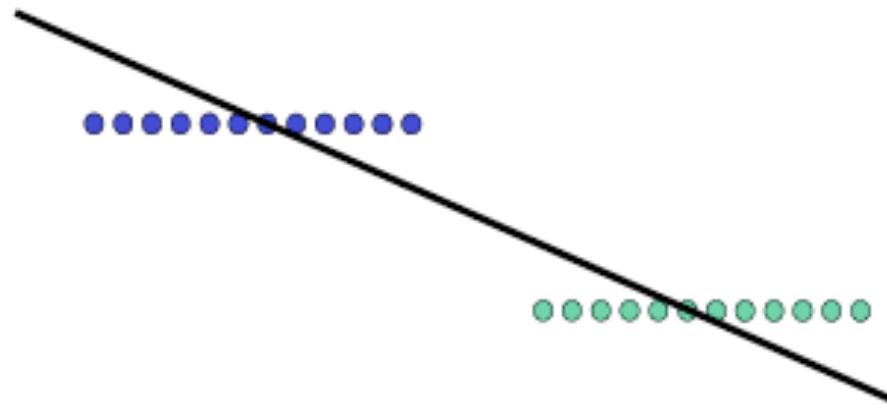
Least square estimation is very sensitive to outliers! :
Few outliers can GREATLY skew the result.



Outliers are not the only problem

Multiple structures can also skew results.

The fitting procedure implicitly assumes ONE instance



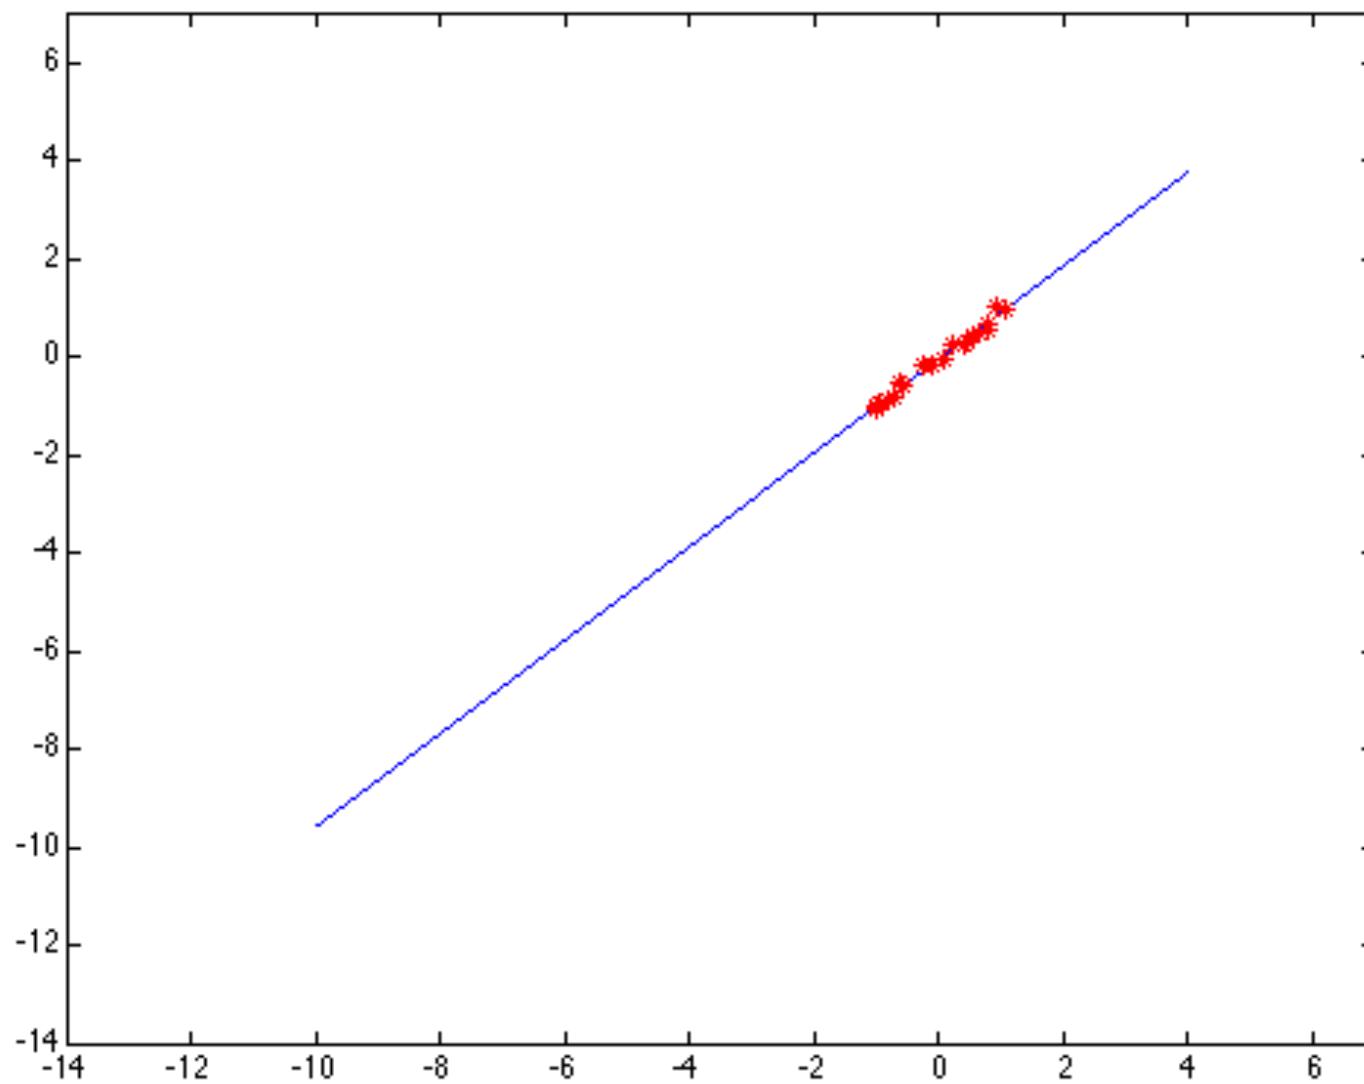
Robust Estimation

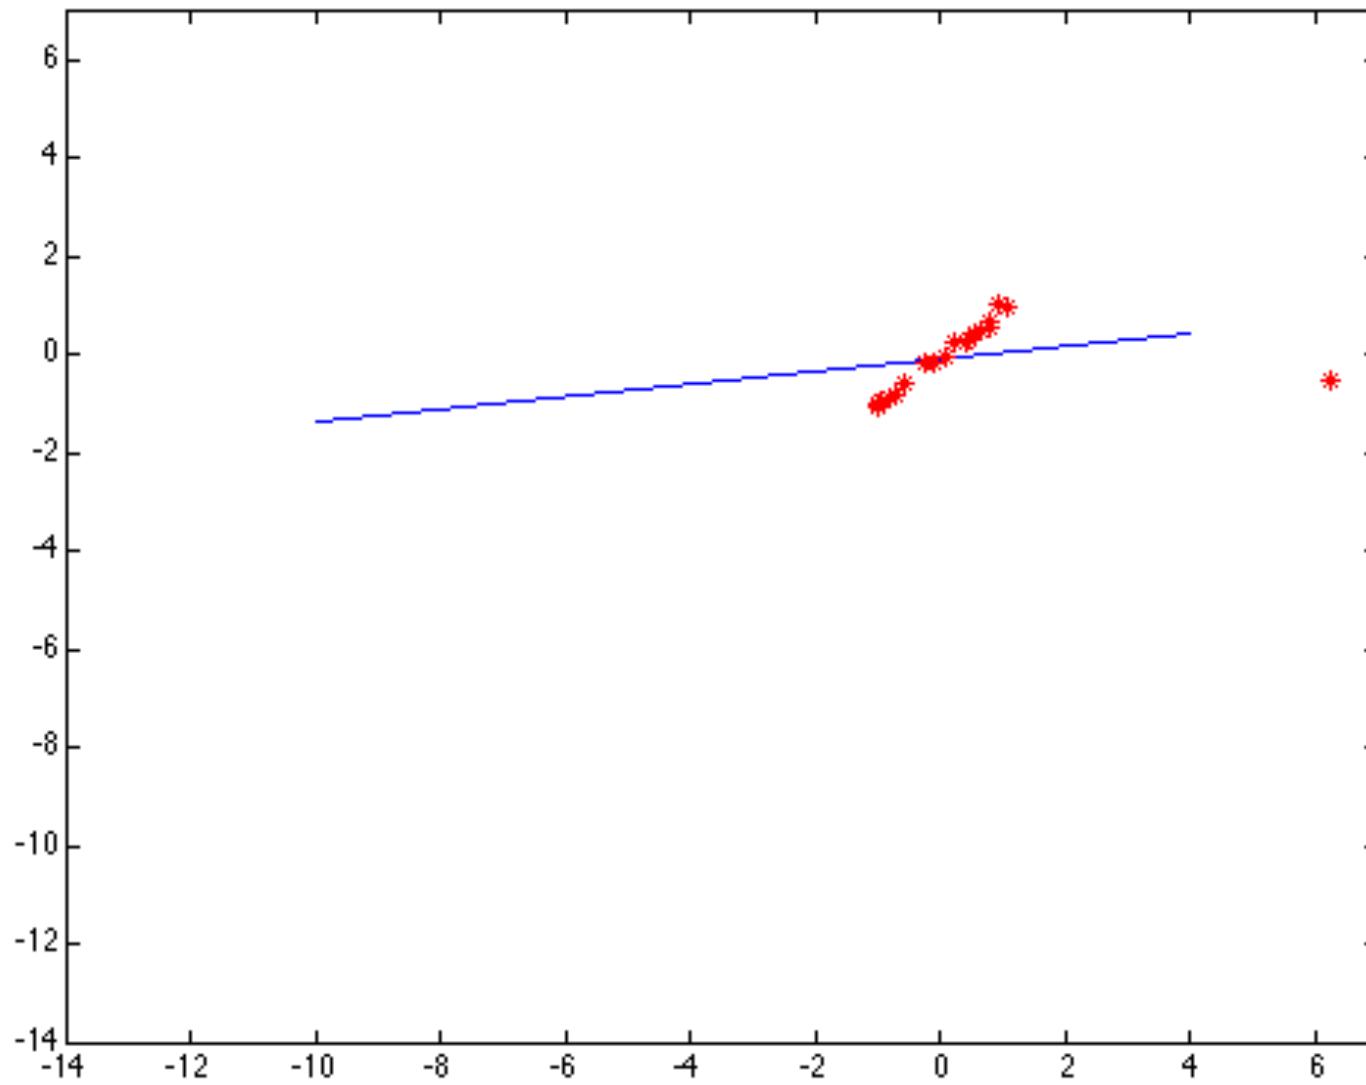
Two steps:

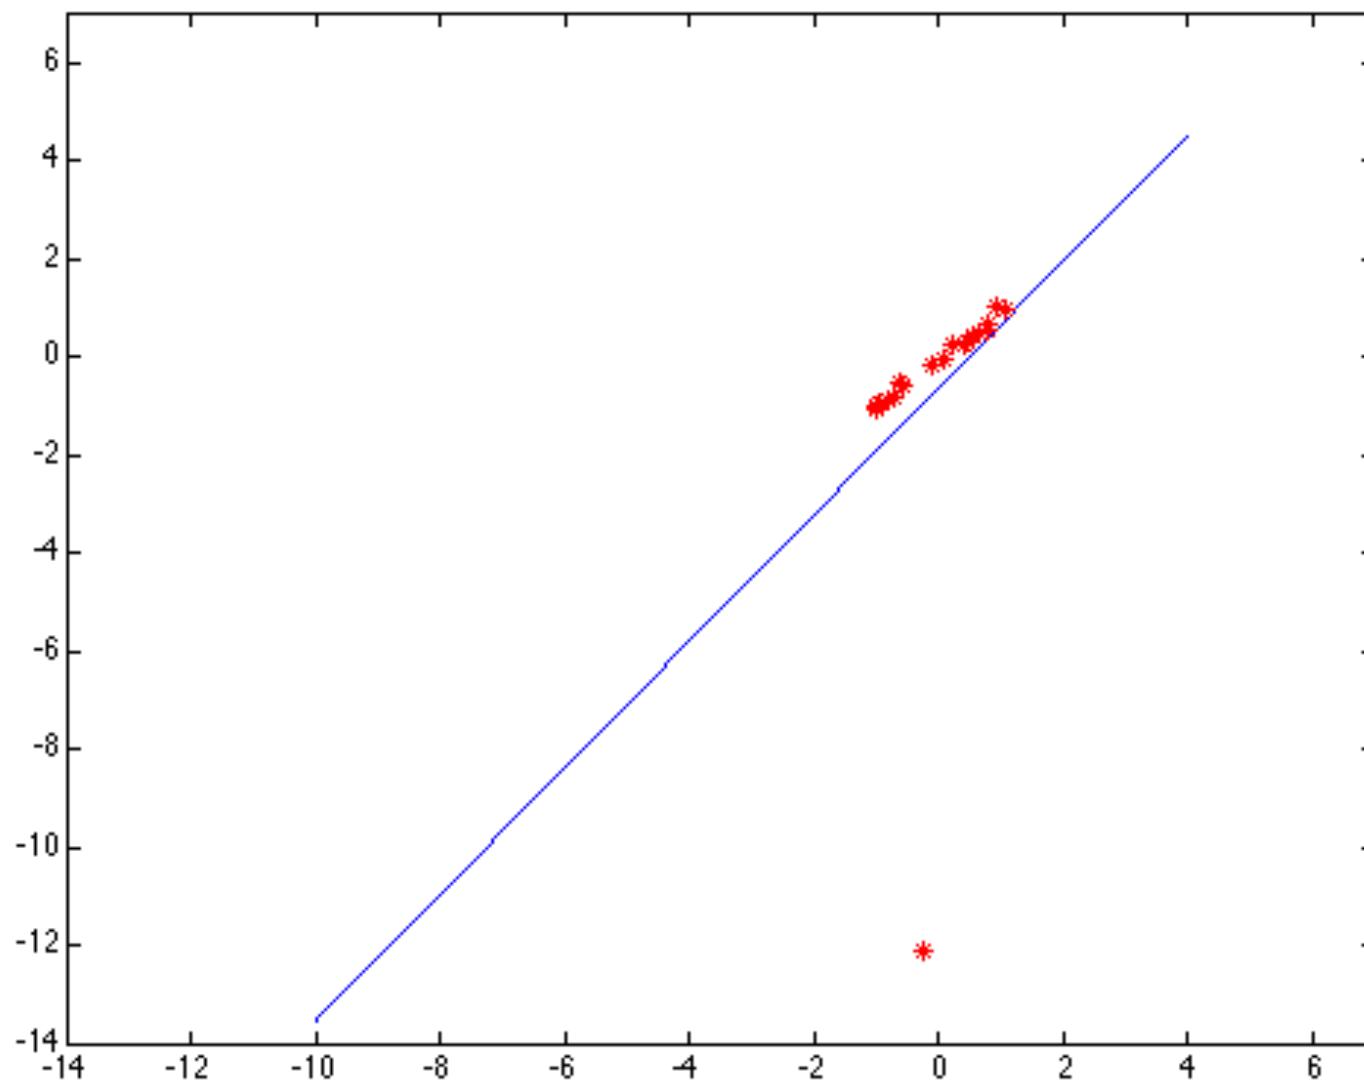
Classify data into INLIERS and OUTLIERS

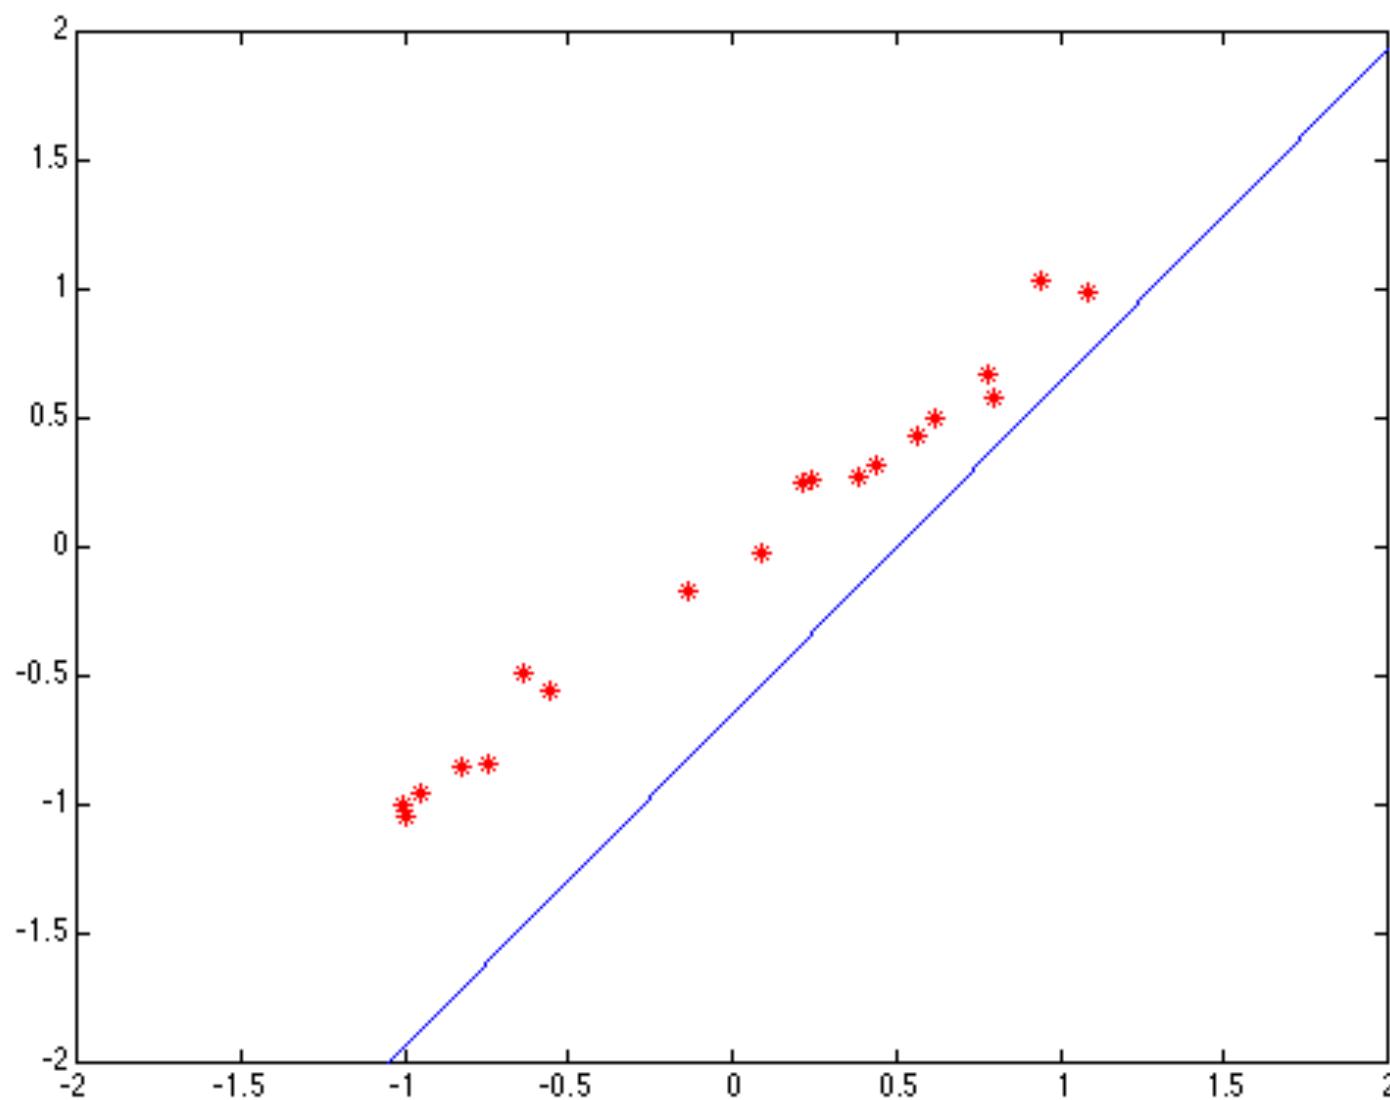
Use only INLIERS to fit the model

RANSAC is and example of this approach.









Robustness and M-estimators

LSE methods are very sensitive to outliers: one bad point can have tremendous effect on the solution.

An M-estimator is used to give different weights to the errors:

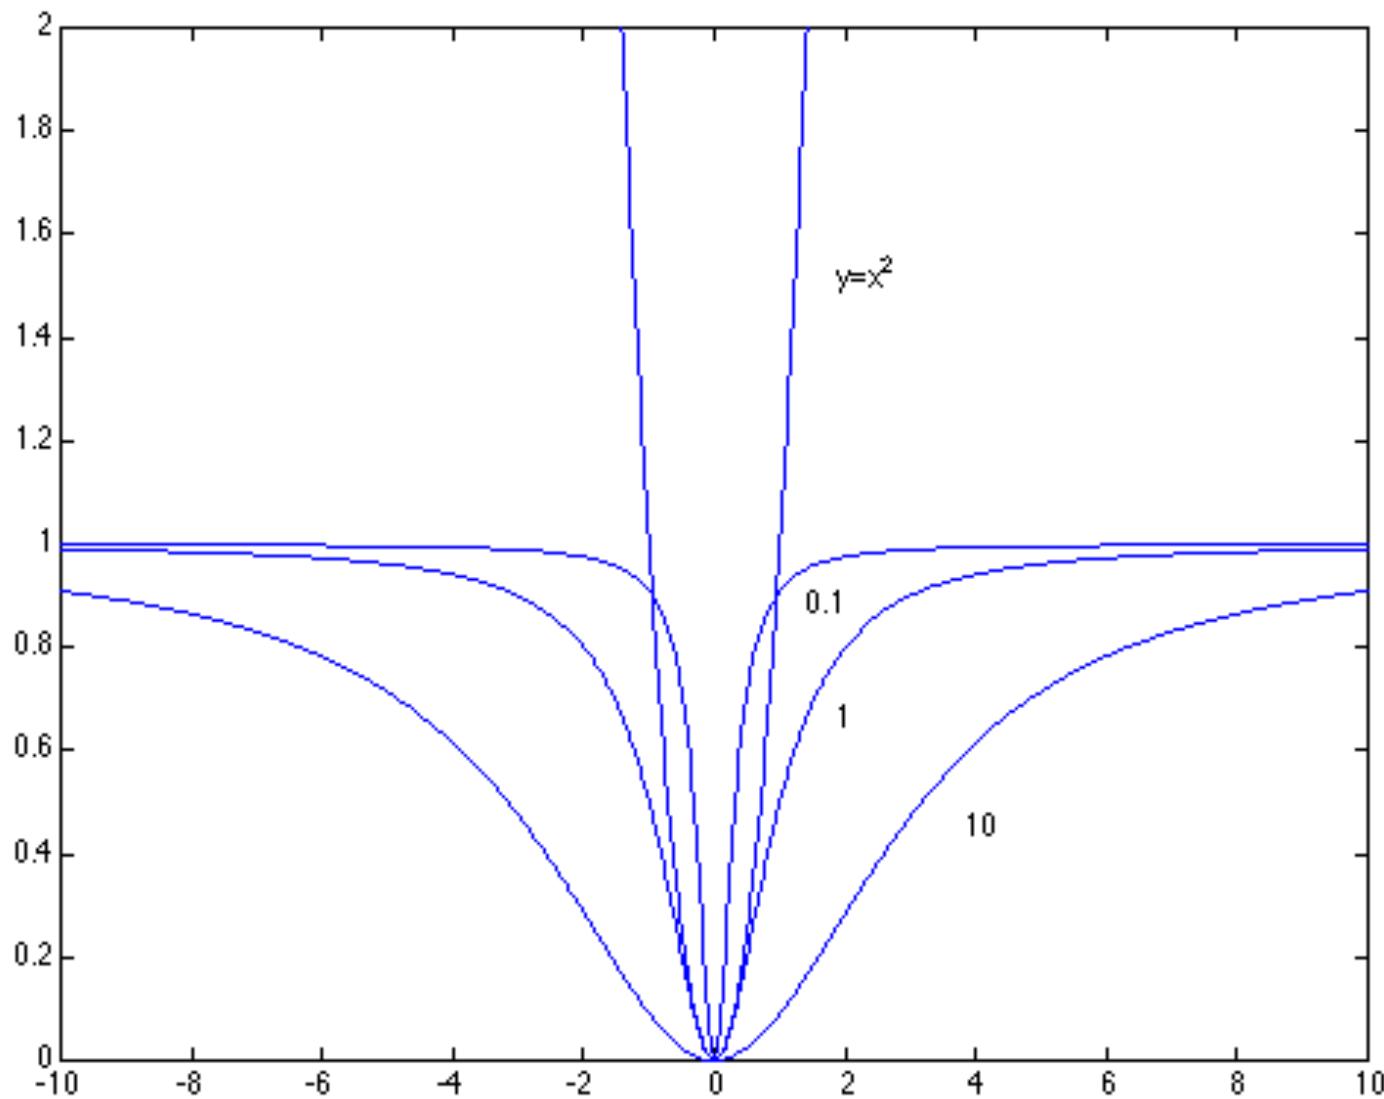
$$\min \sum_i \rho(r_i(x_i, \theta); \sigma)$$

Residual error at
 x_i

Set of parameters
Being fitted

Parameter
Controlling
Error influence

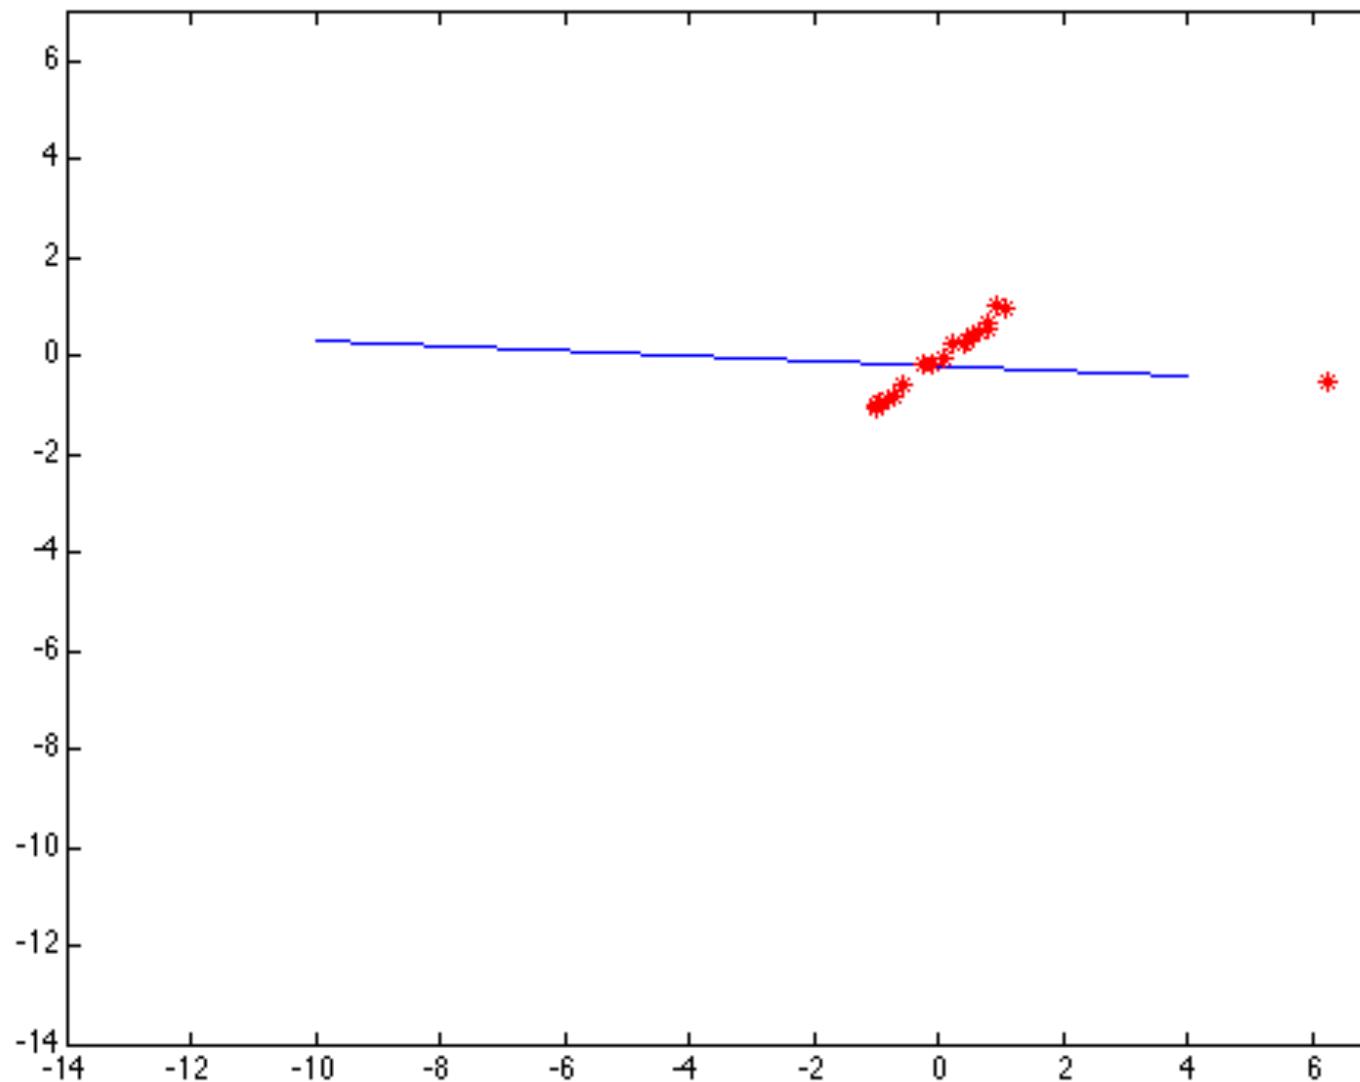
$$\rho(u; \sigma) = \frac{u^2}{u^2 + \sigma^2}$$



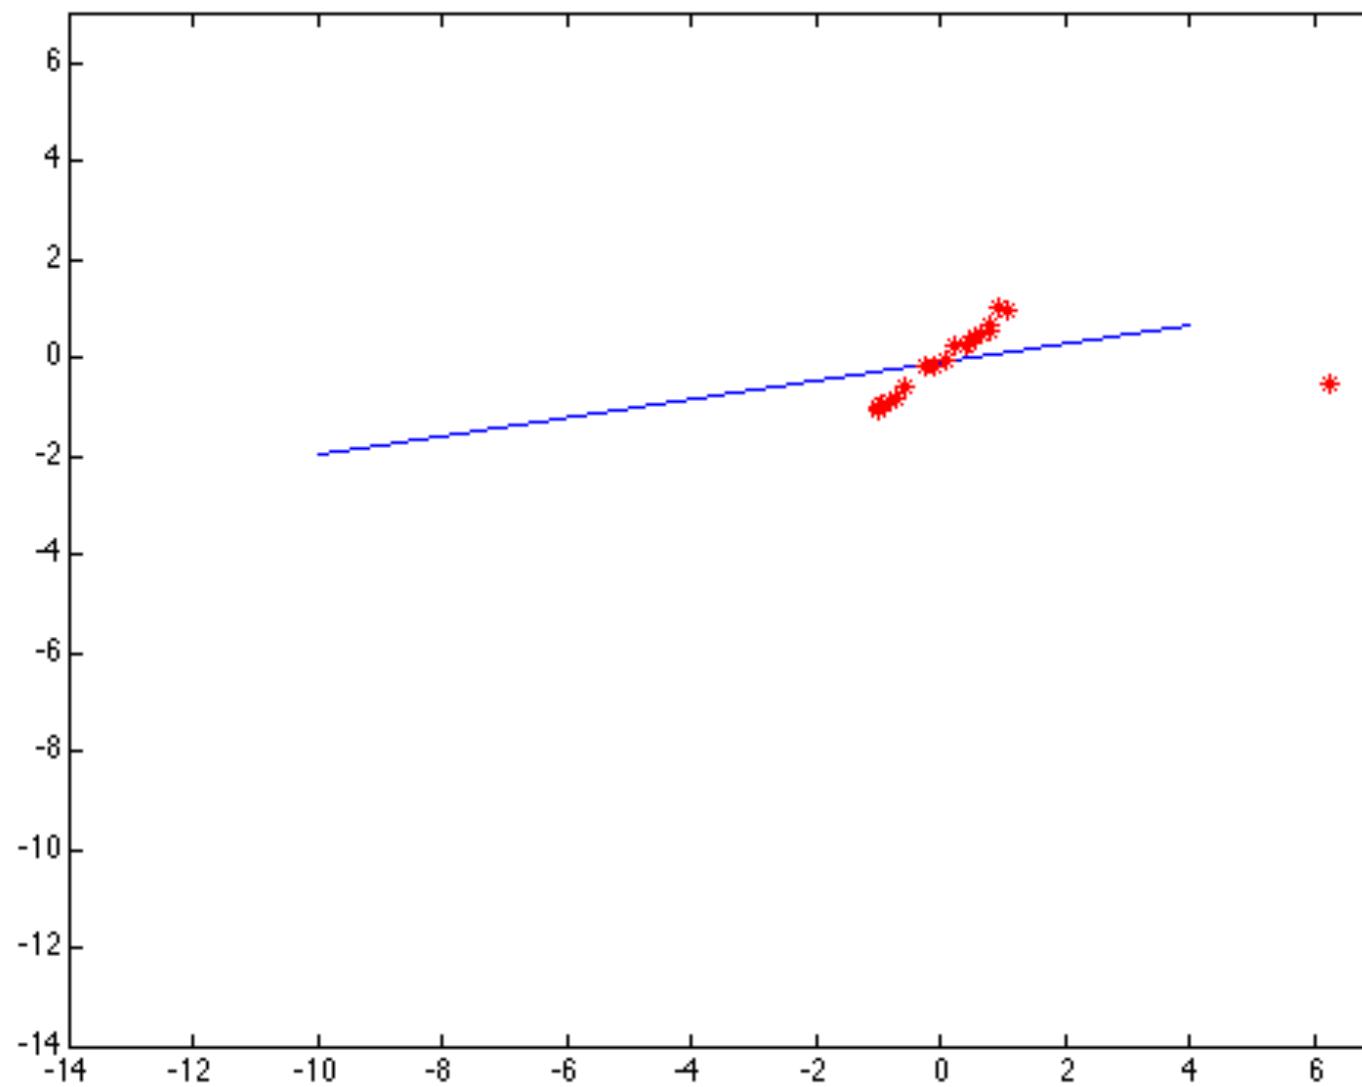
Issues:

1. Minimizing objective is not longer linear:
Solutions must be found interactively
2. Need to decide the scale parameter sigma.

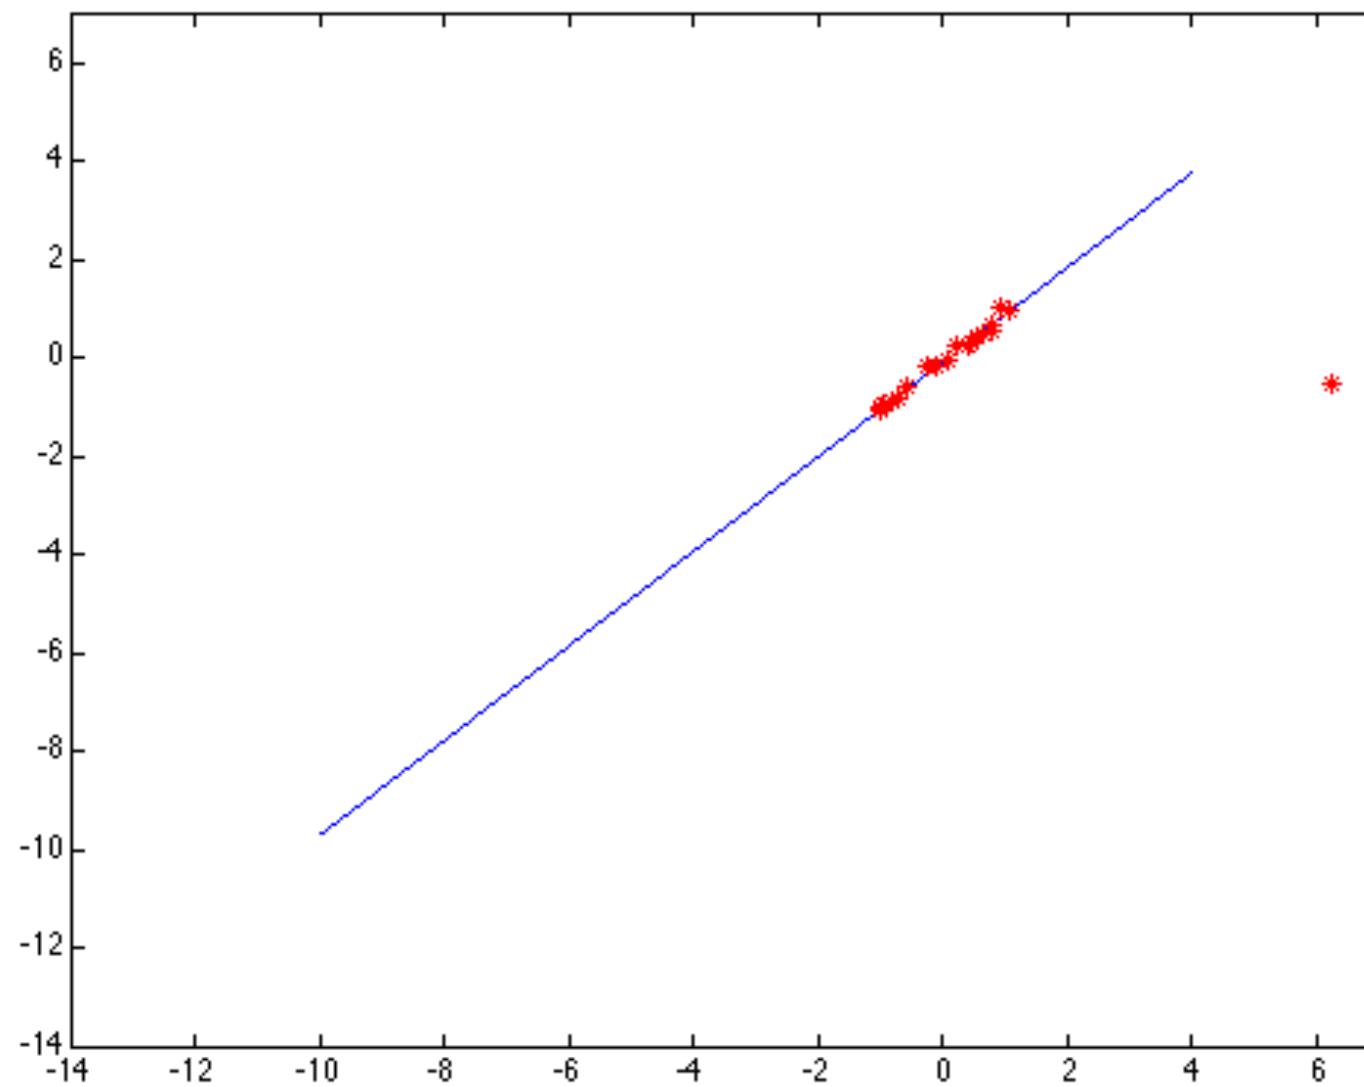
Too small sigma



Too large sigma



Good sigma value:



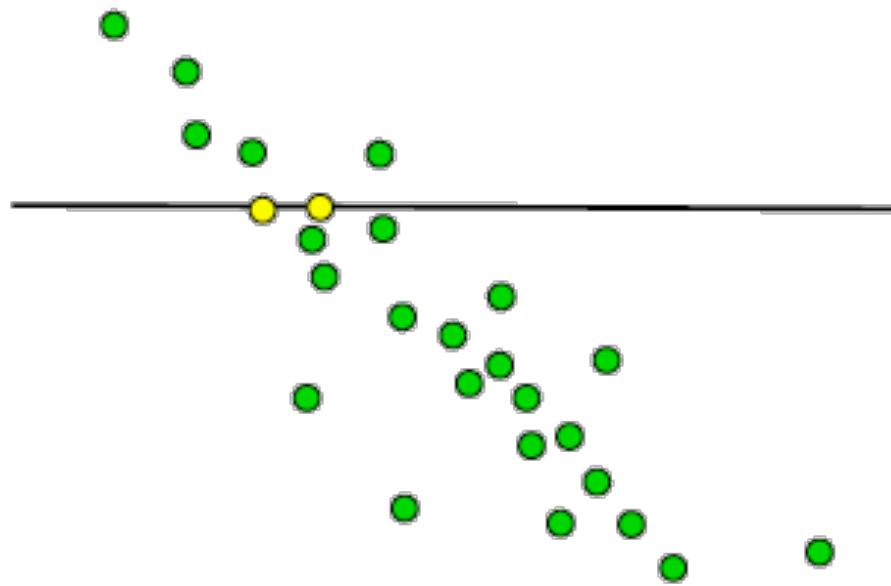
RANSAC

RANDom SAMple Consensus

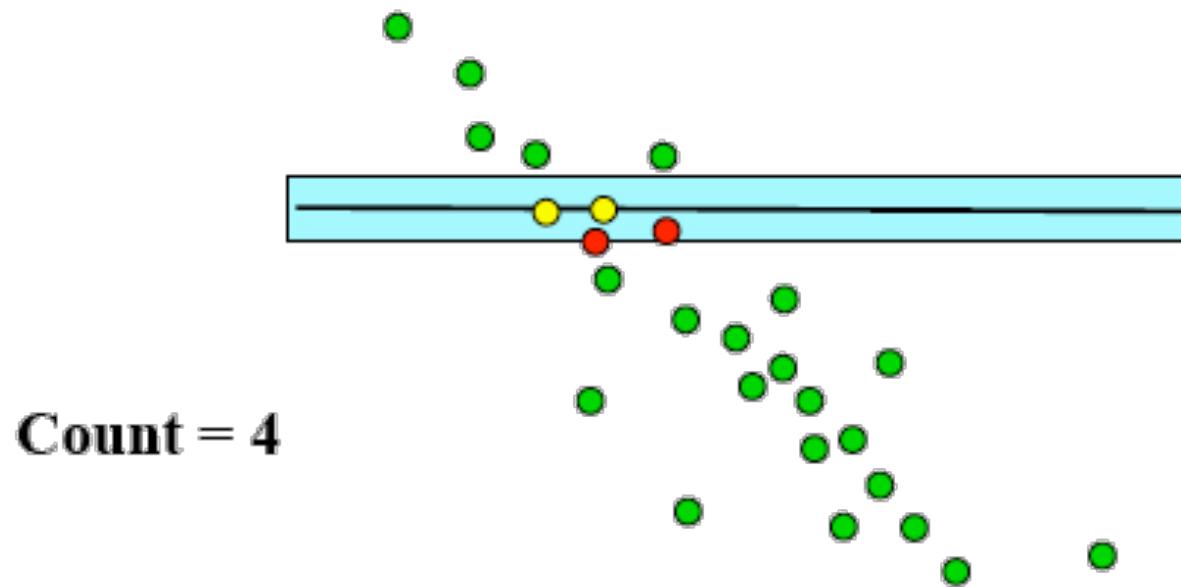
RANSAC procedure (line example)



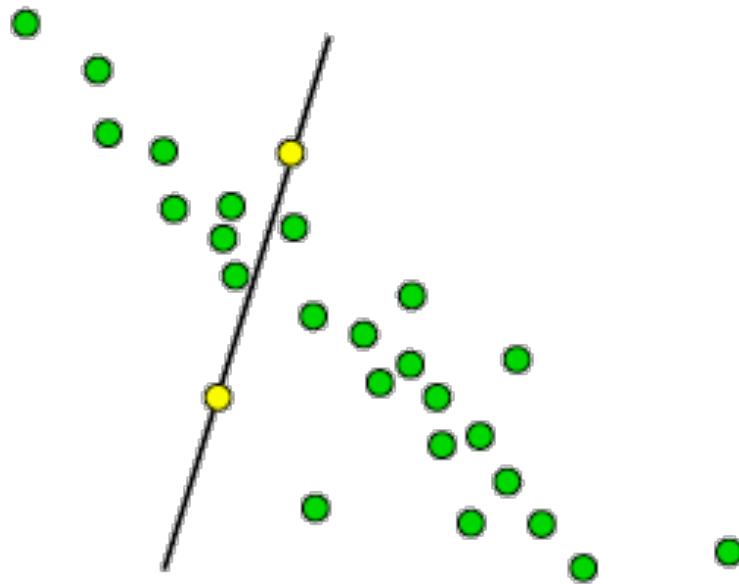
RANSAC procedure (line example)



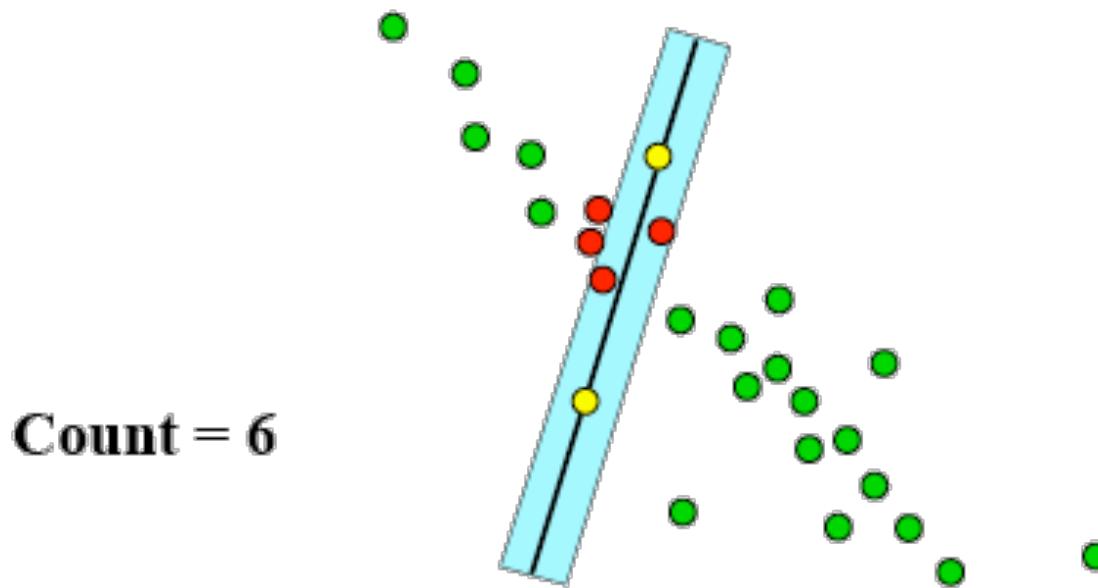
RANSAC procedure (line example)



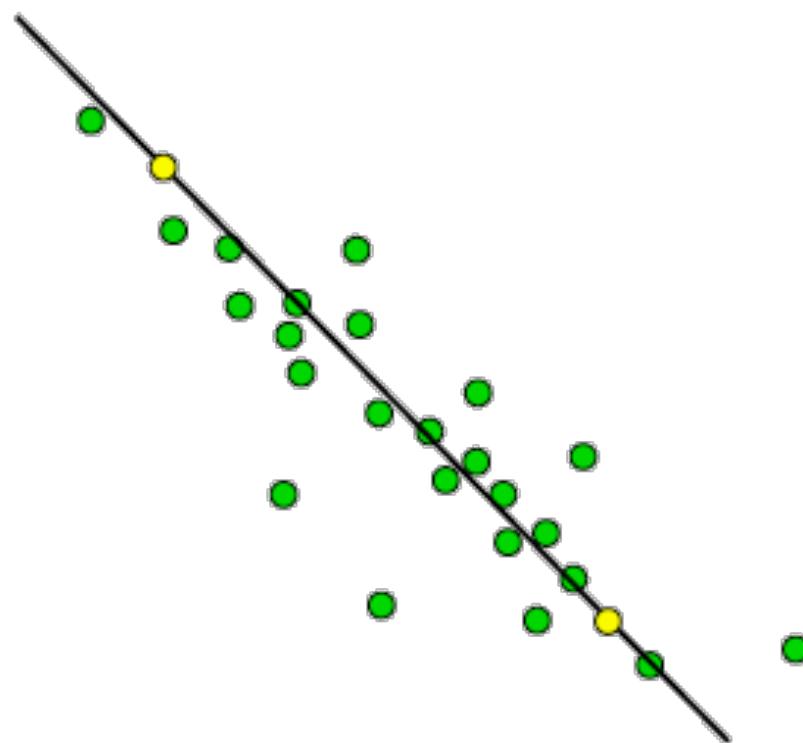
RANSAC procedure (line example)



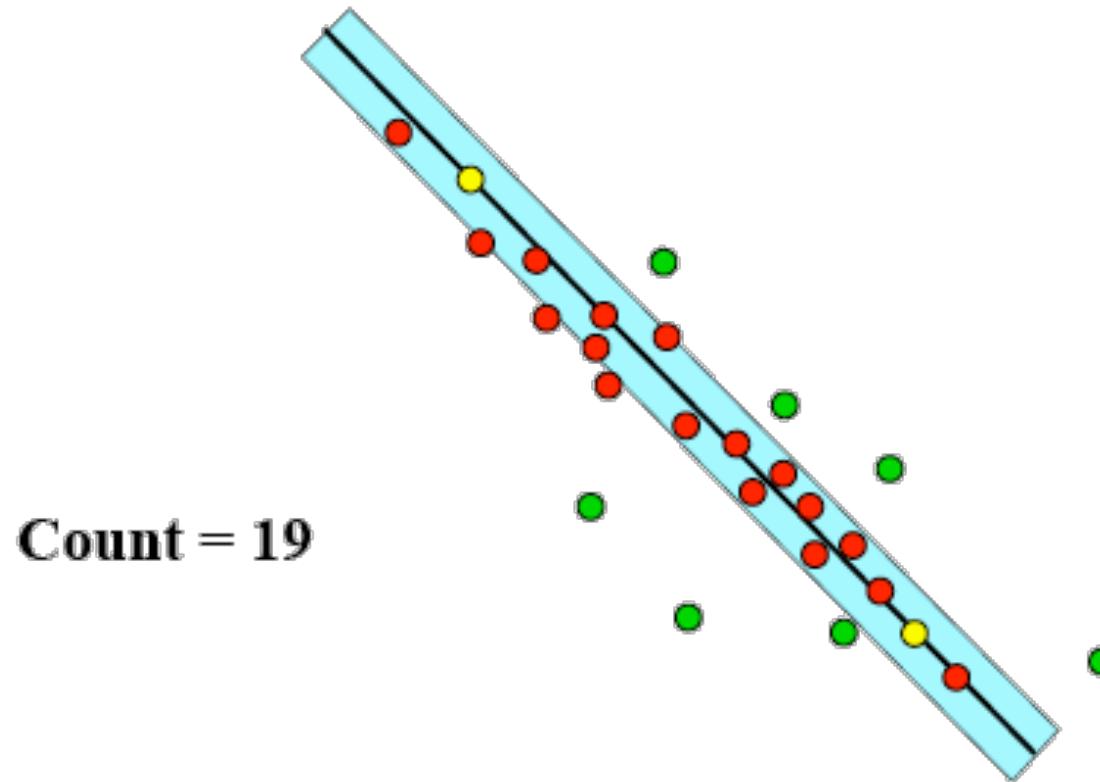
RANSAC procedure (line example)



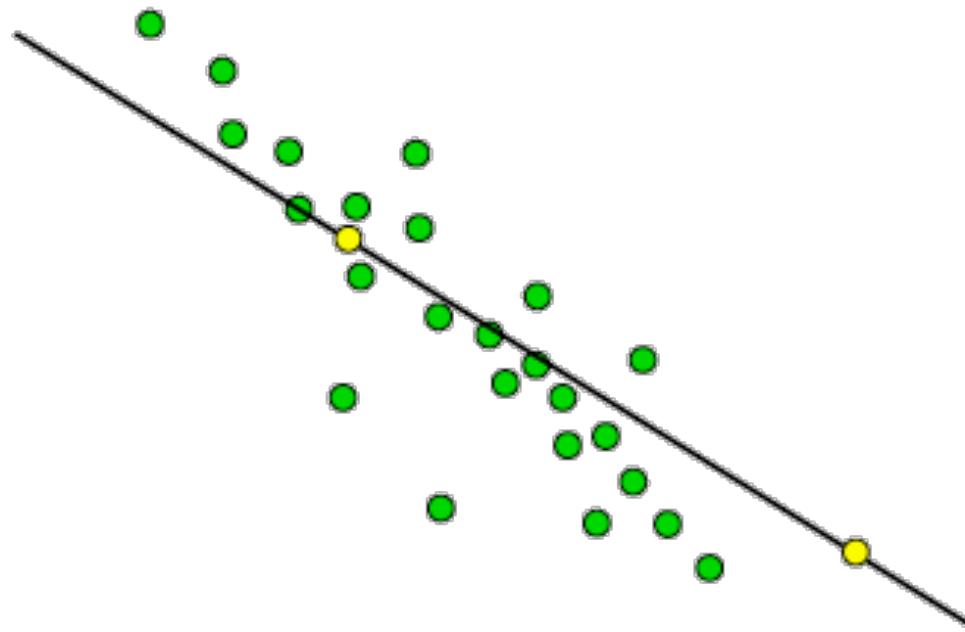
RANSAC procedure (line example)



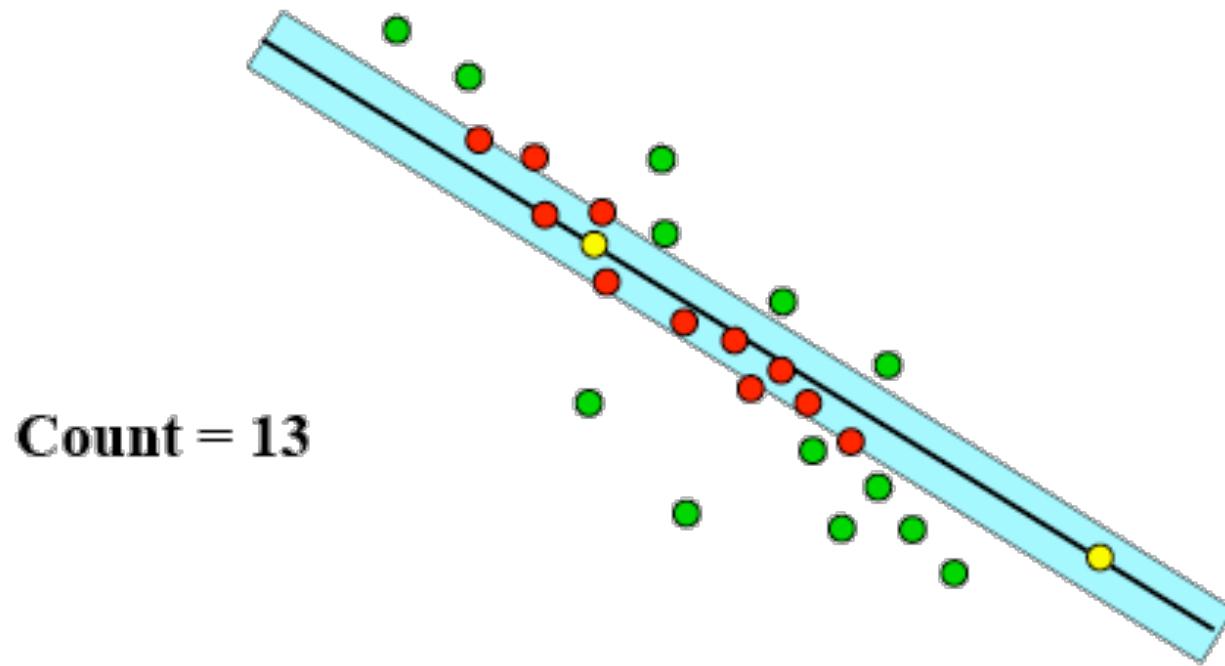
RANSAC procedure (line example)



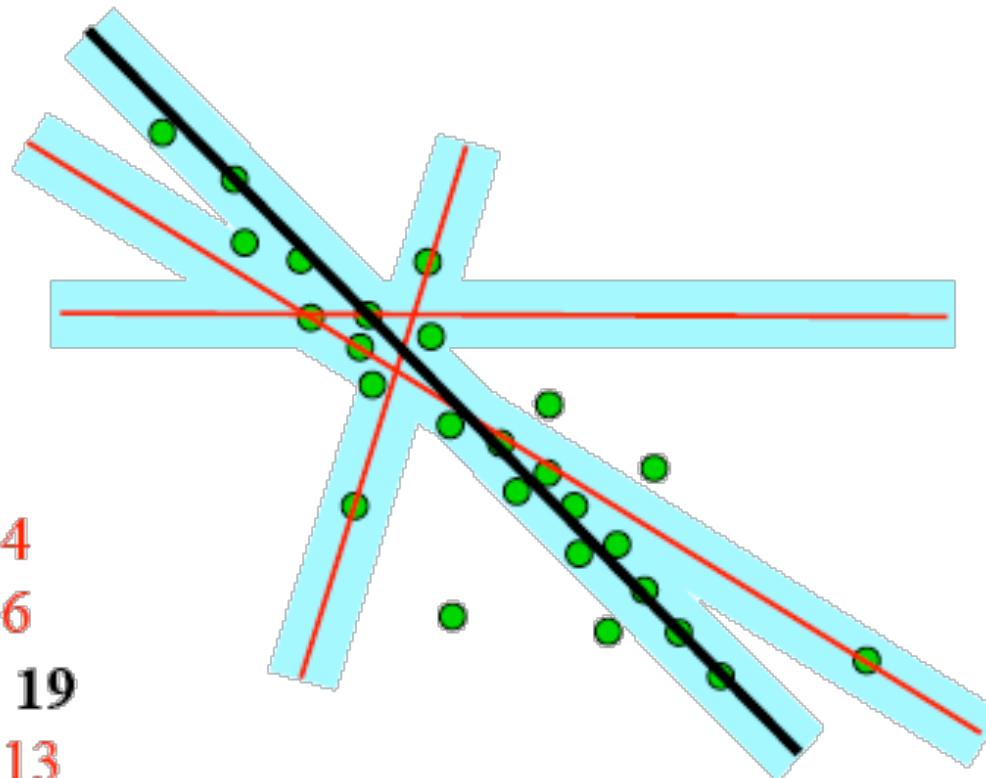
RANSAC procedure (line example)



RANSAC procedure (line example)



RANSAC procedure (line example)



RANSAC

Choose a small subset uniformly at random

Fit to that

Anything that is close to result is signal; all others are noise

Refit

Do this many times and choose the best

ISSUES

- How many times?
 - Often enough that we are likely to have a good line
- How big a subset?
 - Smallest possible
- What does close mean?
 - Depends on the problem
- What is a good line?
 - One where the number of nearby points is so big it is unlikely to be all outliers

Algorithm 15.4: RANSAC: fitting lines using random sample consensus

Determine:

- n — the smallest number of points required
- k — the number of iterations required
- t — the threshold used to identify a point that fits well
- d — the number of nearby points required
to assert a model fits well

Until k iterations have occurred

 Draw a sample of n points from the data

 uniformly and at random

 Fit to that set of n points

 For each data point outside the sample

 Test the distance from the point to the line

 against t ; if the distance from the point to the line
 is less than t , the point is close

 end

 If there are d or more points close to the line

 then there is a good fit. Refit the line using all
 these points.

end

Use the best fit from this collection, using the
fitting error as a criterion

How Many Samples to Choose?

- Probability that a point is an outlier: e
- Number of points in a sample: s
- Number of samples (we want to compute this): N
- Desired probability that we get a good sample: p

$$p = 1 - (1 - (1 - e)^s)^N$$

Let's see why ...

- Probability that a point is an outlier: e
- Number of points in a sample: s
- Number of samples (we want to compute this): N
- Desired probability that we get a good sample: p

$$p = 1 - (1 - (1 - e)^s)^N$$

Probability of choosing one inlier

Let's see why ...

- Probability that a point is an outlier: e
- Number of points in a sample: s
- Number of samples (we want to compute this): N
- Desired probability that we get a good sample: p

$$p = 1 - (1 - (1 - e)^s)^N$$

Probability of choosing s inliers

Let's see why ...

- Probability that a point is an outlier: e
- Number of points in a sample: s
- Number of samples (we want to compute this): N
- Desired probability that we get a good sample: p

$$p = 1 - (1 - (1 - e)^s)^N$$

Probability of that one or more points in the sample were outliers (contaminated sample)

Let's see why ...

- Probability that a point is an outlier: e
- Number of points in a sample: s
- Number of samples (we want to compute this): N
- Desired probability that we get a good sample: p

$$p = 1 - (1 - (1 - e)^s)^N$$

Probability of that N Samples
were contaminated

Let's see why ...

- Probability that a point is an outlier: e
- Number of points in a sample: s
- Number of samples (we want to compute this): N
- Desired probability that we get a good sample: p

$$p = 1 - (1 - (1 - e)^s)^N$$

Probability of that AT LEAST
one sample of N Samples was
NOT contaminated

How Many Samples?

Choose N so that, with probability p , at least one random sample is free from outliers. E.g. $p = 0.99$

$$p = 1 - (1 - (1 - e)^s)^N$$

$$N = \frac{\ln(1 - p)}{\ln(1 - (1 - e)^s)}$$

s	proportion of outliers e							
	5%	10%	20%	25%	30%	40%	50%	
2	2	3	5	6	7	11	17	
3	3	4	7	9	11	19	35	
4	3	5	9	13	17	34	72	
5	4	6	12	17	26	57	146	
6	4	7	16	24	37	97	293	
7	4	8	20	33	54	163	588	
8	5	9	26	44	78	272	1177	

Line example

12 pts: $n = 12$

Sample size: $s = 2$

Outliers 2: $e=1/6 \rightarrow 20\%$

$N = 5$ gives 99% chance of getting a good sample
(trying every possible pair requires 66 trials!)

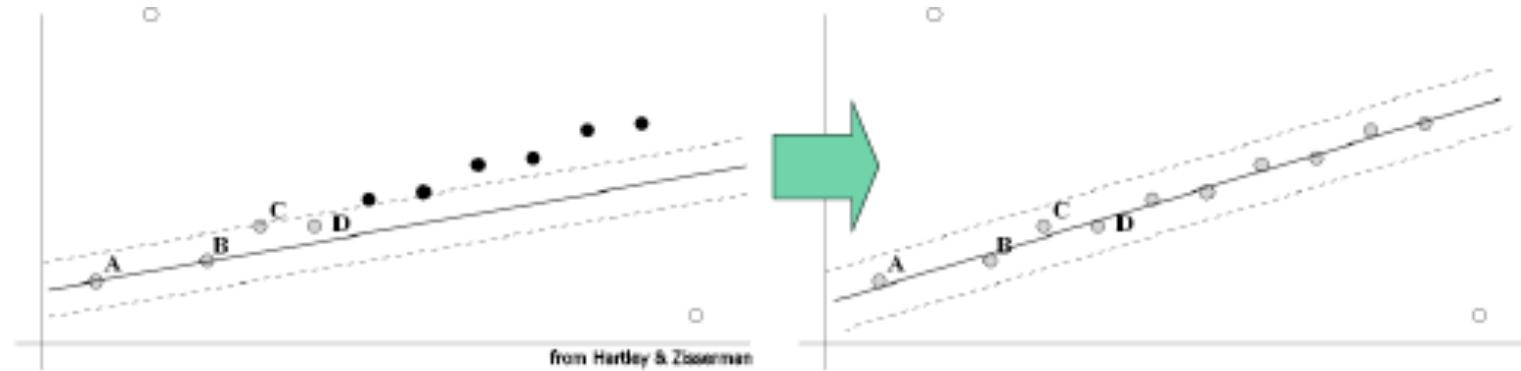
Acceptable Consensus Set

Typically, terminate when inlier ratio reaches expected ratio of inliers

$$T = (1 - e) \times \text{total number of data points}$$

After RANSAC

- RANSAC divides the data into inliers and outliers
- But it computes the estimate with the MINIMAL samples
- We can improve the result by estimating the model using all inliers:
- After RANSAC: estimate once more!



Fitting curves other than lines

In principle, an easy generalization

The probability of obtaining a point, given a curve, is given by a negative exponential of distance squared

- In practice, rather hard
 - It is generally difficult to compute the distance between a point and a curve

Active Contours





Deformable Contours

They are also called
Snakes
Active contours

Think of a snake as an elastic band:
of arbitrary shape
sensitive to image gradient
that can wiggle in the image
represented as a necklace of points

Active Contour Models

An important class of algorithms to find boundaries
Usually does not use prior knowledge of the shape
Poses the problem as an “optimization” problem

Introduction

How can we find the boundary of an object in an image?

One approach could be:

- Find edges

- Link the edges

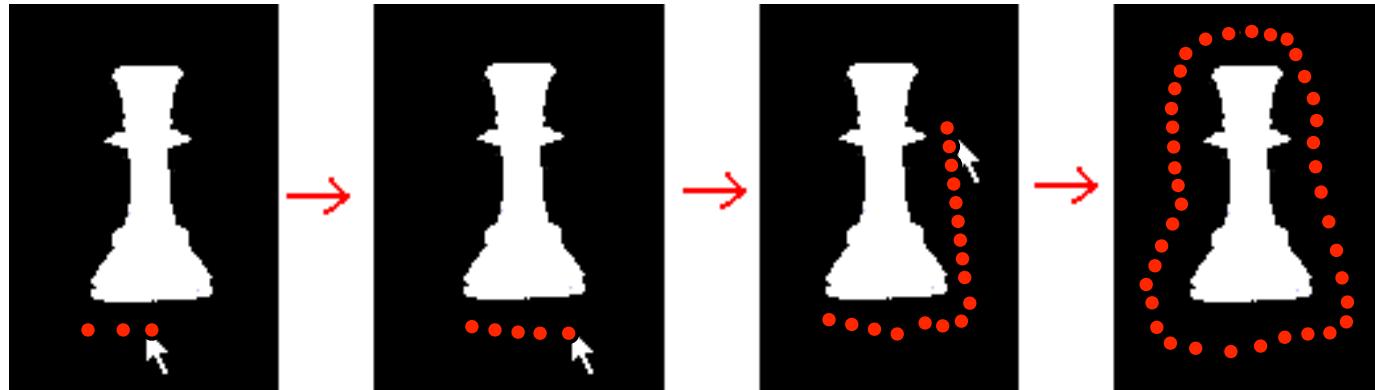
Another possibility is to search for “smooth” boundaries:

- The boundary should “match” the image

- Can iteratively “improve”

Main Idea:

“Drop” a snake



Let the snake “wiggle”, attracted by image gradient, until it glues itself against a contour

The Energy Functional

Associate to each possible shape and location of the snake a value E .

Values should be s.t. the image contour to be detected has the minimum value.

E is called the **energy** of the snake.

Keep wiggling the snake towards smaller values of E .

Energy Functional Design

We need a function that given a snake state, associates to it an Energy value E .

The function should be designed so that the snake moves towards the contour that we are seeking!

What moves the snake?

“Forces” applied to its points

Snake Energy

The total energy of the snake is defined as:

$$E_{total} = E_{internal} + E_{external}$$

The internal energy encourages smoothness

The external energy encourages closeness to edges

Forces moving the snake (External)

It needs to be attracted to contours:

Edge pixels must “pull” the snake points.

The stronger the edge, the stronger the pull.

The force is proportional to $|\nabla|$

Forces preserving the snake (Internal)

The snake should not break apart!

Points on the snake must stay close to each other

Each point on the snake pulls its neighbors

The farther the neighbors, the stronger the force

The force is proportional to the distance $|P_i - P_{i-1}|$

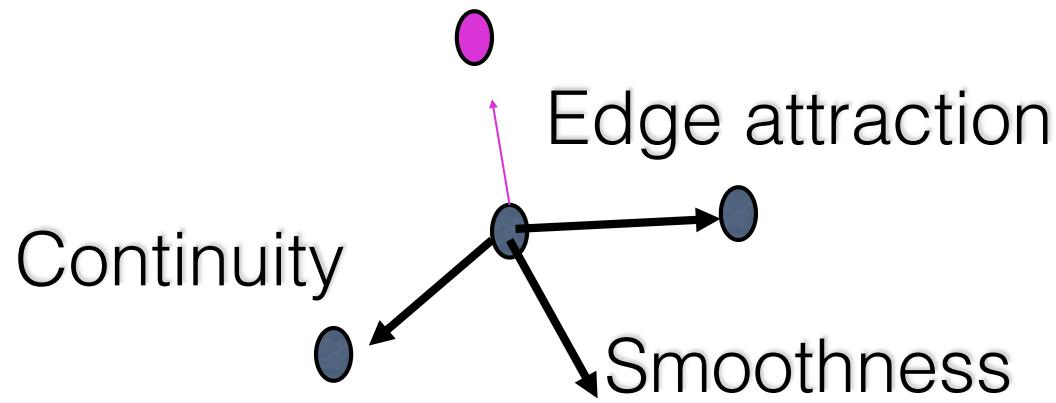
Forces preserving the snake (Internal)

The snake should avoid “oscillations”

Penalize high curvature

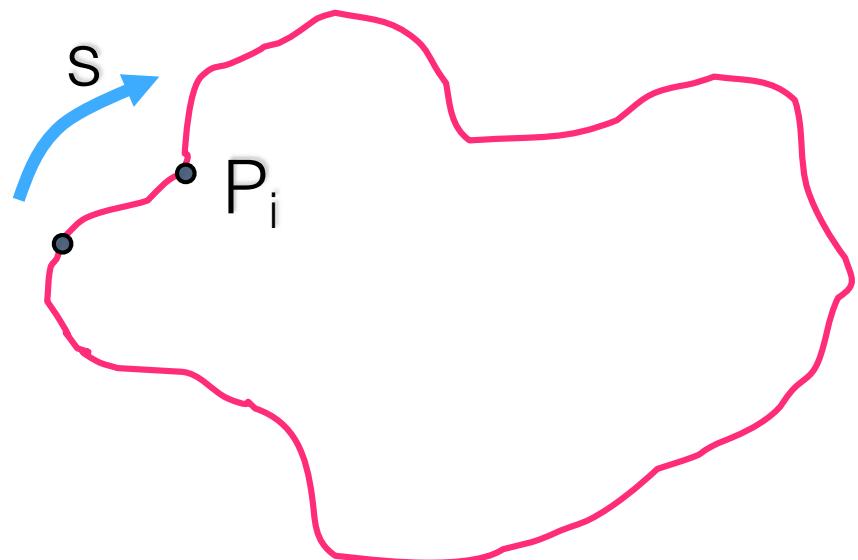
Force proportional to snake curvature

Snake Forces



Snake “state”: Contour Parametrization

Consider a contour parametrization $c=c(s)$ where s is the “arc length”



Each point P_i on the contour has coordinates $(x_i(s), y_i(s))$

$$E = \int E_{int}(s) + E_{ext}(s) ds$$

Snake Energy Functional

Given a snake with N points p_1, p_2, \dots, p_N

$$E = \sum_{i=1}^N a_i E_c(p_i) + b_i E_s(p_i) + c_i E_g(p_i)$$

“Continuity” “Smoothness” “Edgeness”

a_i, b_i, c_i are “weights” to control influence

Continuity Term

Given a snake with N points p_1, p_2, \dots, p_N

Let d be the average distance between points

Distance between points should be kept close to average

Define the continuity term of the Energy Functional:

$$E_c(p_i) = (d - |p_i - p_{i-1}|)^2$$

$$p_i = [x_i \ y_i]$$

$$E_c = \left(d - \sqrt{(x_i - x_{i-1})^2 + (y_i - y_{i-1})^2} \right)^2$$

Smoothness Term

Given a snake with N points p_1, p_2, \dots, p_N

Curvature should be kept small

Define the smoothness term of the Energy Functional:

$$E_s(p_i) = \underbrace{|p_{i-1} - 2p_i + p_{i+1}|^2}_{\text{Second derivative}}$$

Second derivative

$$E_s = (x_{i-1} - 2x_i + x_{i+1})^2 + (y_{i-1} - 2y_i + y_{i+1})^2$$

Edgeness Term

Given a snake with N points p_1, p_2, \dots, p_N

Define the edgeness term of the Energy Functional:

$$E_g(p_i) = -|\nabla I(p_i)|$$

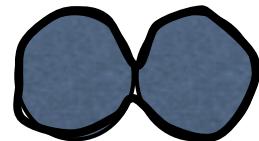
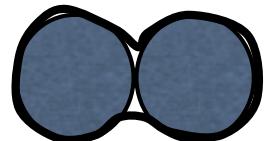
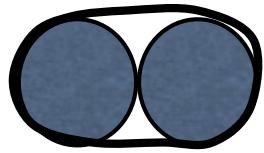
$$\nabla I(p_i) = [G_x(p_i) \ G_y(p_i)]$$

$$|\nabla I(p_i)| = \sqrt{G_x(p_i)^2 + G_y(p_i)^2}$$

Magnitude of the gradient should be LARGE

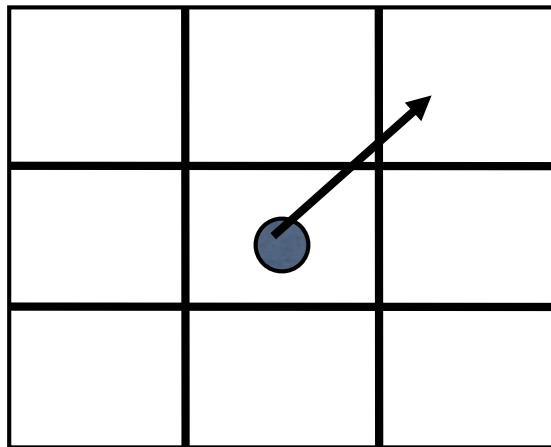
Relative Weighting

The weights control the smoothness and stiffness of the snake



Greedy Algorithm

Each point moves within a small window to minimize the energy



Compute the new energy for each candidate location
Move the point to the one with the minimum value

Keeping corners ...

Before starting a new iteration:

Search for “corners”:

max curvature

large gradient

Corner points should not contribute to the energy (set $b_i = 0$)

Implementation Considerations

To avoid numerical problems, the terms of the energy function should be normalized.

E_c and E_s are normalized by their maximum in the neighborhood

E_g is normalized as $|\nabla| - m| / (M - m)$

M and m are the max and min value of the gradient magnitude in the neighborhood

Snake Algorithm

Input:

gray scale image I

a chain of points p_1, p_2, \dots, p_N

f is the fraction of points that must move to start a new iteration

$U(p)$ is a neighborhood around p

d is the average distance between snake points.

Snake Algorithm

1. While the fraction of moved points > f
 1. For $i=1,2,\dots,N$
 1. find a point in $U(p_i)$ s.t. the energy is minimum,
 2. move p_i to this location
 2. For $i=1,2,\dots,N$
 1. Estimate the curvature $k = |p_{i-1} - 2p_i + p_{i+1}|$
 2. Look for local max, and set $b_{\max} = 0$
 3. Update d

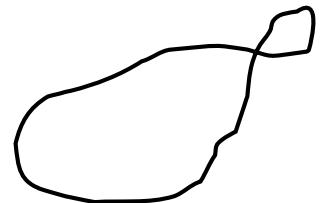
Problems with Snakes

Smoothness does not always capture all prior knowledge

User must define the weights

Snakes might oversmooth boundaries

Not trivial to prevent curve self intersecting



Region Segmentation

Regions and Edges

Ideally, regions are bounded by closed contours

We could “fill” closed contours to obtain regions

We could “trace” regions to obtain edges

Unfortunately, these procedures rarely produce satisfactory results.



Regions and Edges

Edges are found based on DIFFERENCES between values of adjacent pixels.

Regions are found based on SIMILARITIES between values of adjacent pixels.

Region Segmentation

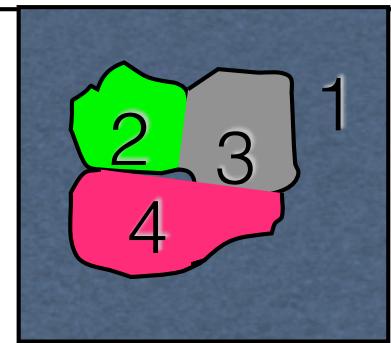
A segmentation is a partition R_1, R_2, \dots, R_n s.t.:

$$\bigvee_{i=1}^n R_i = I$$

$$R_i \wedge R_j = \emptyset \text{ If } i \neq j$$

$\text{Pred}(R_i) = \text{TRUE}$ For all i

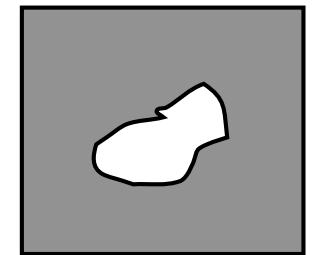
$\text{Pred}(R_i \vee R_j) = \text{FALSE}$ for all R_i adjacent R_j



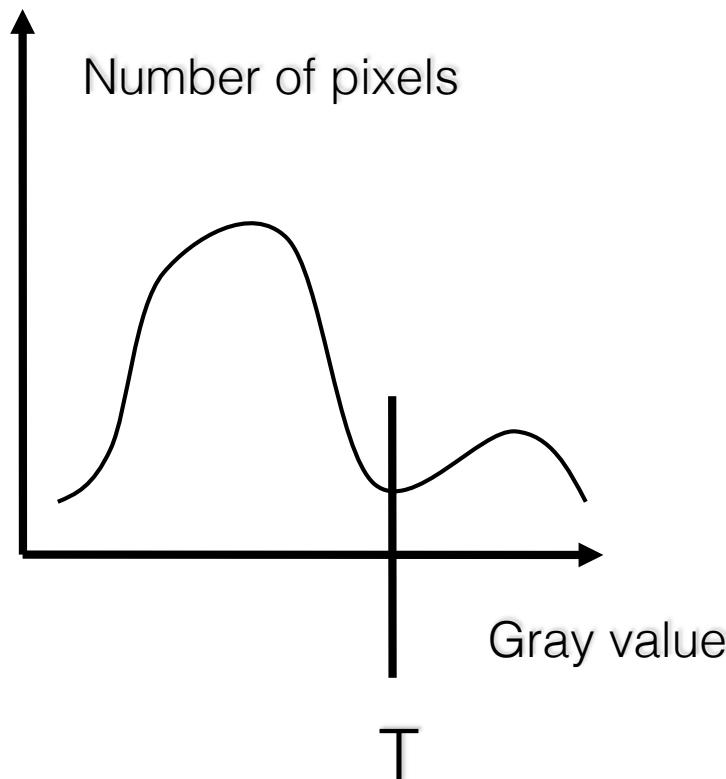
Where “Pred” is a function that evaluates similarities of the pixels in the region

Histogram-based Segmentation

Ex: bright object on dark background:



Histogram



Select threshold
Create binary image:
 $I(x,y) < T \rightarrow O(x,y) = 0$
 $I(x,y) > T \rightarrow O(x,y) = 1$

How do we select a Threshold?

Automatic thresholding

P-tile method

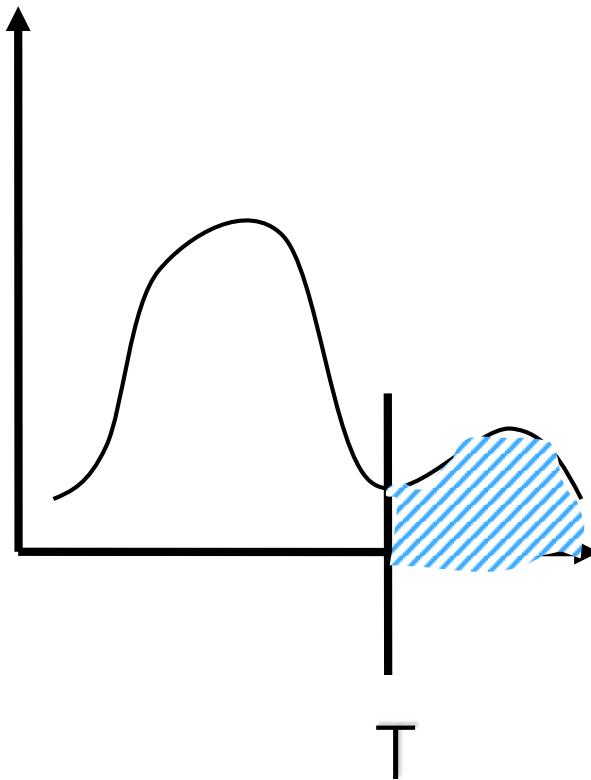
Mode method

Peakiness detection

Iterative algorithm

P-Tile Method

If the size of the object is approx. known, pick T s.t. the area under the histogram corresponds to the size of the object:

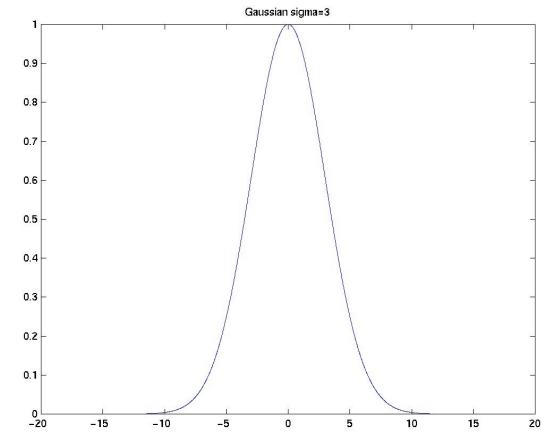


Mode Method

Model each region as “constant” + $N(0, \sigma_i)$:

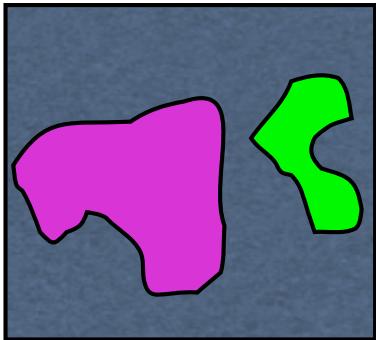
If $(x, y) \in R_i$ then, $I(x, y) = \mu_i + n_i(x, y)$

$$p(n_i) = \frac{1}{\sqrt{2\pi}\sigma_i} e^{-\frac{1}{2}\frac{n_i^2}{\sigma_i^2}}$$

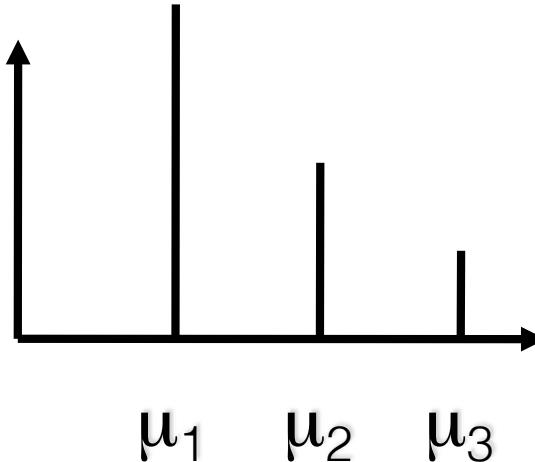


$$E(n_i) = 0 \quad E(n_i^2) = \sigma_i^2$$

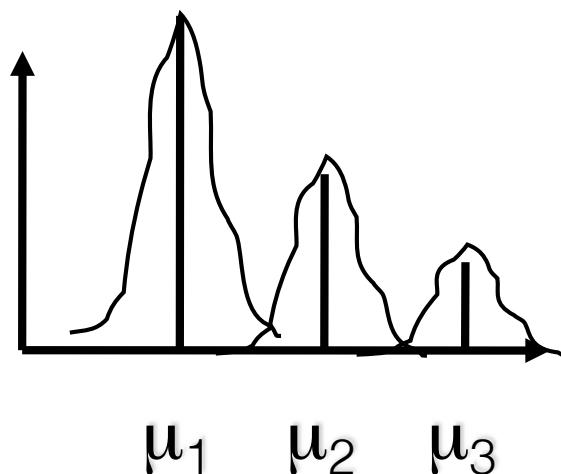
Example: Image with 3 regions



Ideal histogram:



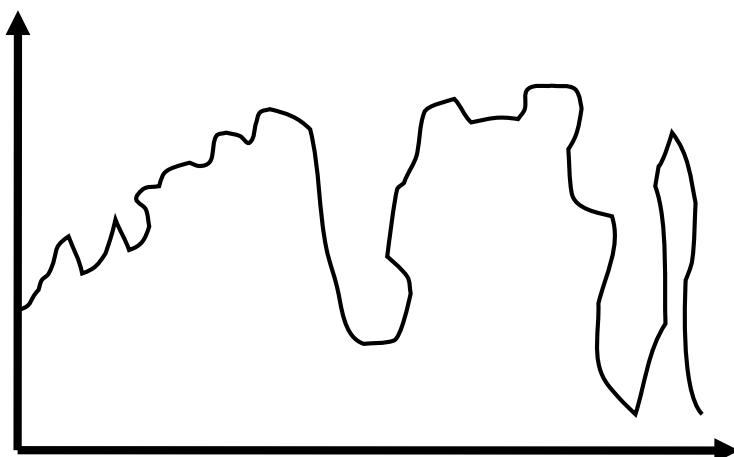
Add noise:



The valleys are good places for thresholding to separate regions.

Finding the peaks and valleys

It is a not trivial problem:



“Peakiness” Detection Algorithm

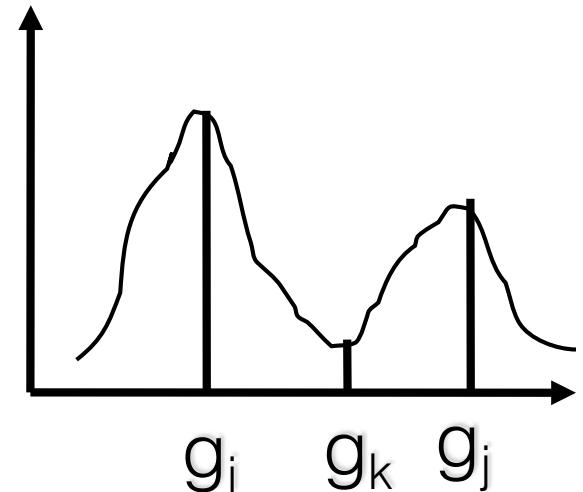
Find the two **HIGHEST LOCAL MAXIMA** at a **MINIMUM DISTANCE APART**: g_i and g_j

Find **lowest point** between them: g_k

Measure “peakiness”:

$$\min(H(g_i), H(g_j))/H(g_k)$$

Find (g_i, g_j, g_k) with highest peakiness



Iterative Threshold Algorithm

1. Select initial $T = T_0$
Ex: T_0 = average intensity
2. Partition image into regions R_1 and R_2 using T
3. Calculate the mean values of R_1 and R_2 , μ_1 and μ_2
Select a new threshold $T = \frac{1}{2}(\mu_1 + \mu_2)$
1. Repeat 2-4 until μ_1 and μ_2 do not change.

Algorithm MEAN SHIFT

A non-parametric technique

Finds the peak of a given histogram

It is based on Robust Statistics

See: “**Robust Analysis of Feature Space: Color Image Segmentation,**” by D. Comaniciu and P. Meer, CVPR 1997, pp. 750-755.

Algorithm MEAN SHIFT to find histogram PEAK

1. Choose a window size
2. Choose the initial location of the search window
3. Compute the mean location in the search window
4. Center the window at the location computed in 3
5. Repeat steps 3 and 4 until convergence.

Mathematical Justification

Assume first that the histogram is unimodal

Let y be the possible gray values

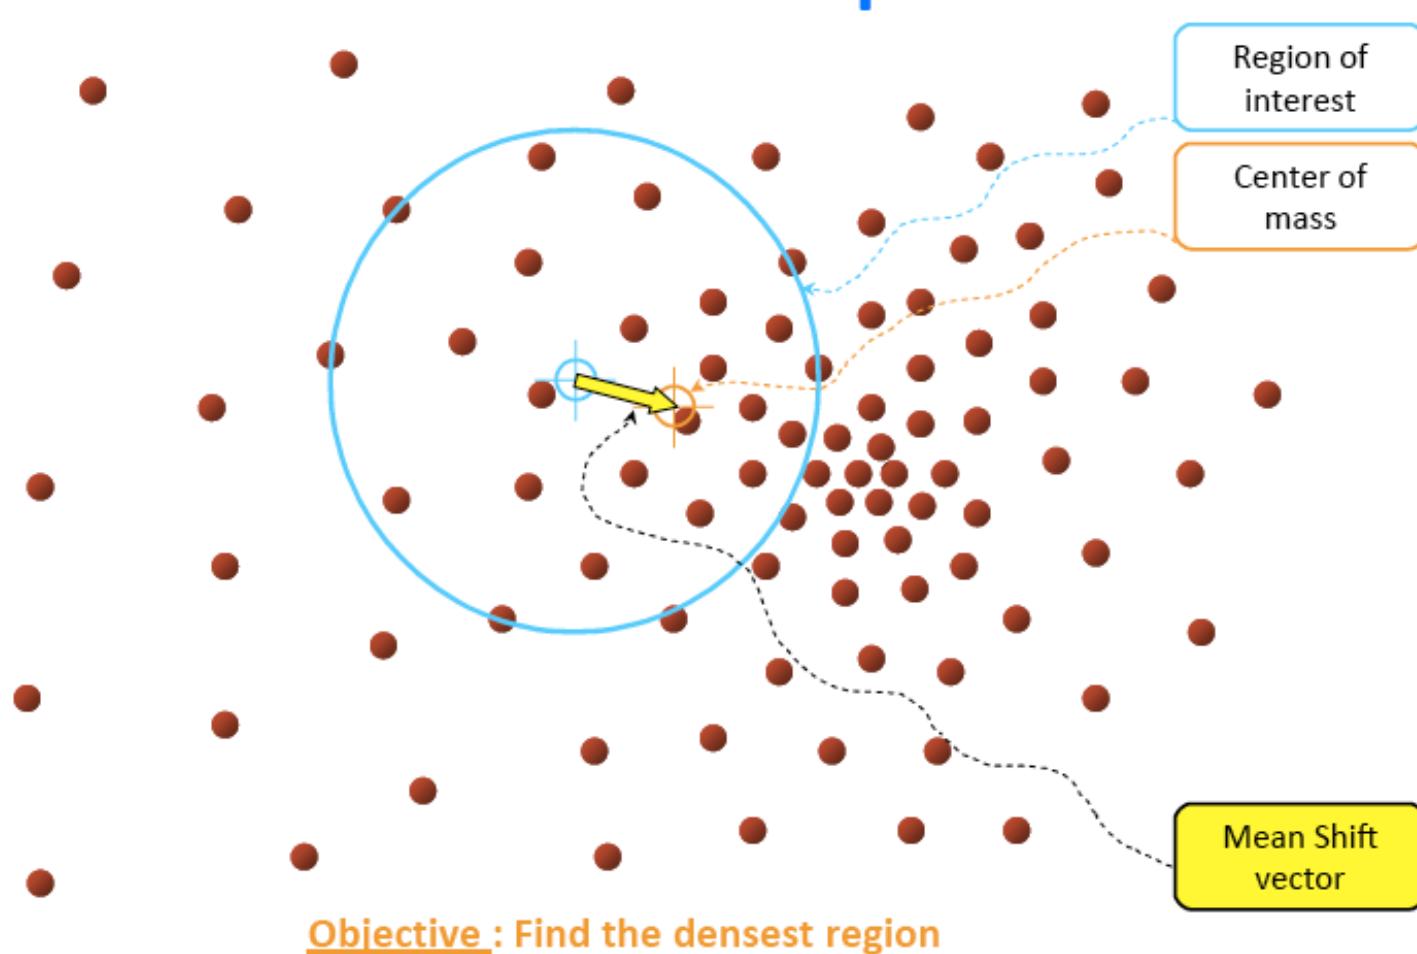
Let $p(y)$ be the normalized histogram \sim pdf

$$p(y) = \frac{\text{pixels with value } y}{\text{pixels in image}}$$

Mean Shift

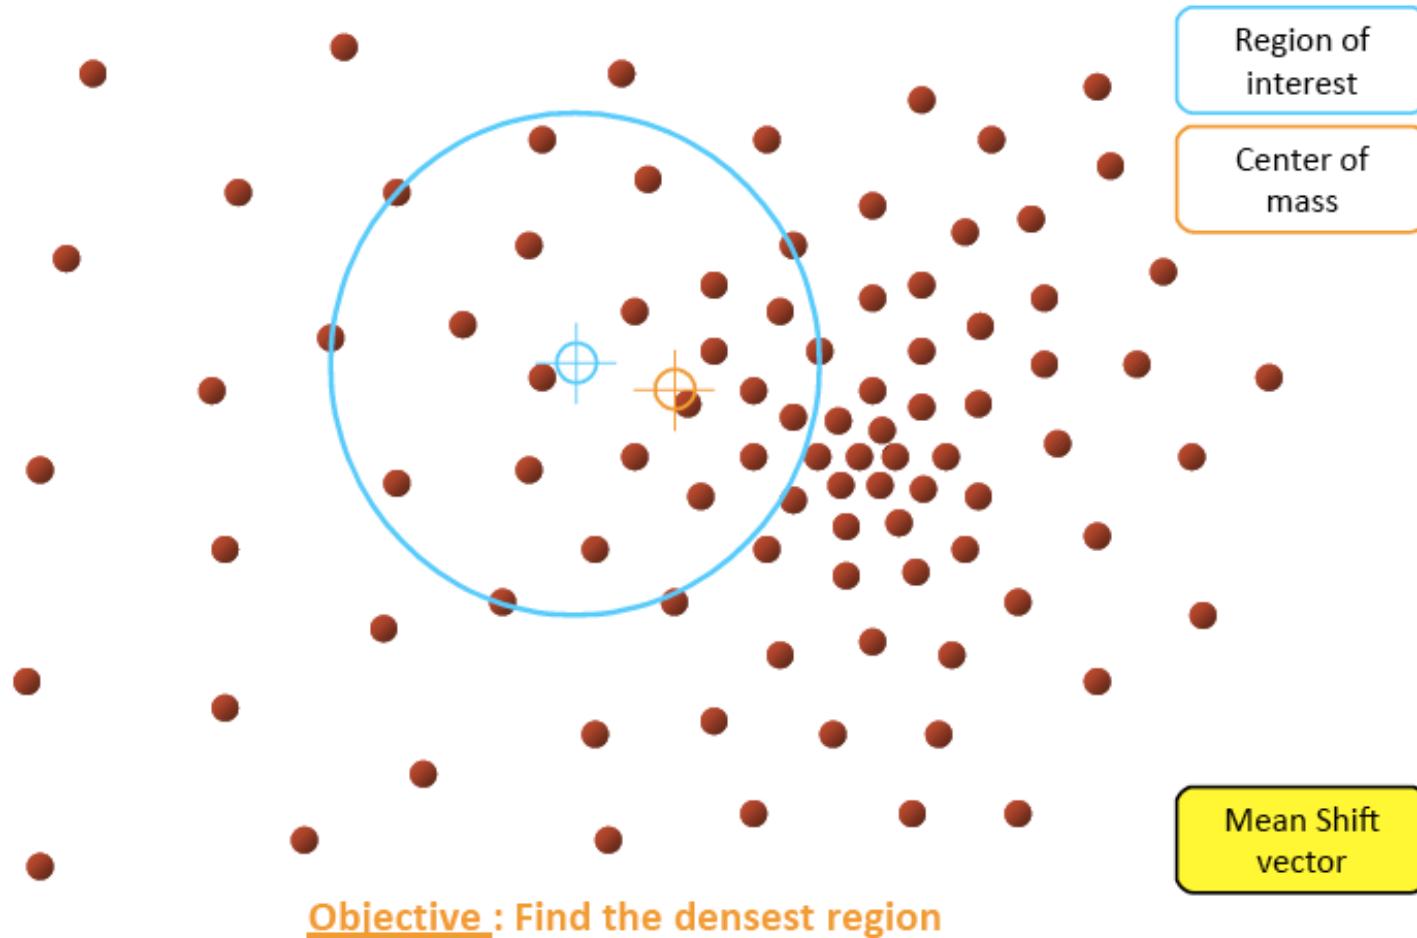
It is a hill-climbing algorithm that seeks modes of a non-parametric density represented by samples.

Intuitive Description



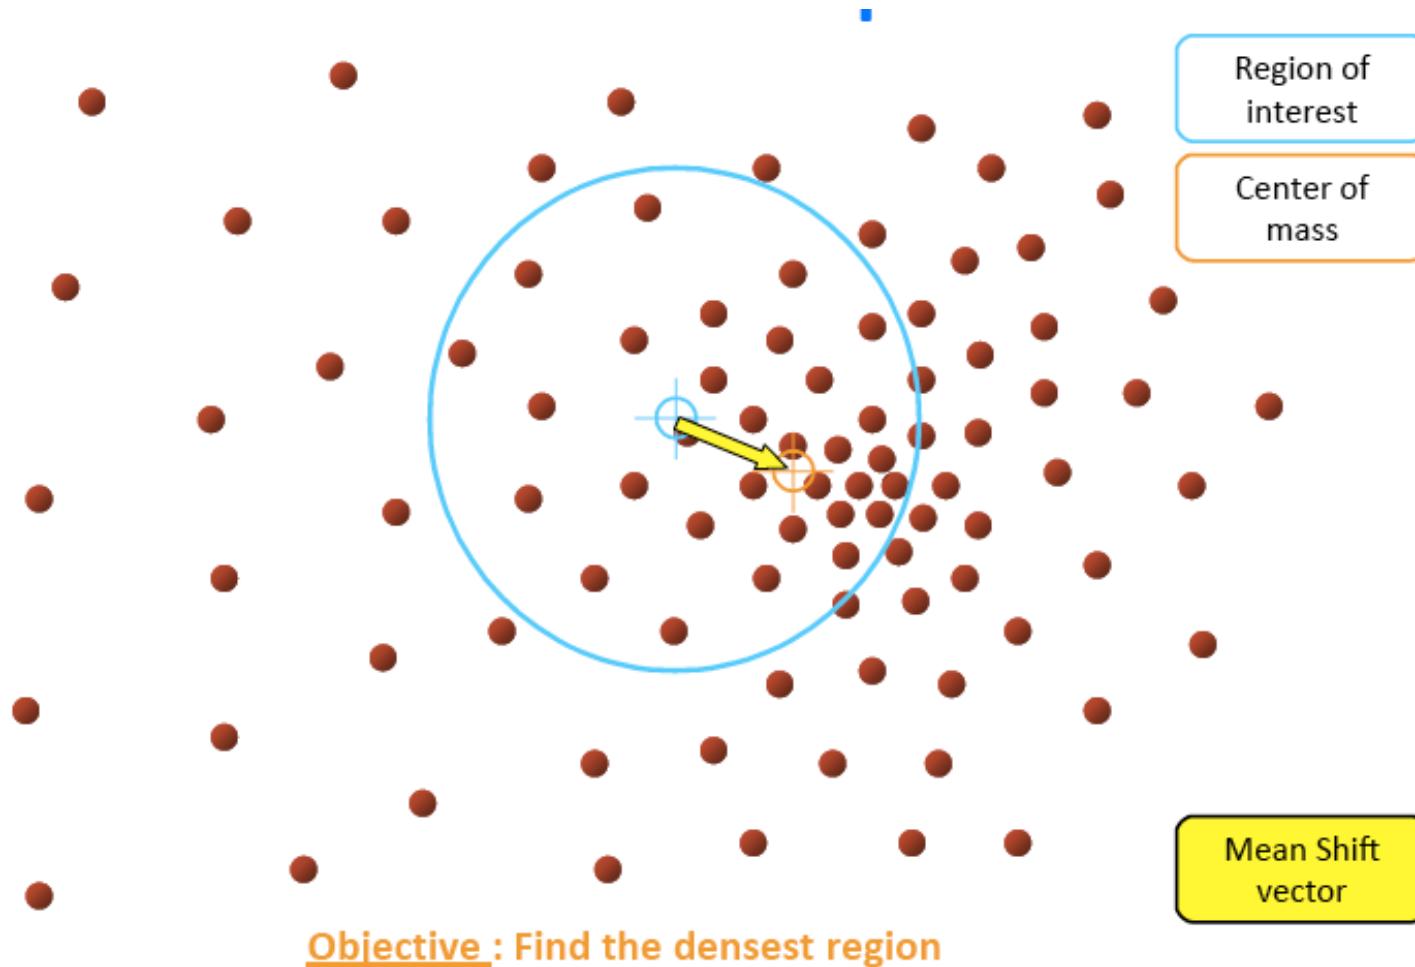
Ukrainitz&Sarel, Weizmann

Intuitive Description



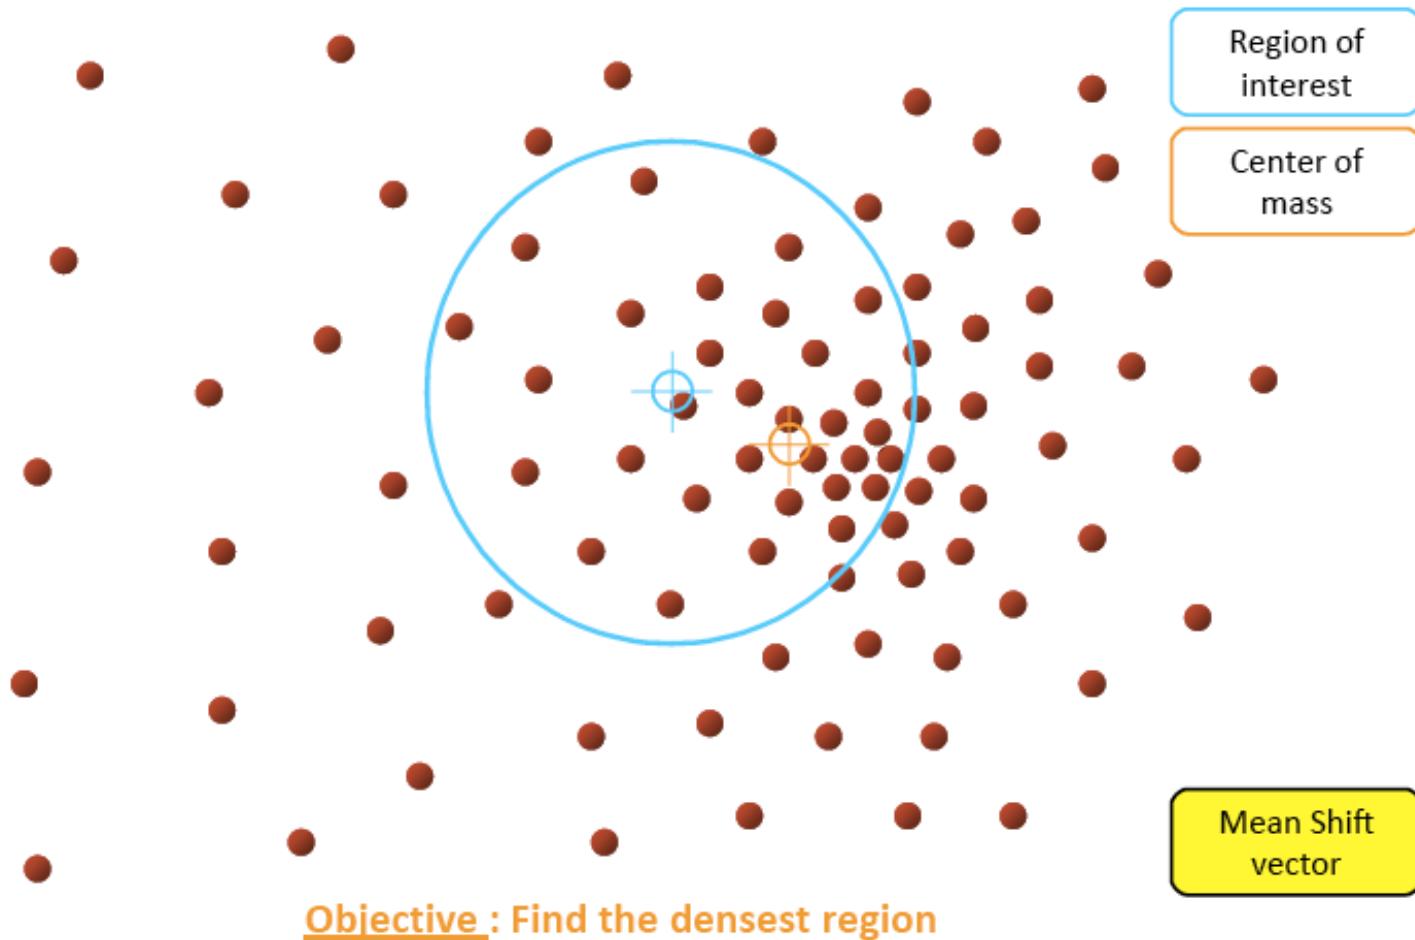
Ukrainitz&Sarel, Weizmann

Intuitive Description



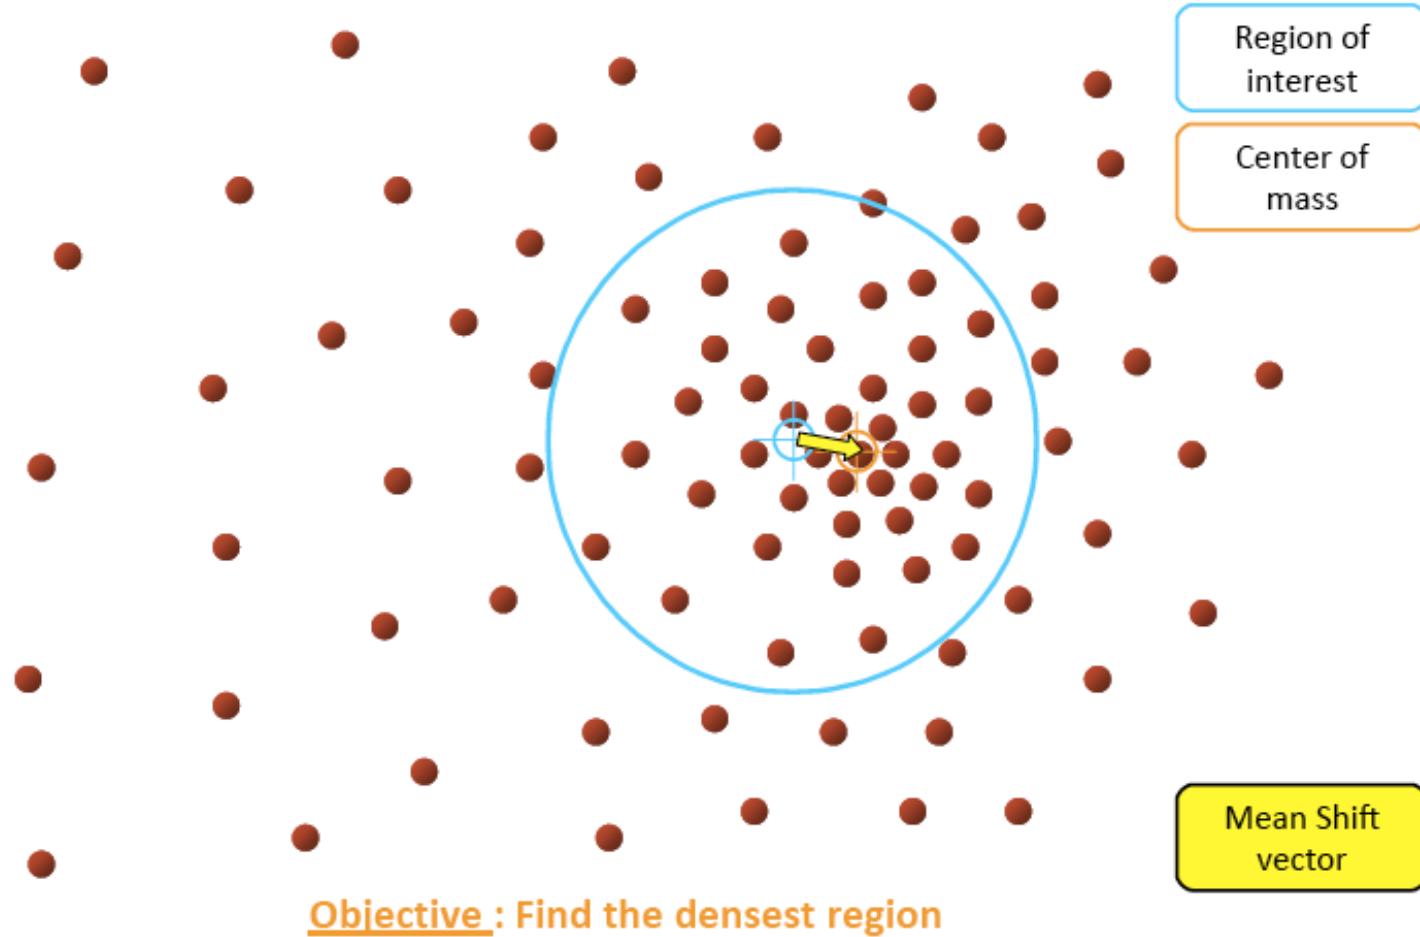
Ukrainitz&Sarel, Weizmann

Intuitive Description



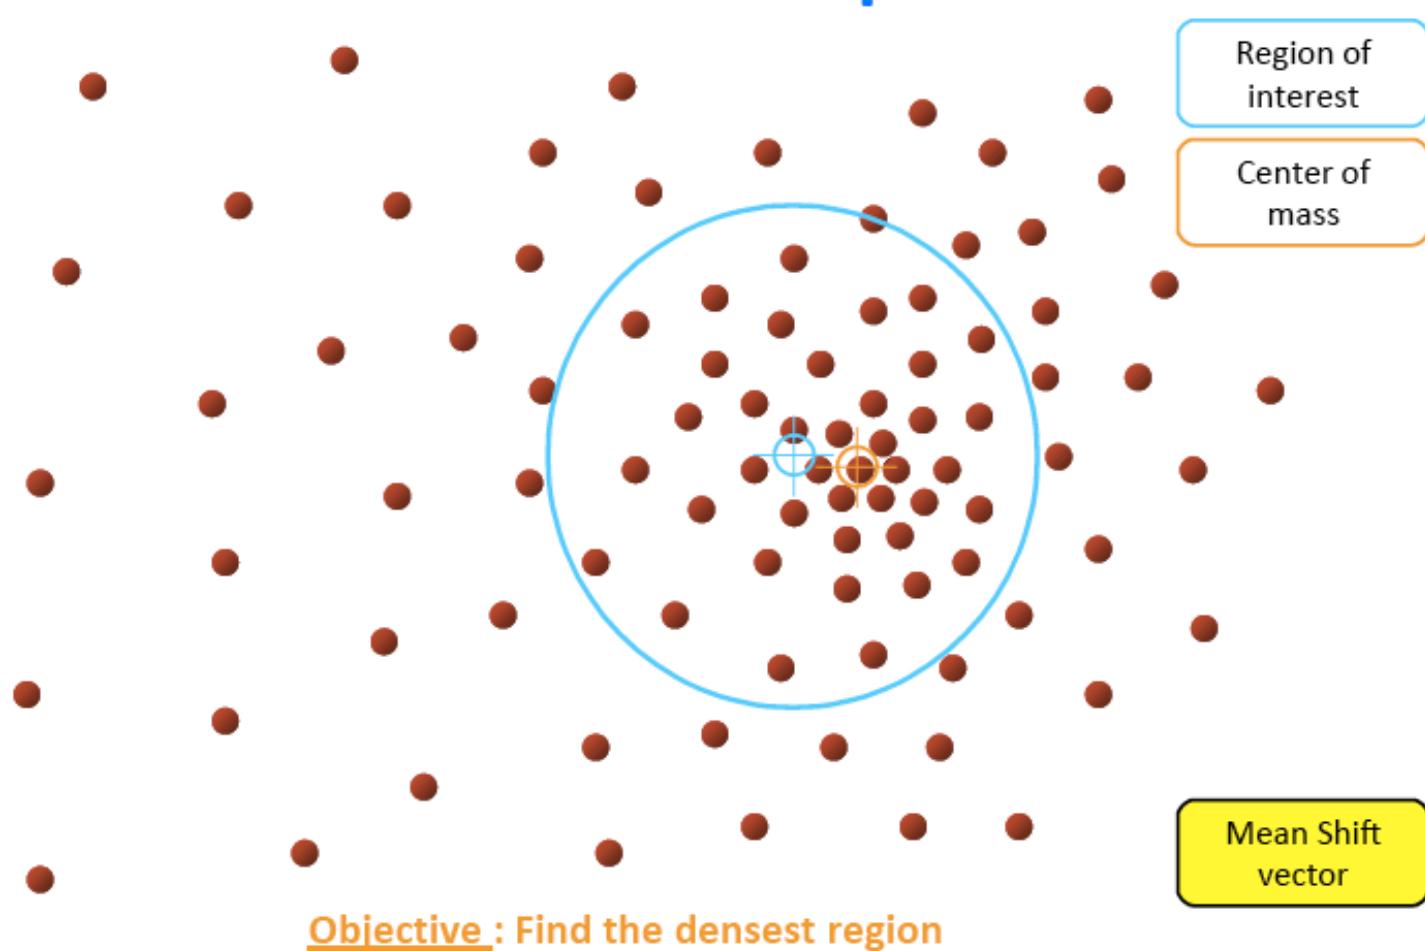
Ukrainitz&Sarel, Weizmann

Intuitive Description



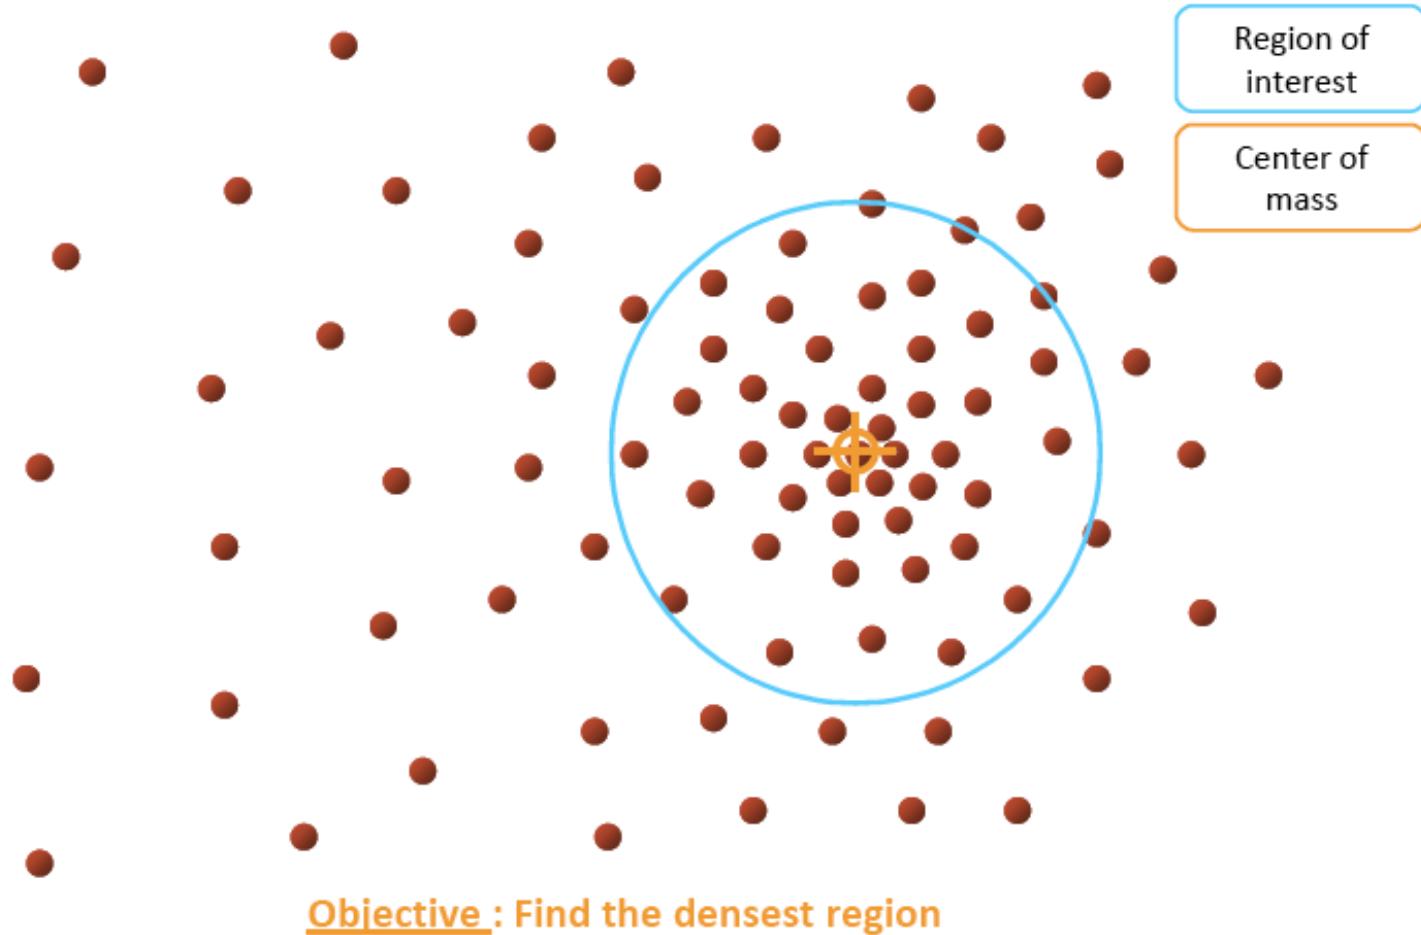
Ukrainitz&Sarel, Weizmann

Intuitive Description

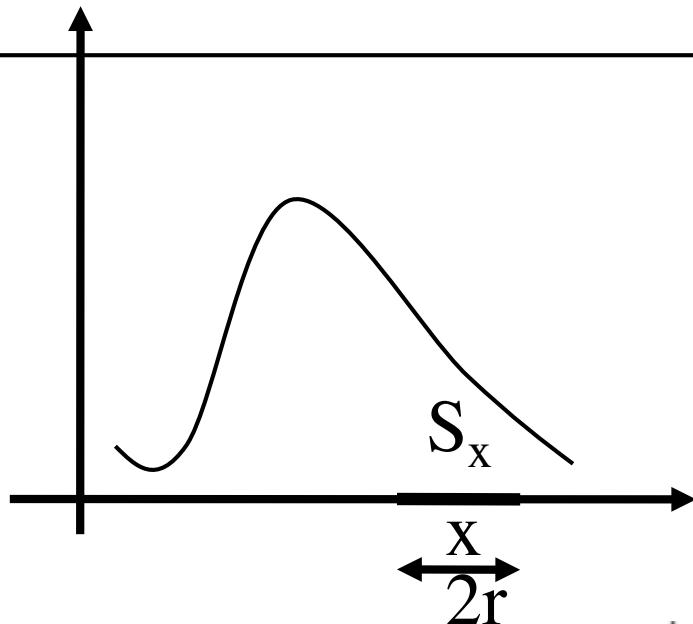


Ukrainitz&Sarel, Weizmann

Intuitive Description



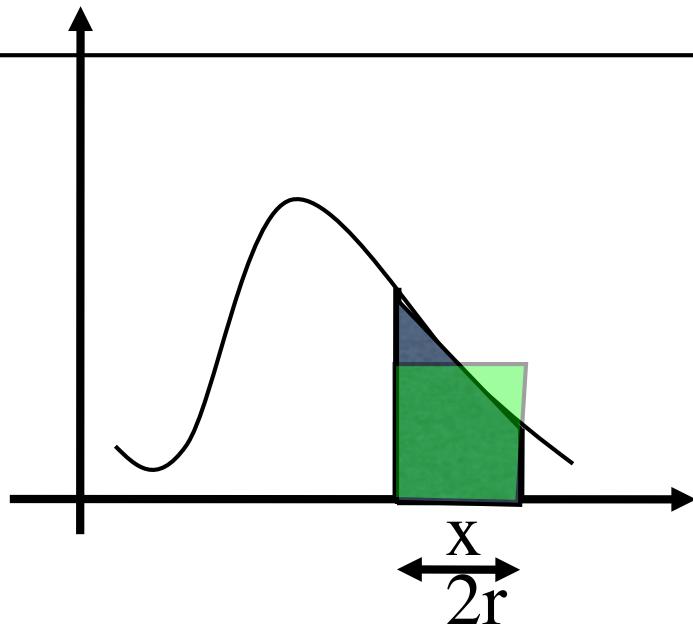
Ukrainitz&Sarel, Weizmann



Define $z = y - x$

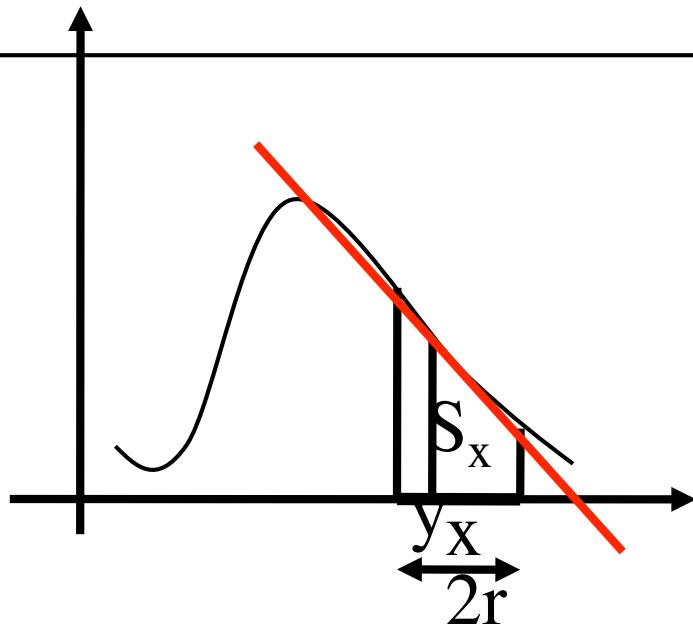
$$E[z|S_x] = \int_{x-r}^{x+r} (y - x)p(y|S_x)dy$$

$$E[z|S_x] = \int_{x-r}^{x+r} (y - x) \frac{p(y)}{p(y \in S_x)} dy$$



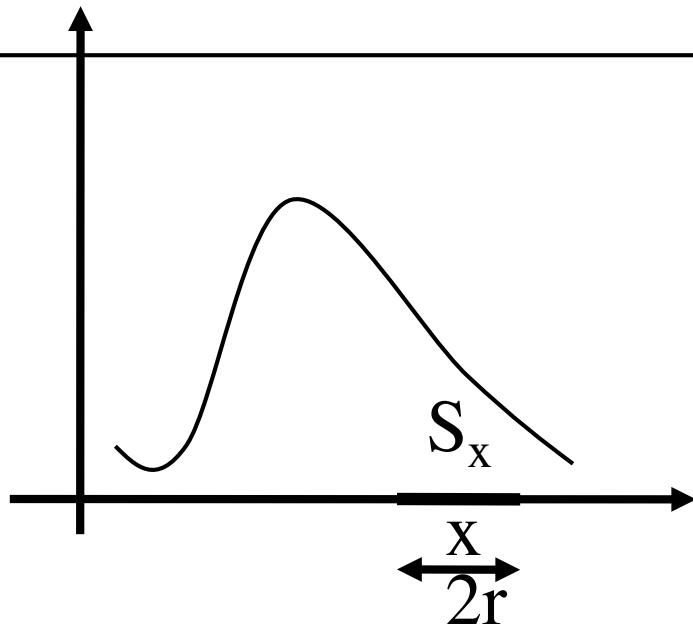
Define $z = y - x$

$$p(y \in S_x) = \int_{x-r}^{x+r} p(y) dy \approx p(x)2r$$



Define $z = y - x$

$$p(y) \approx p(x) + (y - x)p'(x)$$

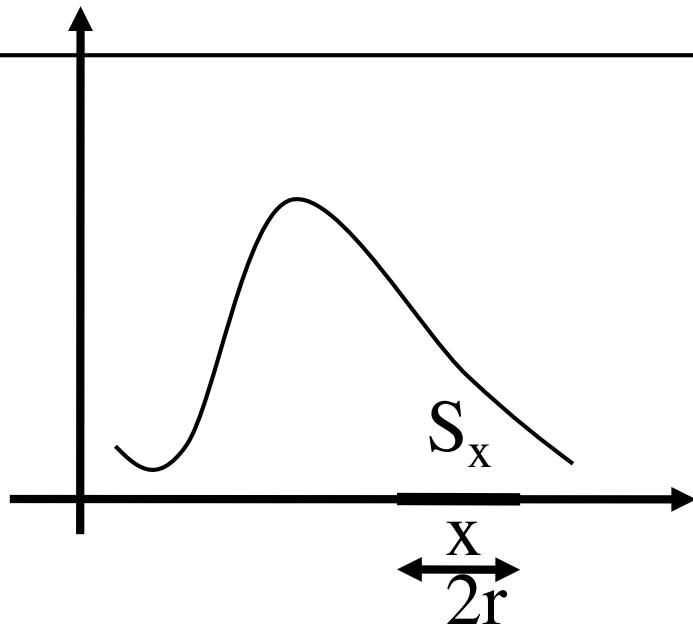


Define $z = y - x$

$$E[z|S_x] = \int_{x-r}^{x+r} (y - x) \frac{p(y)}{p(y \in S_x)} dy$$

$$p(y \in S_x) \approx p(x)2r$$

$$E[z|S_x] \approx \frac{1}{2rp(x)} \int_{x-r}^{x+r} (y - x)p(y)dy$$

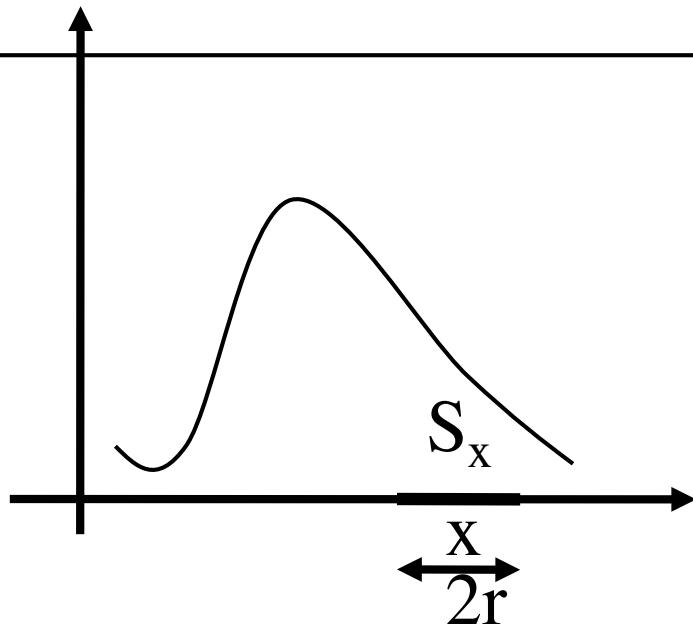


Define $z = y - x$

$$E[z|S_x] \approx \frac{1}{2rp(x)} \int_{x-r}^{x+r} (y - x)p(y)dy$$

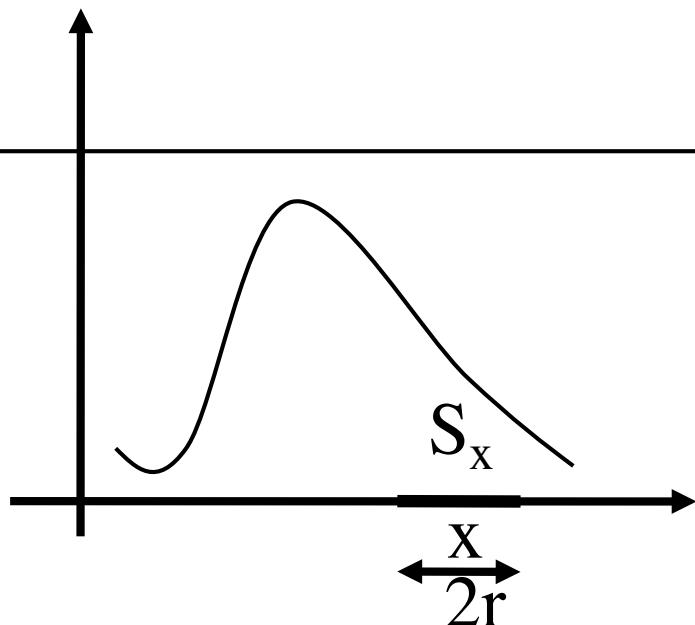
$$p(y) \approx p(x) + (y - x)p'(x)$$

$$E[z|S_x] \approx \frac{1}{2rp(x)} \int_{x-r}^{x+r} (y - x)[p(x) + (y - x)p'(x)]dy$$



Define $z = y - x$

$$\begin{aligned}
 E[z|S_x] &\approx \frac{1}{2rp(x)} \int_{x-r}^{x+r} (y-x)[p(x) + (y-x)p'(x)]dy \\
 &= \frac{1}{2rp(x)} \int_{-r}^{+r} (z)[p(x) + (z)p'(x)]dz \\
 &= \frac{1}{2rp(x)} \left[\frac{1}{2}z^2p(x) + \frac{1}{3}z^3p'(x) \right]_{-r}^{+r} \\
 &= \frac{1}{2rp(x)} \frac{2}{3}r^3p'(x) = \frac{r^2p'(x)}{3p(x)}
 \end{aligned}$$



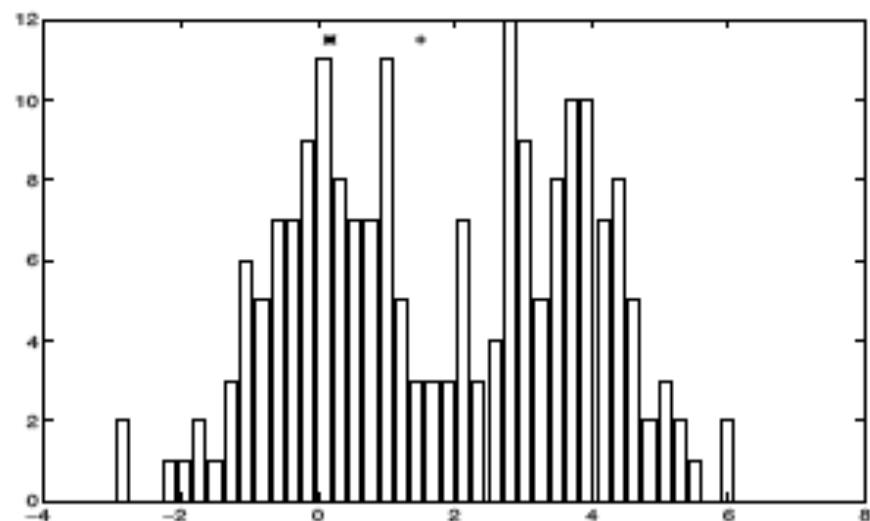
Define $z = y - x$

$$E(z|S_x) = \mu \approx \frac{r^2 p'(x)}{3p(x)}$$

μ is the “mean shift vector”: the vector between the local mean $E[y|y \text{ in } S_x]$ and the center of the window x

- μ is proportional to $p'(x)$
- The factor of proportionality is reciprocal to $p(x)$
- Peaks have large $p(x)$ but small $p'(x)$
- If not at the peak, μ is in the direction of the peak

Example: Comaniciu & Meer



200 data pts
 $N(0,1)$ & $N(3.5,1)$

Mode detector:
Center of shortest window
with half of the data points

Figure 1: An example of the mean shift algorithm.

Initial Mode	Initial Mean	Final Mean
1.5024	1.4149	0.1741

Initial window size = 3.2828

Example

#	0		3	5	3			
gray	0		2	3	4	5	6	7

$$N = 0 + 1 + 3 + 5 + 3 + 1 + 1 + 1 = 15$$

$$R = 1.5, x_0 = 5$$

$$4 \leq y \leq 6$$

$$\mu = (4*3 + 5*1 + 6*1) / (3 + 1 + 1) - 5 = 4.6 - 5 = -0.4$$

$$R = 1.5, x_1 = 4.6$$

Example

#	0		3	5	3			
gray	0		2	3	4	5	6	7

$$N = 0 + 1 + 3 + 5 + 3 + 1 + 1 + 1 = 15$$

$$R = 1.5, x_1 = 4.6$$

$$3 \leq y \leq 6$$

$$\mu = (3*5 + 4*3 + 5*1 + 6*1) / (5 + 3 + 1 + 1) - 4.6 = 3.8 - 4.6 = -0.8$$

$$R = 1.5, x_2 = 3.8$$

Example

#	0	1	3	5	3	1	1	1
gray	0	1	2	3	4	5	6	7

$$N = 0 + 1 + 3 + 5 + 3 + 1 + 1 + 1 = 15$$

$$R = 1.5, x_2 = 3.8$$

$$2 \leq y \leq 5$$

$$\mu = (2*3 + 3*5 + 4*3 + 5*1) / (3 + 5 + 3 + 1) - 3.8 = 3.16 - 3.8 = -0.3$$

$$R=1.5, x_3 = 3.16$$

Example

#	0	1	3	5	3	1	1	1
gray	0	1	2	3	4	5	6	7

$$N = 0 + 1 + 3 + 5 + 3 + 1 + 1 + 1 = 15$$

$$R = 1.5, x3 = 3.16$$

$$2 \leq y \leq 5$$

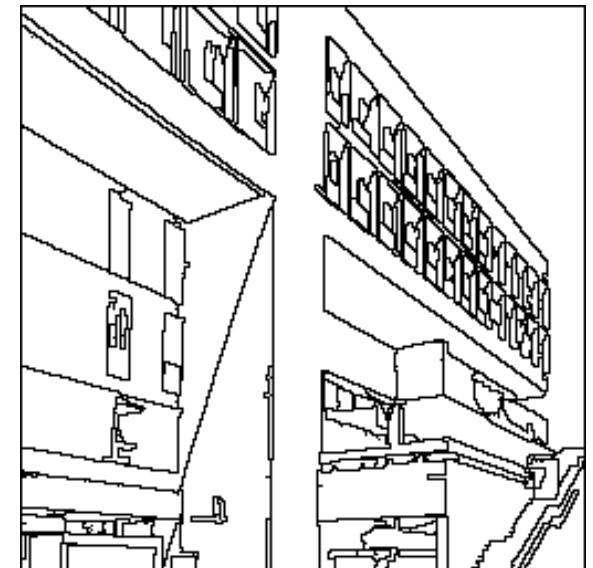
$$\mu = (2*3 + 3*5 + 4*3 + 5*1) / (3 + 5 + 3 + 1) - 3.16 = 3.16 - 3.16 = 0$$

Algorithm MEAN SHIFT for Image Segmentation

1. Find image histogram, choose window size
2. Choose initial location of search window:
 1. Randomly select a number M of image pixels
 2. Find the average value in a 3×3 window for each of these pixels
 3. Set the center of the window to the value with largest histogram count.
3. Apply mean shift to find the window peak.
4. Remove pixels in the window from the image and the histogram.
5. Repeat steps 2 to 4 until no pixels are left.

Example of Mean Shift Segmentation

by D. Comaniciu and P. Meer



Limitations of histogram methods:

Use **GLOBAL** information

Ignore **SPATIAL** relationships among pixels.

SPLIT & MERGE Algorithms

Simple intensity algorithms usually result in too many regions.

Reasons:

- high frequency noise

- Gradual transitions between regions

After segmentation, regions might need refinement:

- Interactively or automatically

- May use domain and/or image process knowledge

Merging Algorithm

Merge **ADJACENT, SIMILAR** regions

What does “similar” mean?

“similar” average values : $|μ_i - μ_j| < T$

“small” spread of gray values: $|g_{max} - g_{min}| < T$

$$g_{max} = \max\{g(x,y) \mid (x,y) \text{ in } R_i \cup R_j\}$$

$$g_{min} = \min\{g(x,y) \mid (x,y) \text{ in } R_i \cup R_j\}$$

Note:

A similar to B, and B similar to C does not imply that A is similar to C.

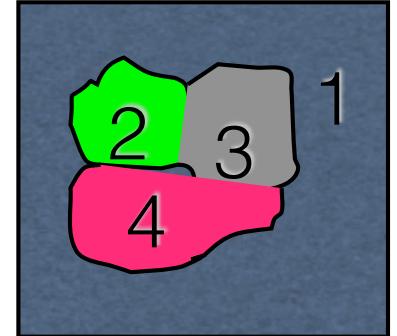
Merging Algorithm

Start with an initial segmentation

Ex:

By thresholding,
nxn (5x5, 7x7, etc) regions
manually selected

Each region has a unique “label”

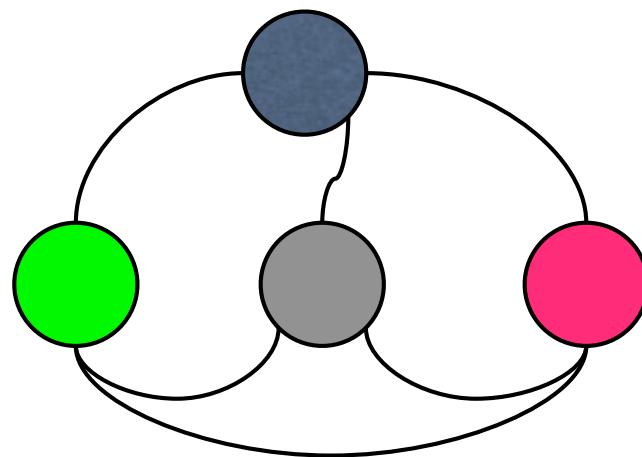
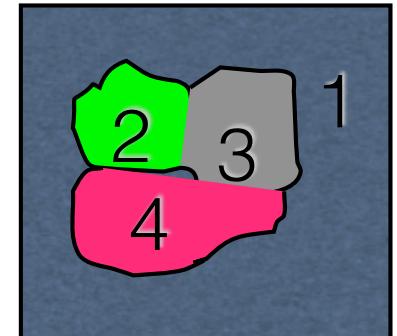


Merging Algorithm

Form the Region Adjacency Graph

Regions are the nodes

Adjacency relations are the links

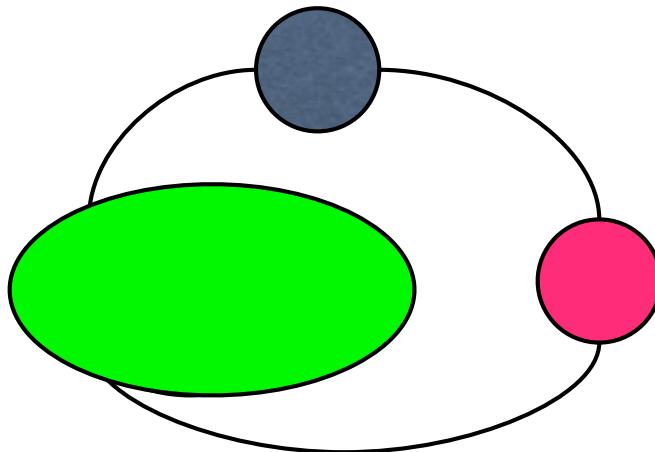
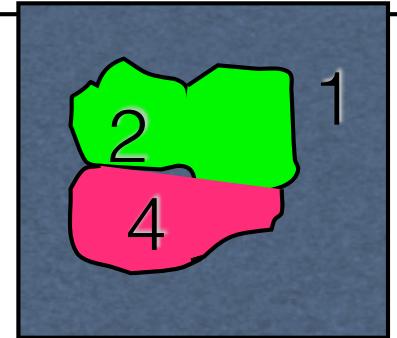


Merging Algorithm

For each region in the image do:

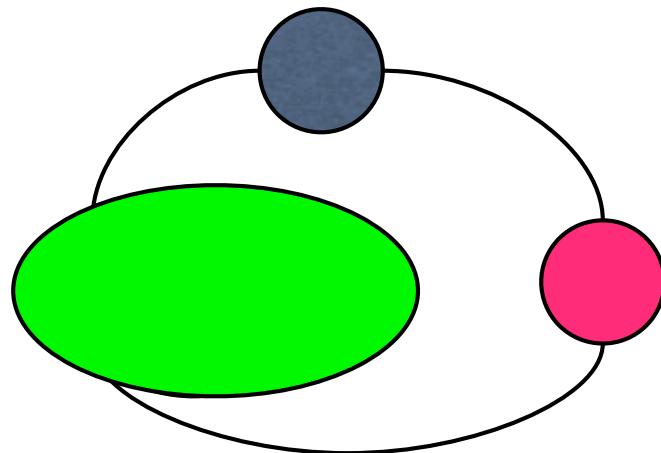
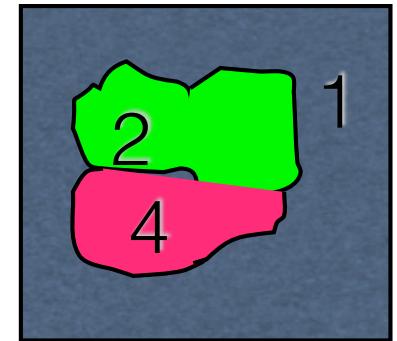
Consider its adjacent regions and test if they are similar

If they are similar, merge them and update the RAG



Merging Algorithm

Repeat the previous step until there are no more merges.



Using Similarity Measures

Compare mean intensities (zero order model)

Compare surface fitting (higher order model)

Model regions as polynomial surfaces:

$$F(x,y;a,m) = \sum_{i+j \leq m} a_{ij} x^i y^j$$

m is the polynomial degree (usually at most 2)

Compute the fitting error:

$$E^2(R,a,m) = \sum_{(x,y) \in R} [I(x,y) - F(x,y;a,m)]^2$$

Merge regions if it decreases the fitting error.

Using Statistics

Model pixels as drawn from probability distributions.

Different regions have different distributions.

Merge regions if their pixels came from the same distribution.

Example: One or Two Regions

Consider two regions:

Let g_1, g_2, \dots, g_{m_1} be the gray values in R_1

Let $g_{m_1+1}, g_{m_1+2}, \dots, g_{m_1+m_2}$ be the gray values in R_2

Assume regions are:

constant + uncorrelated zero mean Gaussian noise

Question:

Should the two regions be merged or not?

Two possible HYPOTHESES:

H_0 :

Both regions belong to the same object and the gray values have distribution $N(\mu_0, \sigma_0)$

H_1 :

Each region belongs to a different object and their gray values have distributions $N(\mu_1, \sigma_1)$ and $N(\mu_2, \sigma_2)$

Assume H_0 is TRUE:

Estimate μ_o and σ_o :

$$\mu_o = \frac{1}{m_1 + m_2} \sum_{g_i \in R_1 \cup R_2} g_i$$

$$\sigma_o^2 = \frac{1}{m_1 + m_2} \sum_{g_i \in R_1 \cup R_2} (g_i - \mu_o)^2$$

“likelihood” of H_o being true:

Compute the probability of independently drawing $g_1, g_2, \dots, g_{m1+m2}$ with distribution $N(\mu_o, \sigma_o)$:

$$P(g_1, g_2, \dots, g_{m1+m2} | H_o) = P(g_1 | H_o)P(g_2 | H_o)\dots p(g_{m1+m2} | H_o)$$

$$P(g_i | H_o) = \frac{1}{\sqrt{2\pi}\sigma_o} e^{\frac{-(g_i - \mu_o)^2}{2\sigma_o^2}}$$

$$\begin{aligned} P(g_1, g_2, \dots, g_{m1+m2} | H_o) &= \prod_{i=1}^{m1+m2} \frac{1}{\sqrt{2\pi}\sigma_o} e^{\frac{-(g_i - \mu_o)^2}{2\sigma_o^2}} \\ &= \frac{1}{(\sqrt{2\pi}\sigma_o)^{m1+m2}} e^{\frac{-1}{2\sigma_o^2} \sum_{i=1}^{m1+m2} (g_i - \mu_o)^2} \\ &= \frac{1}{(\sqrt{2\pi}\sigma_o)^{m1+m2}} e^{\frac{-(m1+m2)}{2}} \end{aligned}$$

Assume H_1 is TRUE:

Estimate μ_1 and σ_1 , μ_2 and σ_2

$$\mu_j = \frac{1}{m_j} \sum_{g_i \in R_j} g_i$$

$$\sigma_j^2 = \frac{1}{m_j} \sum_{g_i \in R_j} (g_i - \mu_j)^2$$

$$j = 1, 2$$

“likelihood” of H_1 being true:

Compute the probability of independently drawing g_1, g_2, \dots, g_{m1} with distribution $N(\mu_1, \sigma_1)$, and $g_{m1+1}, g_{m1+2}, \dots, g_{m1+m2}$ with distribution $N(\mu_2, \sigma_2)$:

$$P(g_1, g_2, \dots, g_{m1+m2} | H_1) = P(g_1 | H_1)P(g_2 | H_1)\dots p(g_{m1+m2} | H_1)$$

$$g_i \in R_1 \rightarrow P(g_i | H_1) = \frac{1}{\sqrt{2\pi}\sigma_1} e^{\frac{-(g_i - \mu_1)^2}{2\sigma_1^2}} \quad g_i \in R_2 \rightarrow P(g_i | H_1) = \frac{1}{\sqrt{2\pi}\sigma_2} e^{\frac{-(g_i - \mu_2)^2}{2\sigma_2^2}}$$

$$\begin{aligned} P(g_1, g_2, \dots, g_{m1+m2} | H_1) &= \prod_{i=1}^{m1} \frac{1}{\sqrt{2\pi}\sigma_1} e^{\frac{-(g_i - \mu_1)^2}{2\sigma_1^2}} \prod_{i=m1+1}^{m1+m2} \frac{1}{\sqrt{2\pi}\sigma_2} e^{\frac{-(g_i - \mu_2)^2}{2\sigma_2^2}} \\ &= \frac{1}{(\sqrt{2\pi})^{m1+m2} \sigma_1^{m1} \sigma_2^{m2}} e^{\frac{-1}{2\sigma_1^2} \sum_{i=1}^{m1} (g_i - \mu_1)^2} e^{\frac{-1}{2\sigma_2^2} \sum_{i=m1+1}^{m1+m2} (g_i - \mu_2)^2} \\ &= \frac{1}{(\sqrt{2\pi})^{m1+m2} \sigma_1^{m1} \sigma_2^{m2}} e^{\frac{-m1}{2}} e^{\frac{-m2}{2}} \end{aligned}$$

Choose the most “likely” Hypothesis

$$L = \frac{P(g_1, \dots, g_{m1+m2} | H_1)}{P(g_1, \dots, g_{m1+m2} | H_0)} = \frac{\sigma_o^{m1+m2}}{\sigma_1^{m1}\sigma_2^{m2}}$$

If $L < 1$, H_0 is more likely than H_1 :

Merge the regions!

Splitting Algorithms

When are regions split?

Split a region if:

- A property is not “constant”

- A predicate is not TRUE

Deciding to split is fairly straight-forward.

Where to split?

This is a difficult problem.

Some approaches:

- Divide it into equal parts along image dimensions.

- Look for strong edges to create boundaries.

Split and Merge Algorithms

Split and merge are often used together:

Start with the initial image and a “predicate”

Test the image with the predicate:

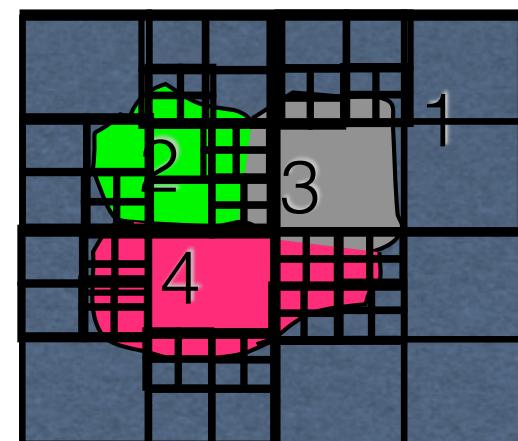
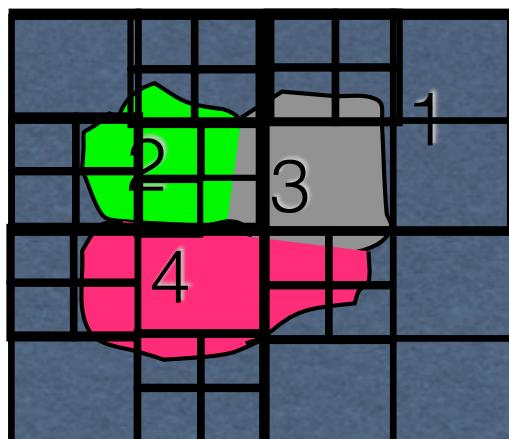
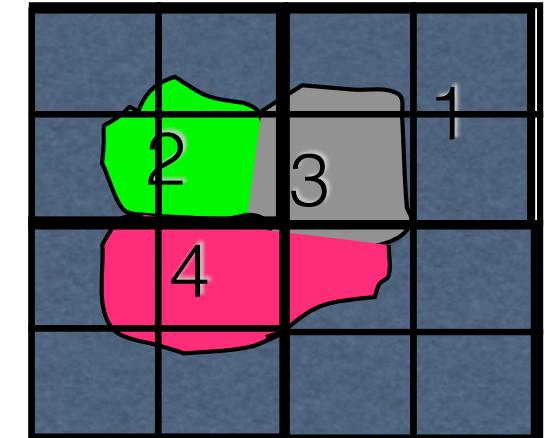
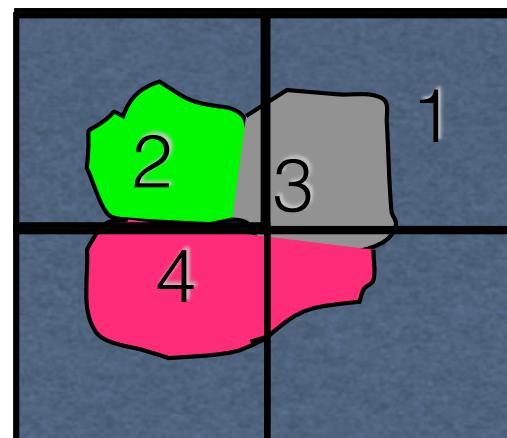
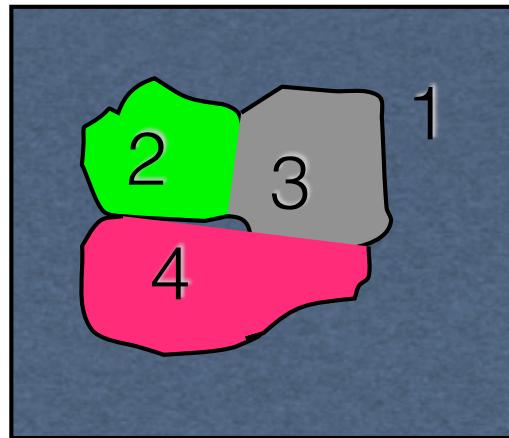
If it doesn't satisfy it, split image into quarters;

Repeat for each sub-region until there are no more splits.

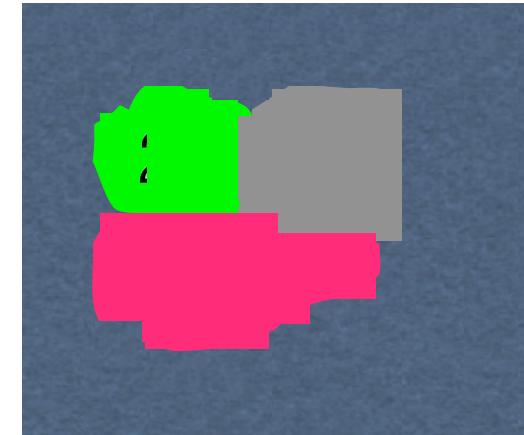
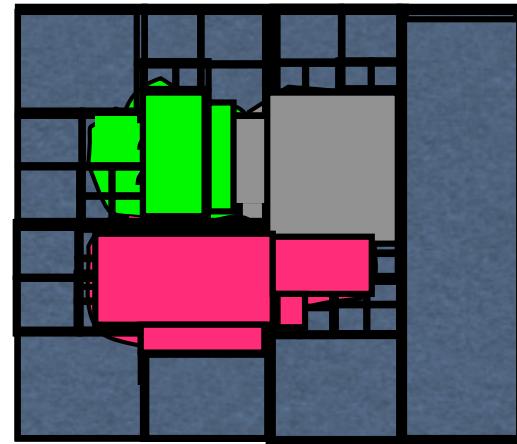
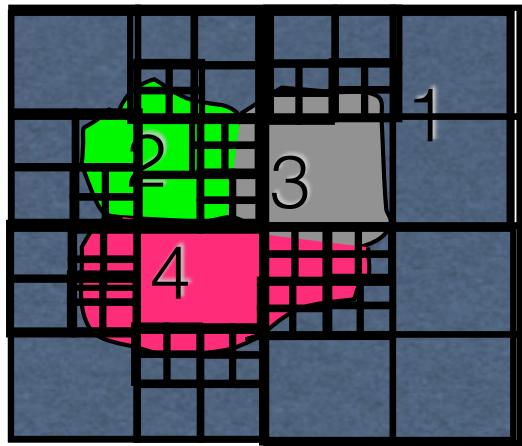
Test adjacent regions with the predicate:

If they satisfy it: merge them.

Split ...



Merge ...



Quadtree Representation

Quadtrees:

Trees where nodes have 4 children

Build quadtree:

Nodes represent regions

Every time a region is split, it's node gives birth to 4 children

Leaves are nodes for uniform regions

Merging:

Siblings that are “similar” can be merged.

Quadtree Representation

