

# EECE 5639 Computer Vision I

---

## Lecture 19

**Tracking: KLT, Camshift, Tracking as Classification, Circulant tracker**

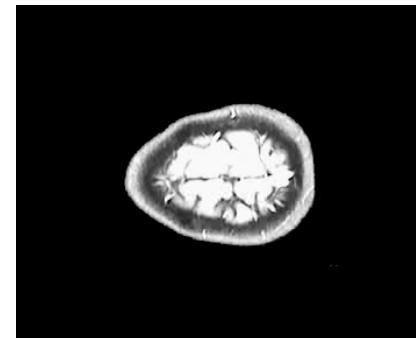
**Project 4 is out. Due April 12.**

## Next Class

**Dynamics-based Tracking**

**SFM**

# Tracking Applications



# Template Matching & LK Tracking

# Basic Template Matching

---

- Assumptions:
  - A snapshot of the target can be used to describe it
  - Target does not change (much) between frames
  - Motion is mostly translational

**The last two assumptions are very restrictive.  
Most current research works on relaxing these.**

# Template Matching

---

Is a **search problem**:

Given a “template” from a previous frame, search for the corresponding region in the current frame

Typically, use constraints to reduce the search space

Need a “matching function” to evaluate similarity

# Some Practical Issues

---

Object shape might not be well described by a patch aligned with the image axes.



We end including lots of the background in the template.

# Some Practical Issues

---

A solution: use a Gaussian windowing function to weight pixels on the object more than pixels in the background



Works best with “compact” targets such as vehicles.

# Some Practical Issues

---

Do not want to search the whole image for the target.

**Solution:** search a bounded region around likely position

Need: a prediction of where to look (more of this later).

Target appearance might change.

**Solution:** use an adaptive template

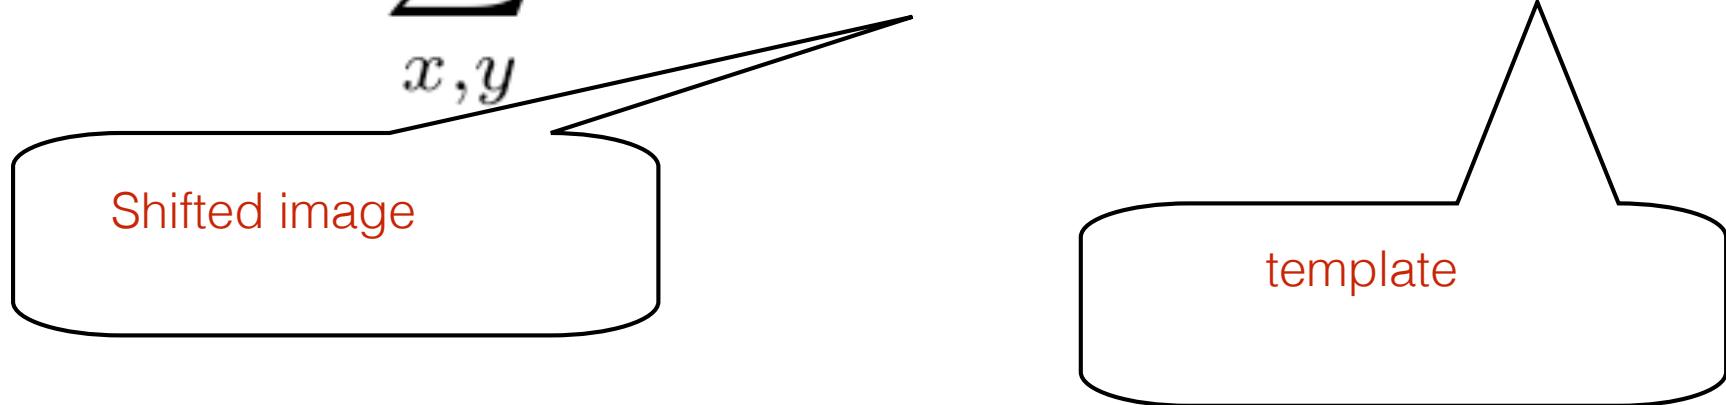
Need: a rule to change it (more of this later).

# Lucas Kanade Tracking

---

Consider the SSD error function:

$$E(u, v) = \sum_{x, y} [I(x + u, y + v) - T(x, y)]^2$$

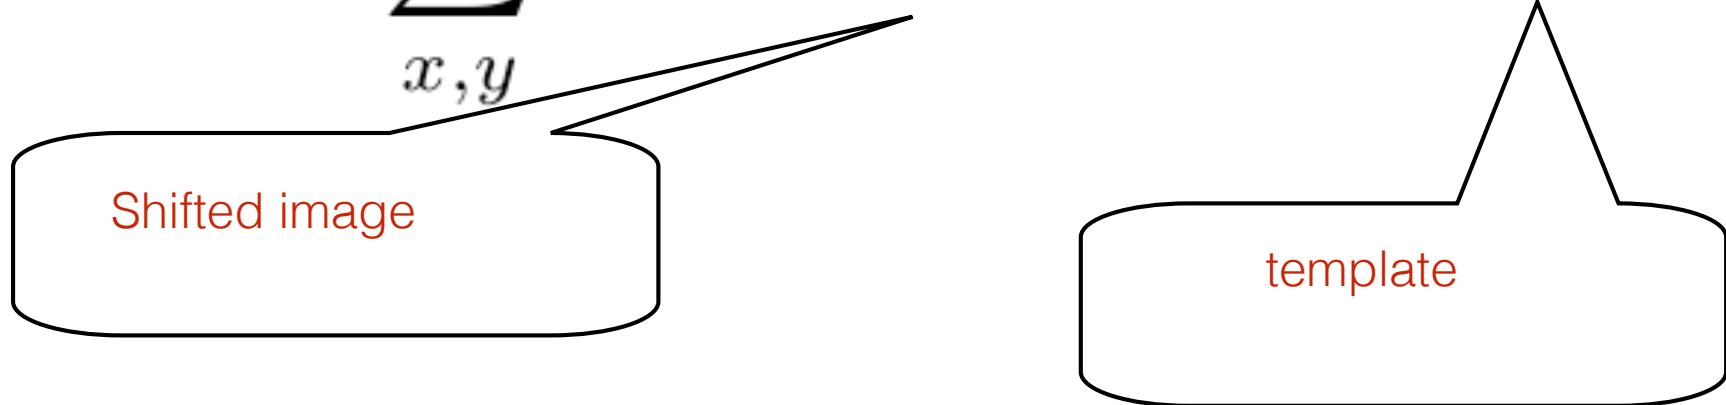


Score when template is at location  $(u, v)$

# Lucas Kanade Tracking

Consider the SSD error function:

$$E(u, v) = \sum_{x, y} [I(x + u, y + v) - T(x, y)]^2$$



Score when template is at location  $(u, v)$

We want to find the location  $(u, v)$  where  $E$  is minimum.

**Approach:** steepest descent.

# Lucas Kanade Tracking

---

$$E(u, v) = \sum_{x,y} [I(x + u, y + v) - T(x, y)]^2$$

Want to compute the gradient of this wrt u and v and set it to 0.

Steps:

- 1) Use a 1st order approximation of E
- 2) Take partial derivatives wrt u and v
- 3) Set them to 0
- 4) Solve a linear system of equations

# First Order Approximation

---

$$\begin{aligned} E(u, v) &= \sum_{x,y} [I(x + u, y + v) - T(x, y)]^2 \\ &\approx \sum_{x,y} [I(x, y) + uI_x(x, y) + vI_y(x, y) - T(x, y)]^2 \\ &\approx \sum_{x,y} [uI_x(x, y) + vI_y(x, y) + D(x, y)]^2 \end{aligned}$$

# Taking Partial Derivatives

---

$$E(u, v) \approx \sum_{x,y} [u I_x(x, y) + v I_y(x, y) + D(x, y)]^2$$

$$\frac{\partial E}{\partial u} = 2 \sum_{x,y} [u I_x(x, y) + v I_y(x, y) + D(x, y)] I_x(x, y)$$

$$\frac{\partial E}{\partial v} = 2 \sum_{x,y} [u I_x(x, y) + v I_y(x, y) + D(x, y)] I_y(x, y)$$

Set them to 0 and solve for u and v

---

$$\sum_{x,y} [uI_x(x,y) + vI_y(x,y) + D(x,y)] I_x(x,y) = 0$$

$$\sum_{x,y} [uI_x(x,y) + vI_y(x,y) + D(x,y)] I_y(x,y) = 0$$

$$\sum_{x,y} \begin{bmatrix} I_x^2 & I_x I_y \\ I_x I_y & I_y^2 \end{bmatrix} \begin{bmatrix} u \\ v \end{bmatrix} = - \sum_{x,y} \begin{bmatrix} I_x D \\ I_y D \end{bmatrix}$$

Again, this is the matrix used for corner detection!

# Lucas Kanade Tracking

---

Traditional LK is typically run on small (5x5) corner like features.

Observation:

Can we use the same approach on larger windows, say a b.box around the target?



# LK Tracking

---

In our previous derivation we made implicit assumptions ...

- 1) brightness constancy assumption
- 2) all points in the template moved the same

# For larger windows that might be a stretch

---



# Fix by changing the motion parameters:

---

Some simple models: affine and projective.

# Remember ...

---

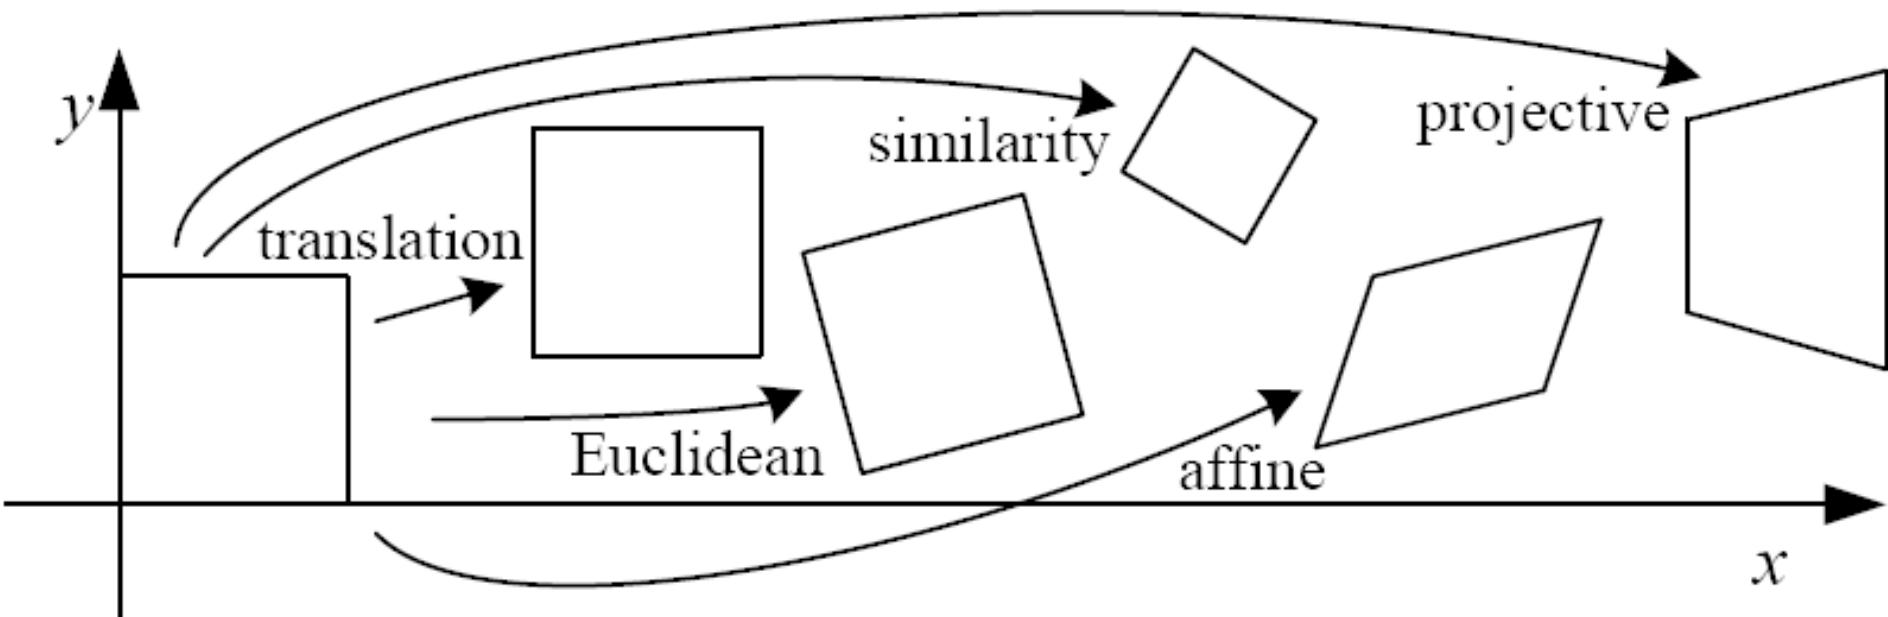
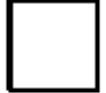
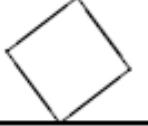
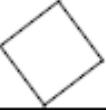


FIGURE 1. Basic set of 2D planar transformations

# Summary of 2D Transformations

Name	Matrix	# D.O.F.	Preserves:	Icon
translation	$[ \ I \   \ t \ ]_{2 \times 3}$	2	orientation + ⋯	
rigid (Euclidean)	$[ \ R \   \ t \ ]_{2 \times 3}$	3	lengths + ⋯	
similarity	$[ \ sR \   \ t \ ]_{2 \times 3}$	4	angles + ⋯	
affine	$[ \ A \ ]_{2 \times 3}$	6	parallelism + ⋯	
projective	$[ \ H \ ]_{3 \times 3}$	8	straight lines	

# Back to LK Tracking ...

---

Consider a warping function  $W$  with parameters  $P$ :  $W([x,y];P)$

Examples:

- 1) Translation  $(u,v)$  (OF)  $P=[u,v]$
- 2) Affine transformation  $P$  has 6 parameters  $(A,b)$

Then, the objective is to minimize the matching score wrt  $P$

$$E(u, v) = \sum_{x,y} [I((W(x, y); P)) - T(x, y)]^2$$

# Chain Rule

---

$$z = f(x, y) \quad \text{Image gray values}$$

$$\begin{aligned} x &= x(t) \\ y &= y(t) \end{aligned} \quad \text{Warping of image coordinates}$$

$$\frac{\partial f}{\partial t} = \frac{\partial f}{\partial x} \frac{\partial x}{\partial t} + \frac{\partial f}{\partial y} \frac{\partial y}{\partial t}$$

# Taylor Expansion

$$I = I(x, y; p_1, p_2, \dots, p_n)$$

After warping with slightly changed parameters:

$$I' = I(x, y; p_1 + \Delta p_1, p_2 + \Delta p_2, \dots, p_n + \Delta p_n)$$

Using Taylor's expansion:

$$I' \approx I(x, y; p_1, p_2, \dots, p_n) + \left[ \begin{array}{c|c} \frac{\partial I}{\partial x} & \frac{\partial x}{\partial p_1} \\ \hline \end{array} \right] \Delta p_1 + \left[ \begin{array}{c|c} \frac{\partial I}{\partial y} & \frac{\partial y}{\partial p_1} \\ \hline \end{array} \right] \Delta p_1$$
$$+ \left[ \begin{array}{c|c} \frac{\partial I}{\partial x} & \frac{\partial x}{\partial p_2} \\ \hline \end{array} \right] \Delta p_2 + \left[ \begin{array}{c|c} \frac{\partial I}{\partial y} & \frac{\partial y}{\partial p_2} \\ \hline \end{array} \right] \Delta p_2$$
$$\dots$$
$$+ \left[ \begin{array}{c|c} \frac{\partial I}{\partial x} & \frac{\partial x}{\partial p_n} \\ \hline \end{array} \right] \Delta p_n + \left[ \begin{array}{c|c} \frac{\partial I}{\partial y} & \frac{\partial y}{\partial p_n} \\ \hline \end{array} \right] \Delta p_n$$

Image Gradient

Warp Jacobian

The diagram illustrates the Taylor expansion of a warped image intensity. The expansion is shown as a sum of terms, each consisting of a partial derivative of the image with respect to position (x or y) multiplied by a corresponding element of the warp Jacobian. The warp Jacobian is represented as a vertical stack of columns, where each column is labeled with a partial derivative of the image with respect to a parameter ( $p_1, p_2, \dots, p_n$ ). The diagram uses green boxes to highlight the partial derivatives and the Jacobian columns. Red arrows point from the text 'Image Gradient' and 'Warp Jacobian' to their respective components in the expansion.

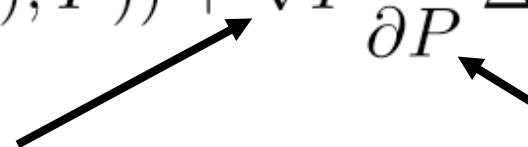
# LK Tracking

---

Use steepest descent to find (local) min.

$$E(P) = \sum_{x,y} [I((W(x,y); P)) - T(x,y)]^2$$

$$E(P + \Delta P) \approx \sum_{x,y} \left[ I((W(x,y); P)) + \nabla I \frac{\partial W}{\partial P} \Delta P - T(x,y) \right]^2$$

  
Image gradient      Warp Jacobian

# Example: Affine warping (6 parameters)

---

$$W = \begin{bmatrix} 1 + p_1 & p_3 & p_5 \\ p_2 & 1 + p_4 & p_6 \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$$

# Warp Jacobian

---

$$W([x, y]; P) = (W_x, W_y)$$

$$\frac{\partial W}{\partial P} = \begin{bmatrix} \frac{\partial W_x}{\partial p_1} & \frac{\partial W_x}{\partial p_2} & \cdots & \frac{\partial W_x}{\partial p_n} \\ \frac{\partial W_y}{\partial p_1} & \frac{\partial W_y}{\partial p_2} & \cdots & \frac{\partial W_y}{\partial p_n} \end{bmatrix}$$

# Warp Jacobian, Affine warping

---

$$W = \begin{bmatrix} 1 + p_1 & p_3 & p_5 \\ p_2 & 1 + p_4 & p_6 \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$$

$$W_x = (1 + p_1)x + p_3y + p_5$$

$$W_y = p_2x + (1 + p_4)y + p_6$$

# Warp Jacobian, Affine warping

---

$$W_x = (1 + p_1)x + p_3y + p_5$$

$$W_y = p_2x + (1 + p_4)y + p_6$$

$$\frac{\partial W}{\partial P} = \begin{bmatrix} \frac{\partial W_x}{\partial p_1} & \frac{\partial W_x}{\partial p_2} & \cdots & \frac{\partial W_x}{\partial p_n} \\ \frac{\partial W_y}{\partial p_1} & \frac{\partial W_y}{\partial p_2} & \cdots & \frac{\partial W_y}{\partial p_n} \end{bmatrix}$$

$$\frac{\partial W}{\partial P} = \begin{bmatrix} x & 0 & y & 0 & 1 & 0 \\ 0 & x & 0 & y & 0 & 1 \end{bmatrix}$$

Taking derivatives of E wrt P and setting them to 0:

---

$$E(P + \Delta P) \approx \sum_{x,y} \left[ I((W(x, y); P)) + \nabla I \frac{\partial W}{\partial P} \Delta P - T(x, y) \right]^2$$

$$\sum_{x,y} \left\{ [\nabla I \frac{\partial W}{\partial P}]^T [(I(W[x, y]; P)) + \nabla I \frac{\partial W}{\partial P} \Delta P - T(x, y)] \right\} = 0$$

- Solve for Delta P;
- Update P:  $P = P + \Delta P$
- Iterate until Delta P is negligible

# Solve for Delta P:

---

$$\sum_{x,y} \left\{ [\nabla I \frac{\partial W}{\partial P}]^T [(I(W[x,y]; P)) + \nabla I \frac{\partial W}{\partial P} \Delta P - T(x,y)] \right\} = 0$$

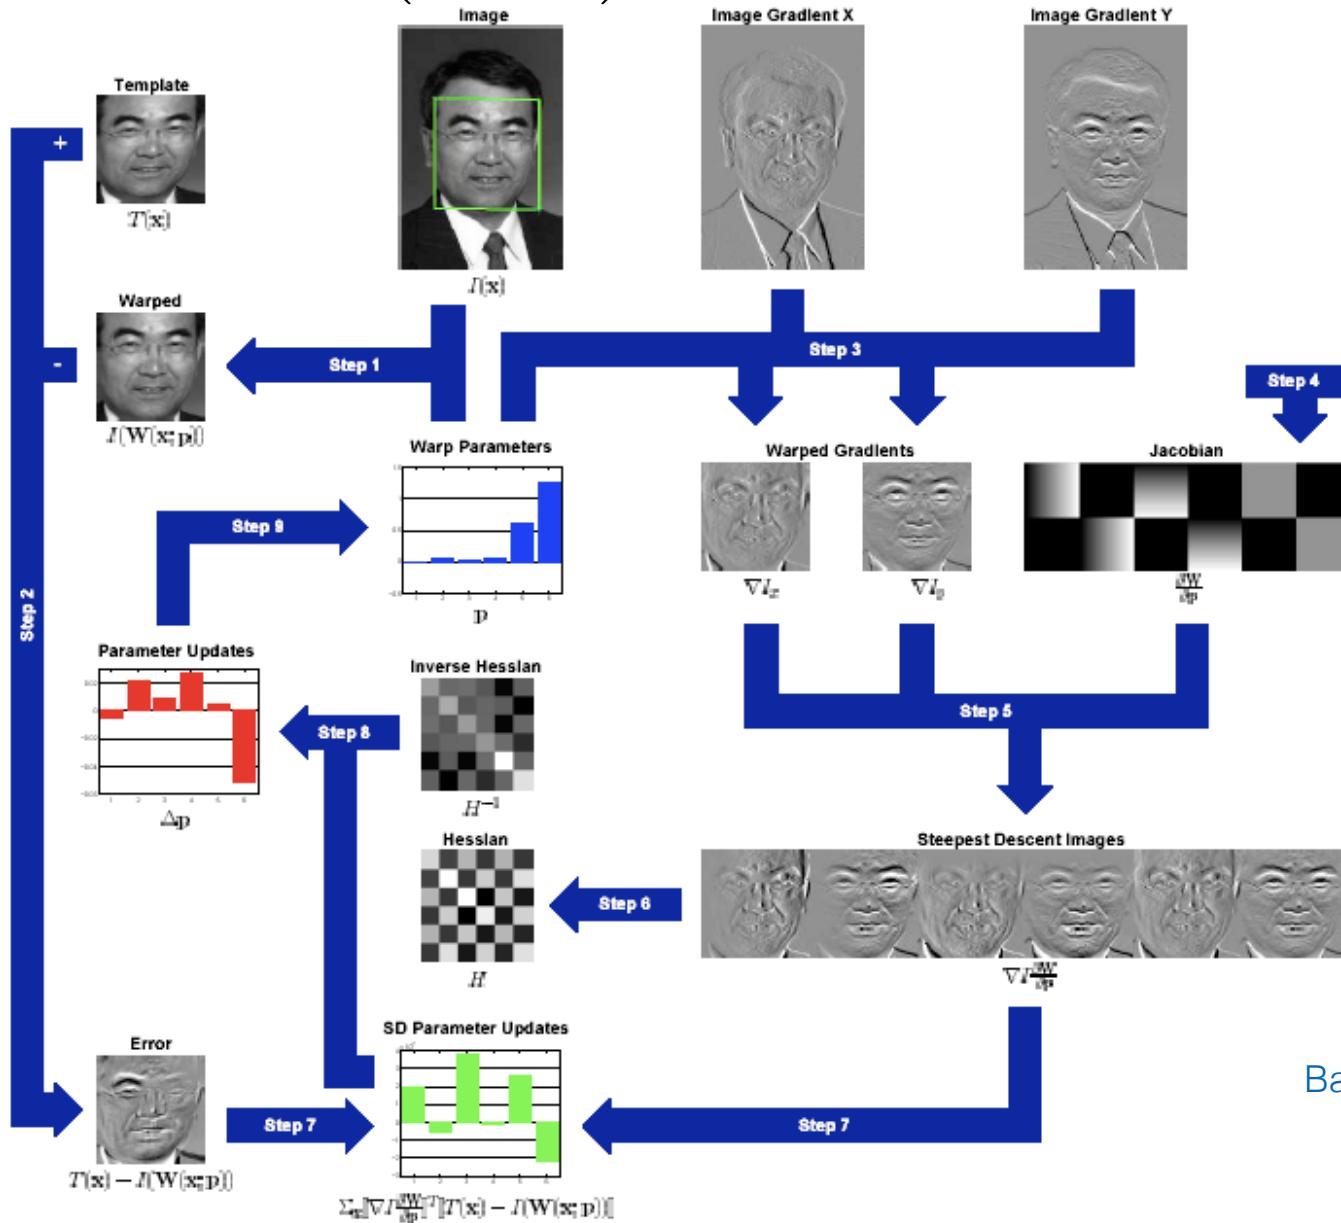
Gauss-Newton Approximation of the Hessian



$$\begin{aligned}\Delta P &= [\sum (\nabla I \frac{\partial W}{\partial P})^T \nabla I \frac{\partial W}{\partial P}]^{-1} \sum (\nabla I \frac{\partial W}{\partial P})^T [T(x,y) - (I(W[x,y]; P))] \\ &= H^{-1} \sum (\nabla I \frac{\partial W}{\partial P})^T [T(x,y) - (I(W[x,y]; P))]\end{aligned}$$

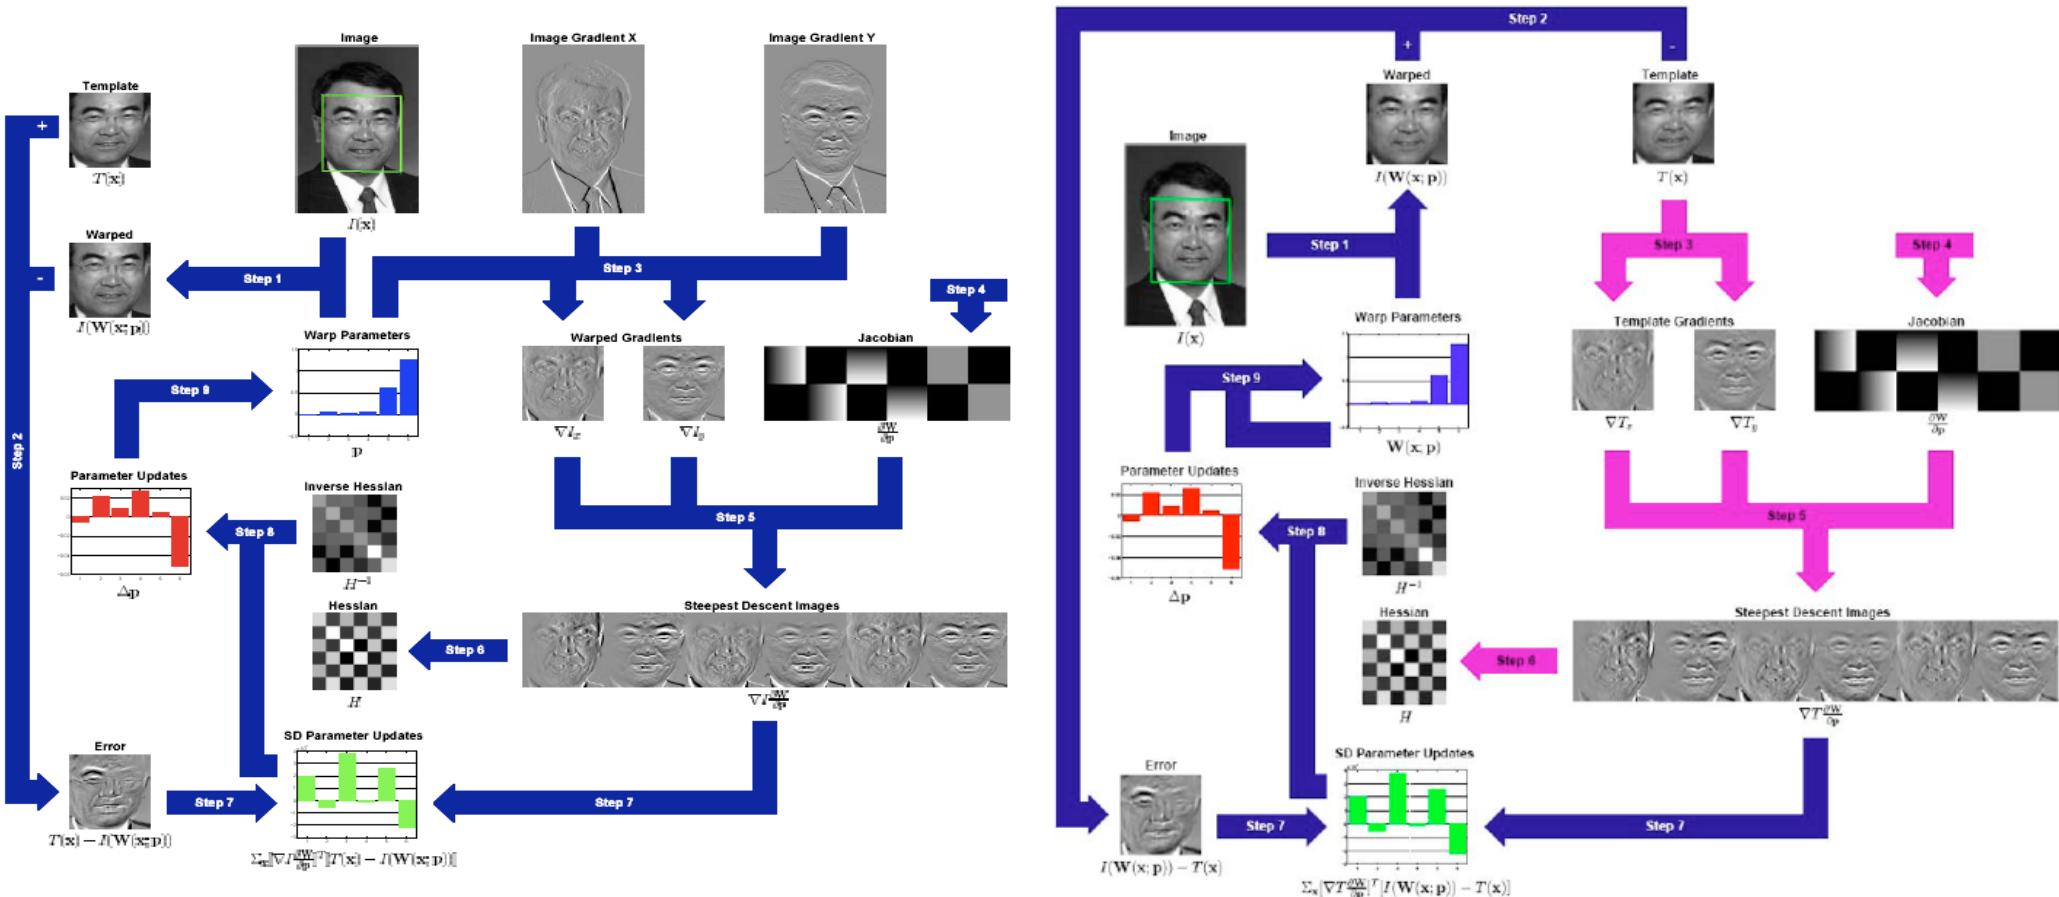
# Algorithm at a glance

$$\Delta P = H^{-1} \sum \left( \nabla I \frac{\partial W}{\partial P} \right)^T [T(x, y) - I(W[x, y]; P)]$$



Baker & Matthews  
IJCV 04

# Inverse Compositional Variant



# Brightness Changes

---

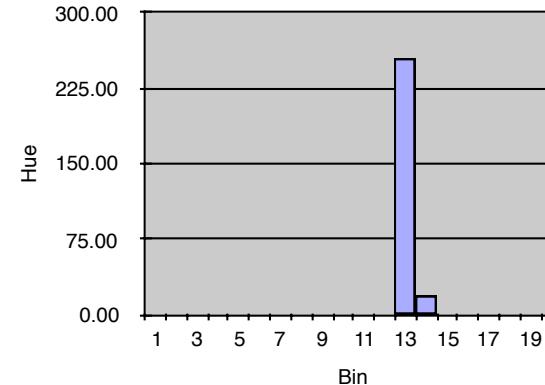
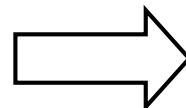
Traditional LK assumes brightness constancy.

We can remove potential error sources by first normalizing the template and the local search window, by taking off the mean and dividing by the std. dev.

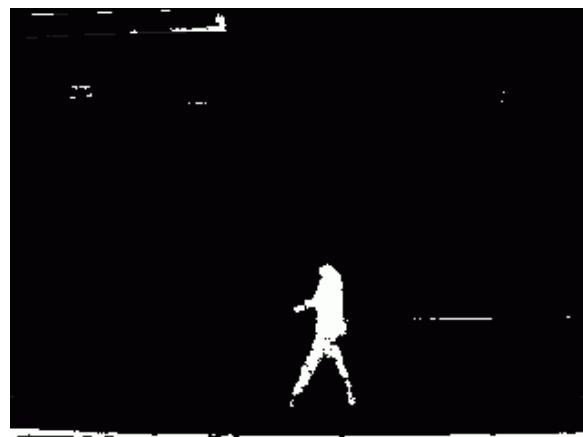
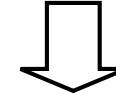
# Color-based tracking



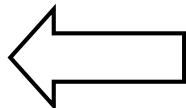
1. Set initial search window



2. Calculate histogram



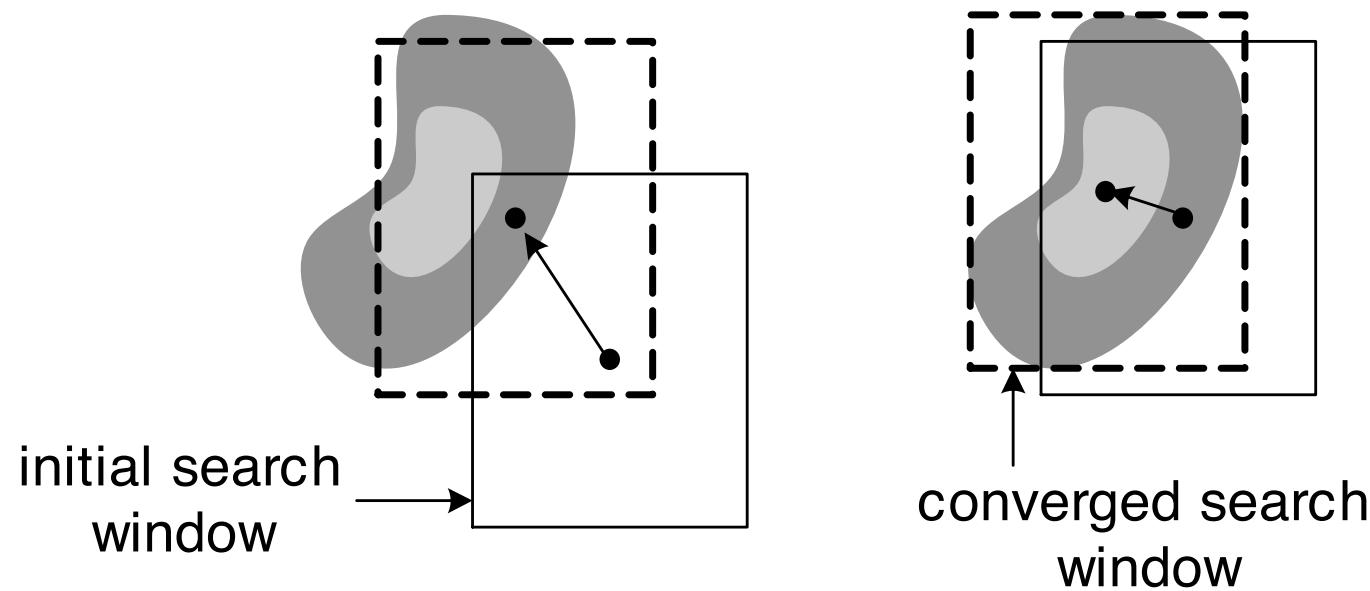
4. Probability image



3. Original image

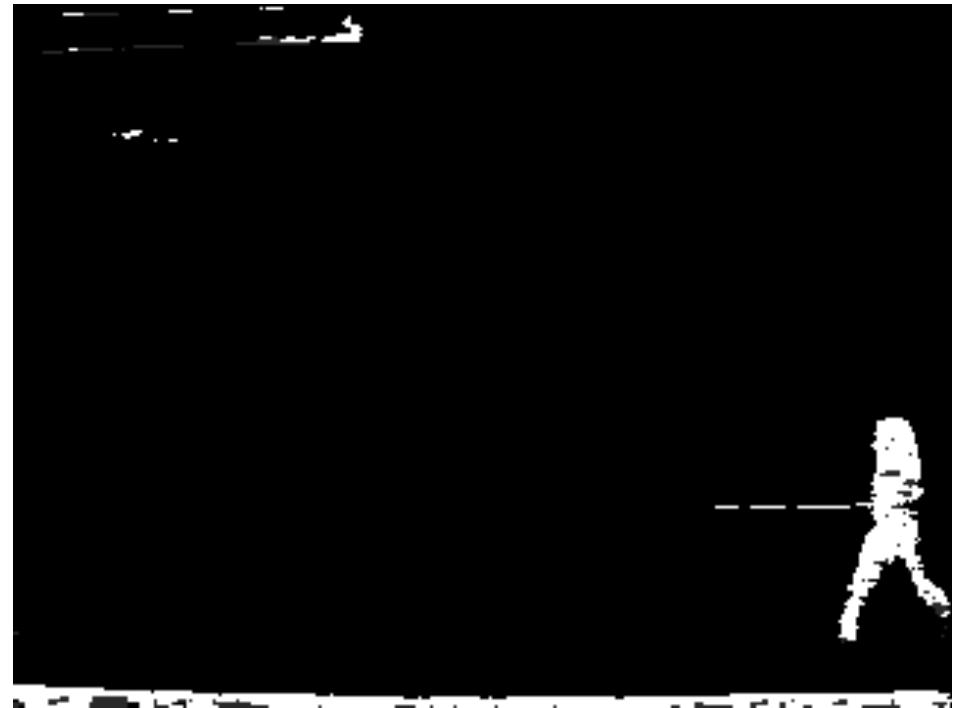
# Finding the most likely blob:

Mean Shift Algorithm: steepest ascent search of the mode:



# Camshift Tracking

---



# Tracking Colored Blobs:Algorithm CAMSHIFT

(Continuously Adaptive Mean Shift)

---

## Application:

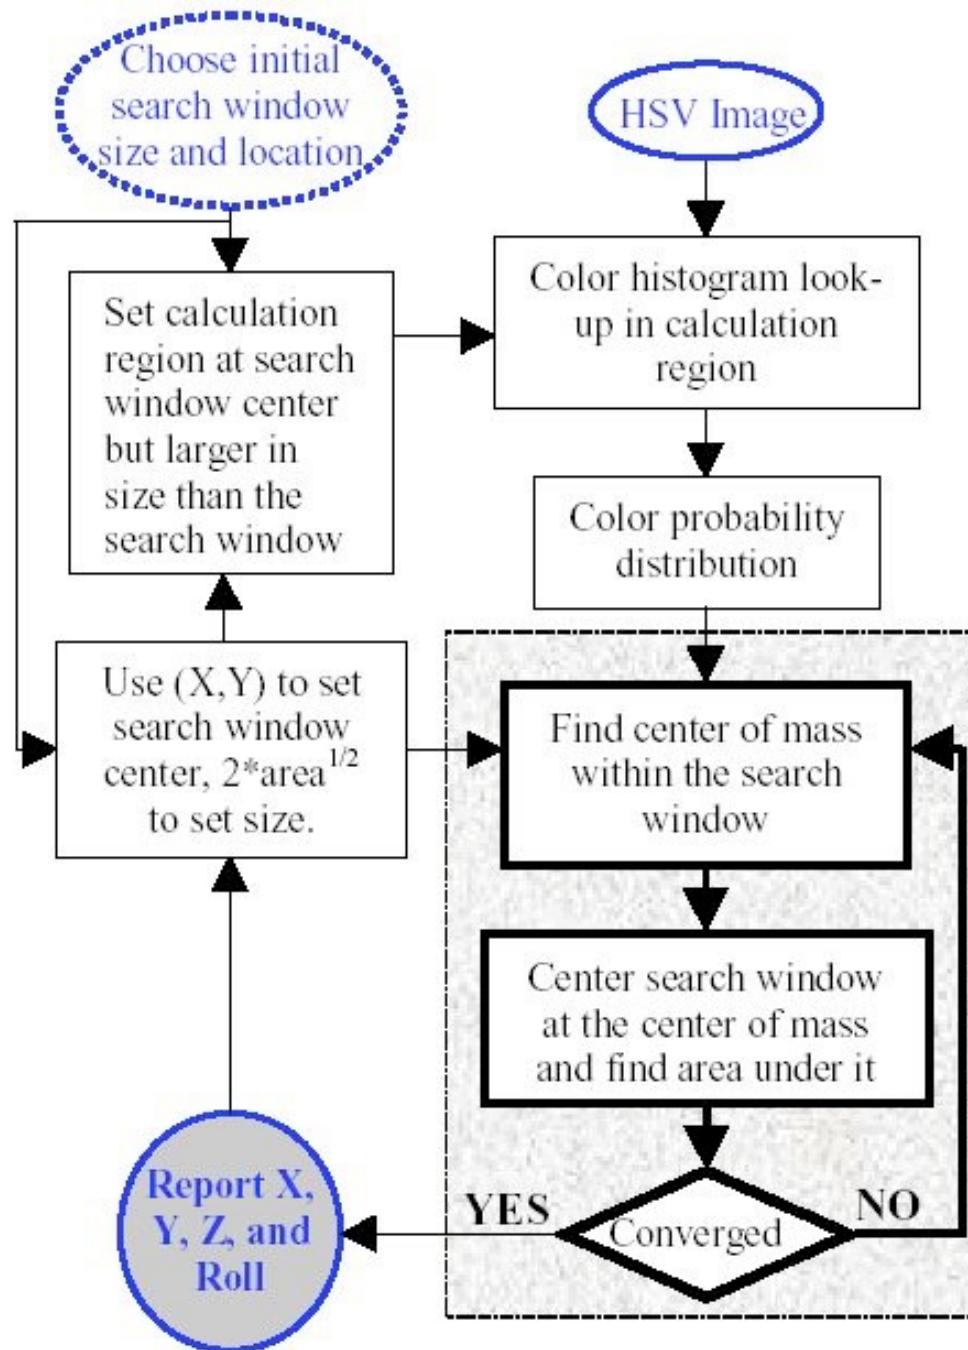
Tracking faces for hci, teleconferencing, etc

## Approach:

Find pixels that have the same color distribution across frames by looking at their histogram in HSV color space

Based on “Mean Shift” algorithm

Adapts the size of the window at each frame

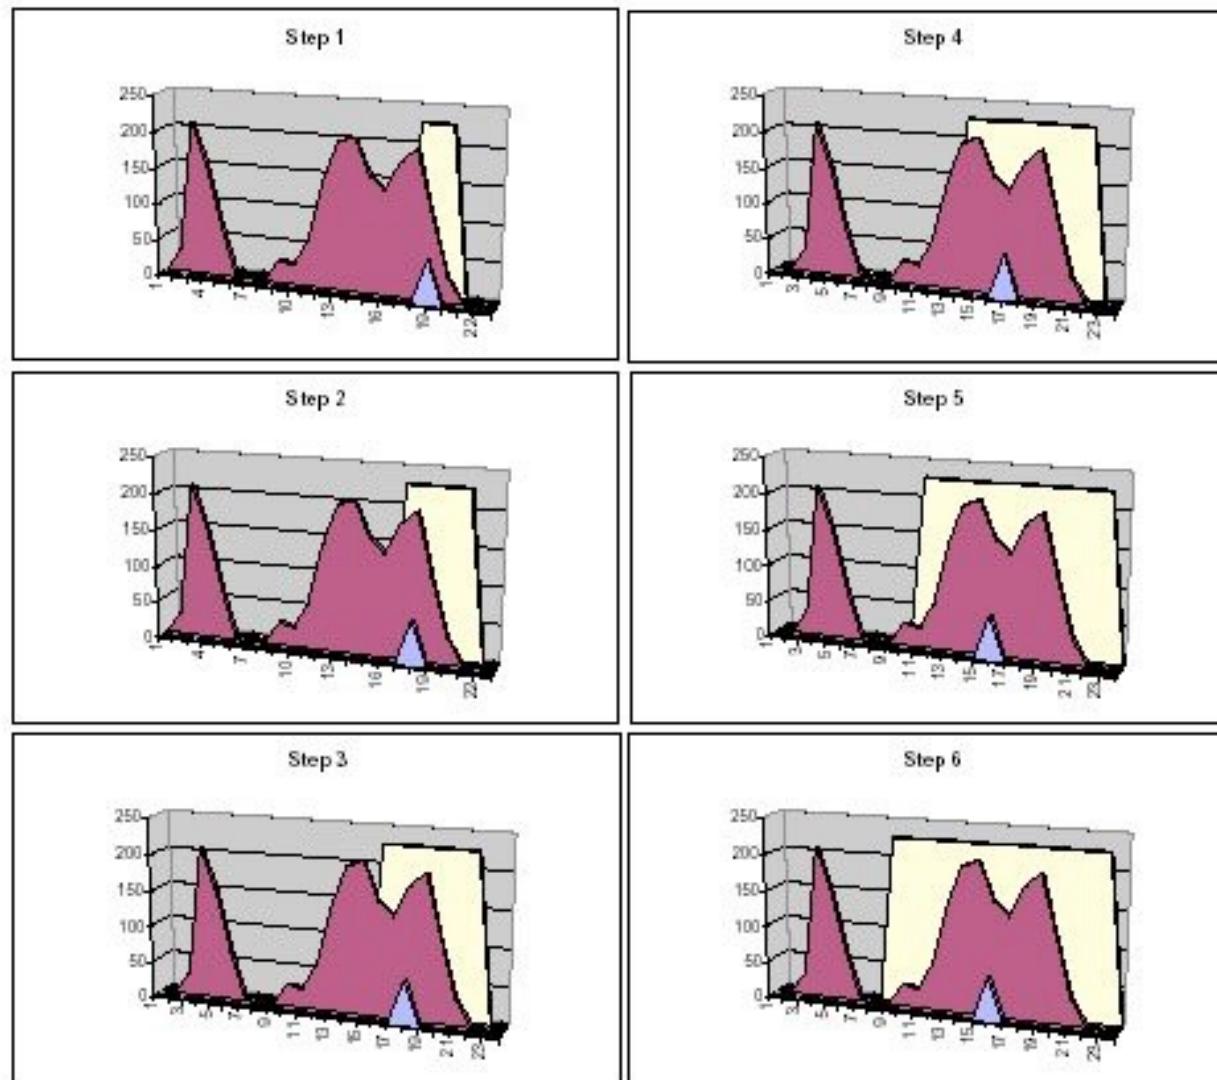


# Algorithm CAMSHIFT

---

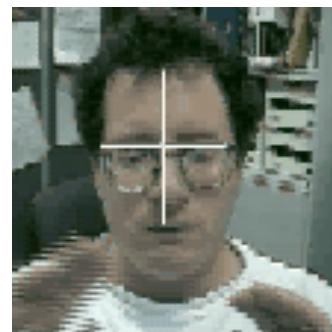
1. Choose initial window
2. Mean Shift (one or more iterations)
3. Record zeroth moment of the window  $M_{00} = \sum I(x,y)$
4. Set window size  $s = 2 * (M_{00}/256)^{1/2}$
5. Repeat steps 2 to 4 until convergence.

# Camshaft Windows Iterations



# EXAMPLES

---



Tracking as  
Classification

# Online Adaptation

---

Goal: long-term tracking

The more invariant is the appearance model to variations in lighting and geometry, the less specific it is in representing a particular object.

Then, there is danger of getting confused by surroundings

Online adaptation allows us to retain specificity at each time frame while maintaining some generality about appearance changes .

# Tracking as Classification (Collins & Liu '03)

---

“Online Selection of Discriminative Tracking Features” (iccv 03)

Target tracking posed as a binary classification problem:  
discriminate between foreground (target) and background

This point of view brings in a wide range of classification techniques to  
the tracking problem.

# Motivation

---

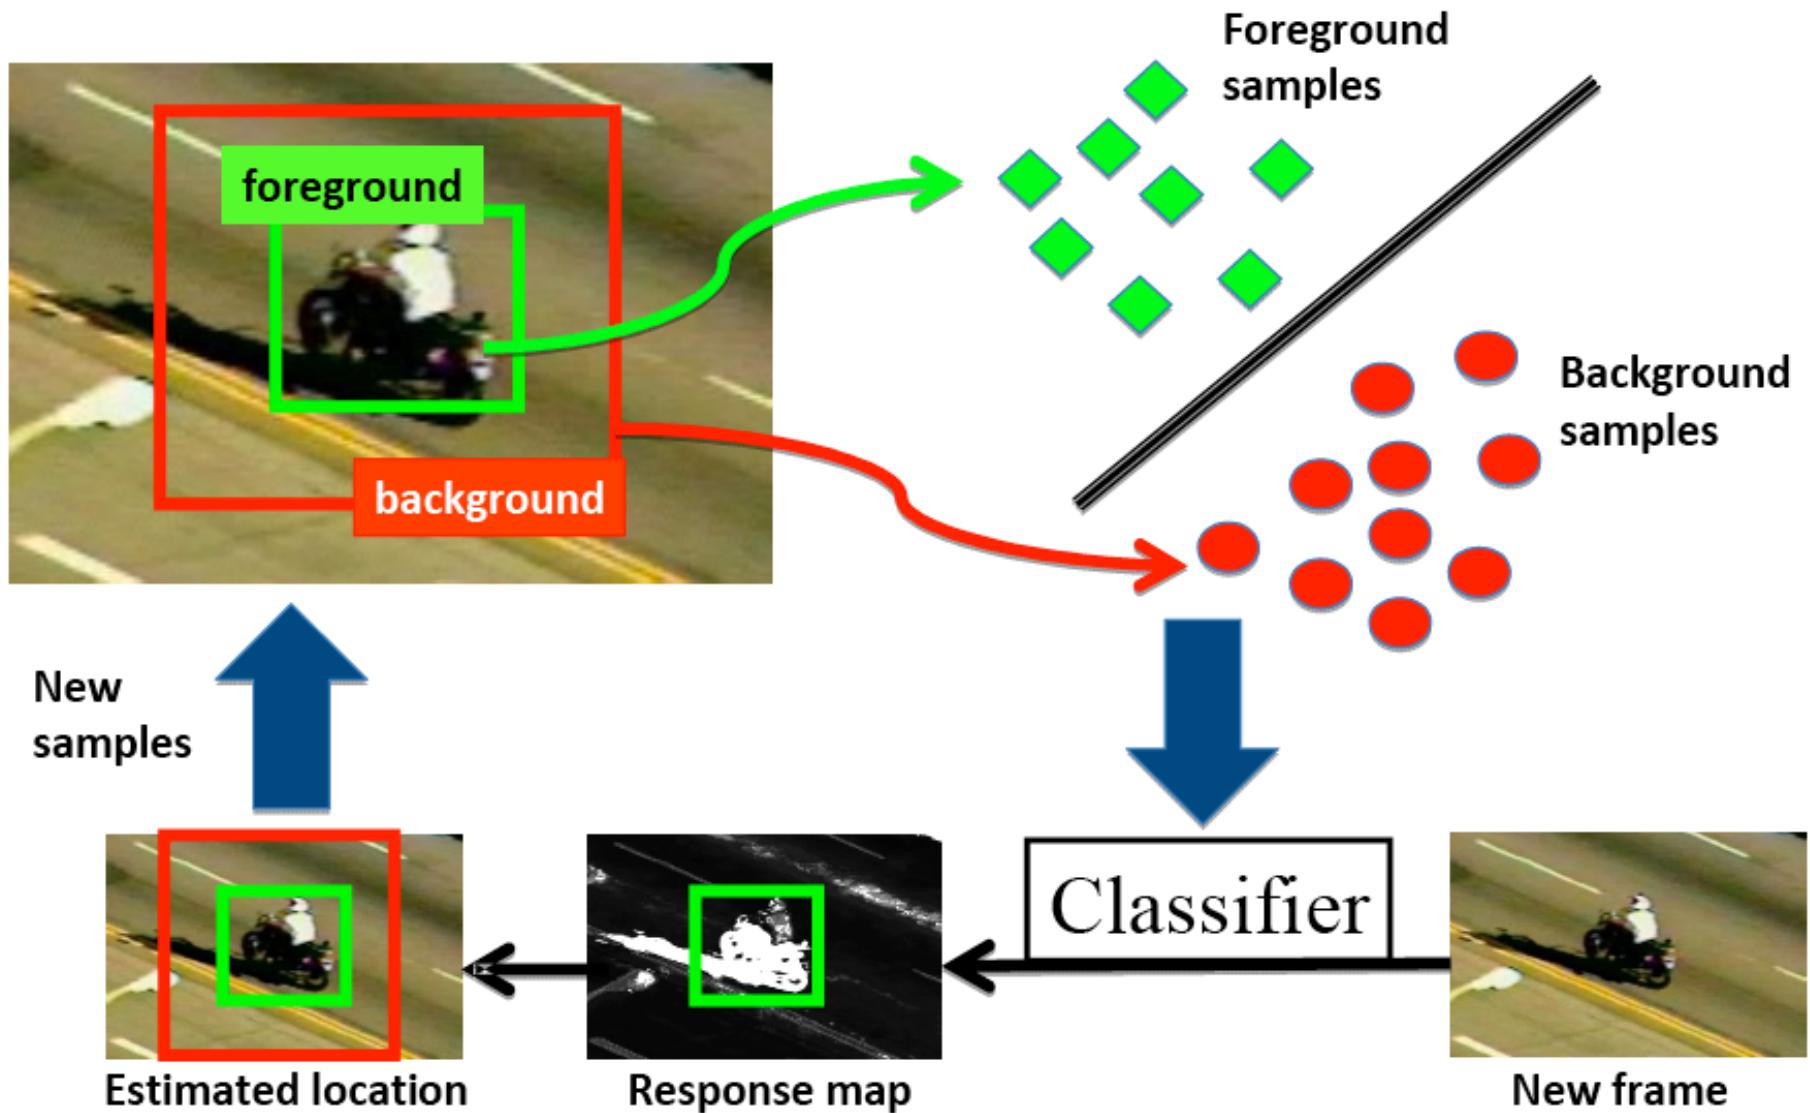
Tracking success/failure is highly correlated with our ability to distinguish the target appearance from the background (think camouflage)



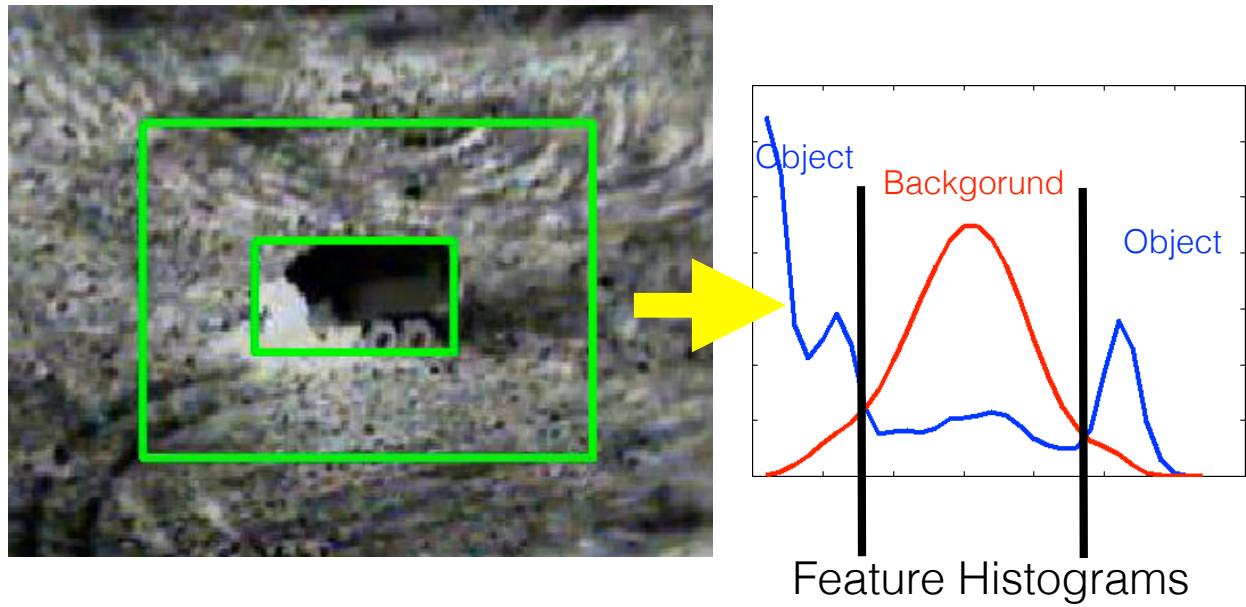
Suggestions:

Explicitly seek “good” features that discriminate  
Continuously adapt the features

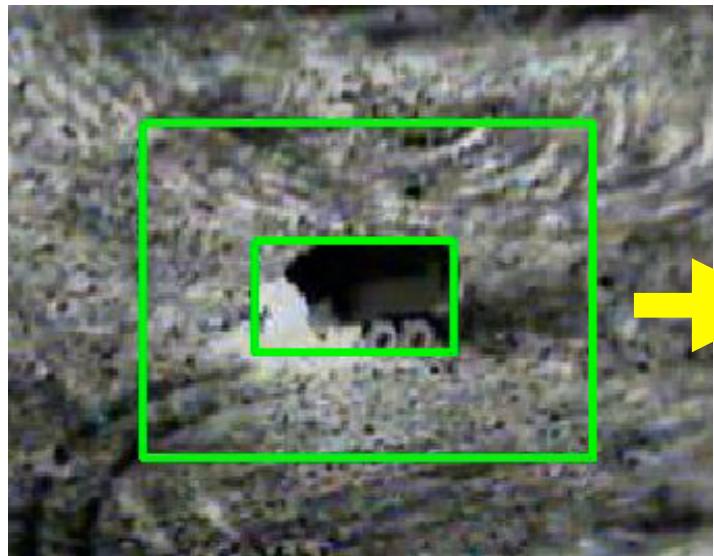
# Overview



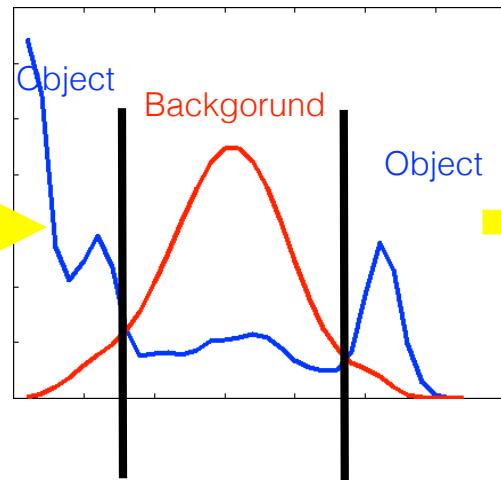
# Evaluation of Feature Discriminant Power



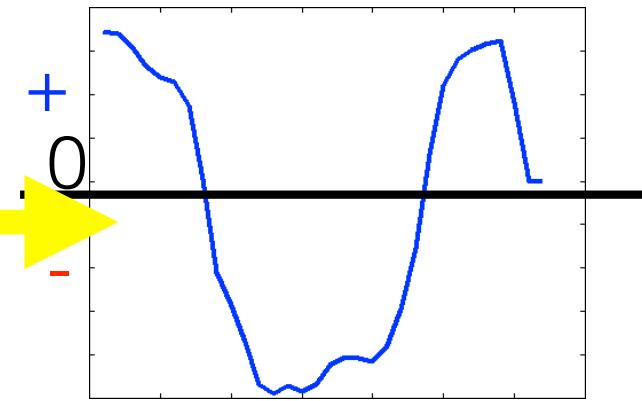
# Evaluation of Feature Discriminant Power



Feature Histograms

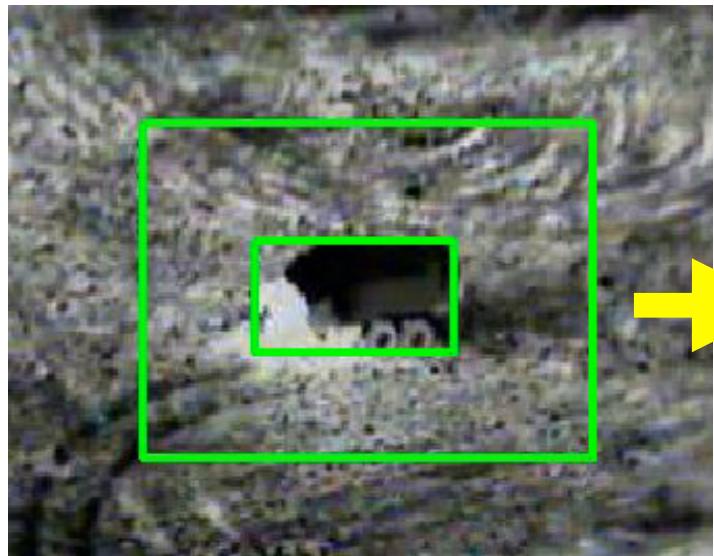


Log Likelihood Ratio

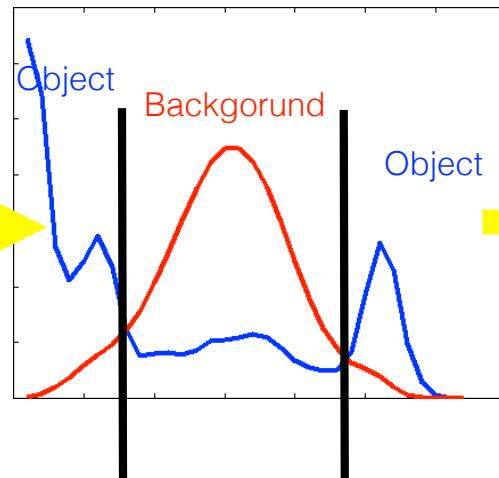


$$L(i) = \log \frac{\max \{ p(i), \delta \}}{\max \{ q(i), \delta \}}$$

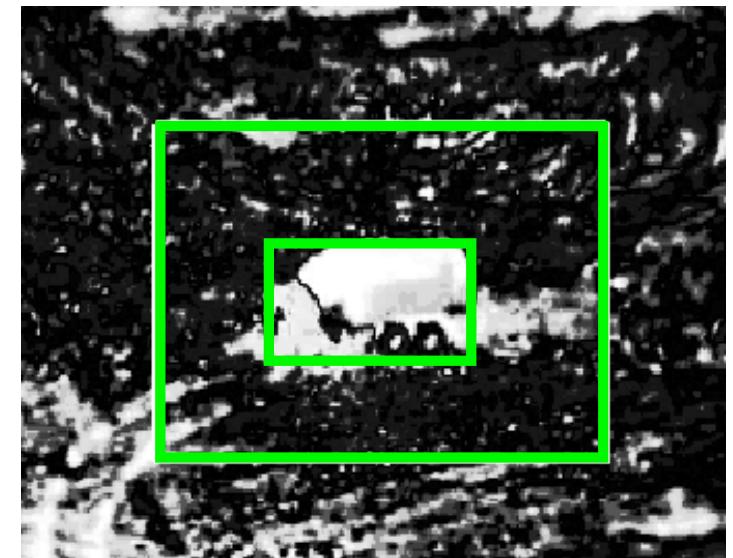
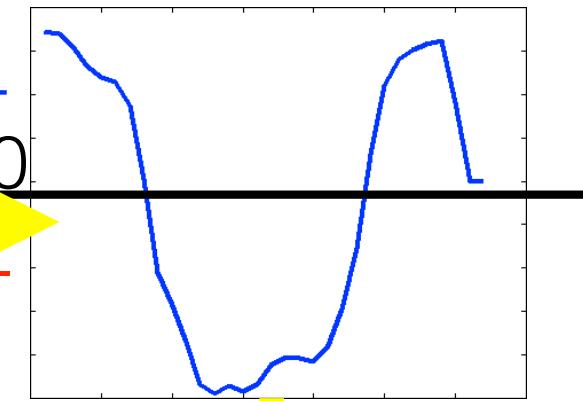
# Evaluation of Feature Discriminant Power



Feature Histograms

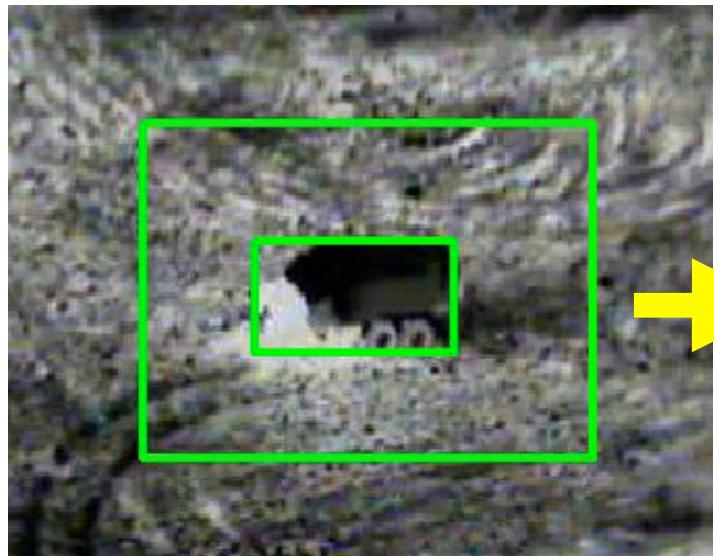


Log Likelihood Ratio

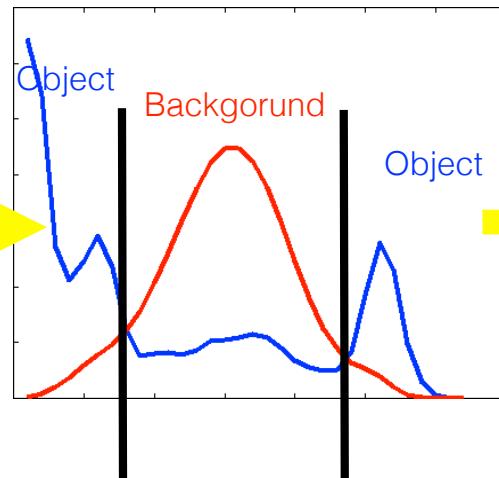


Likelihood Image

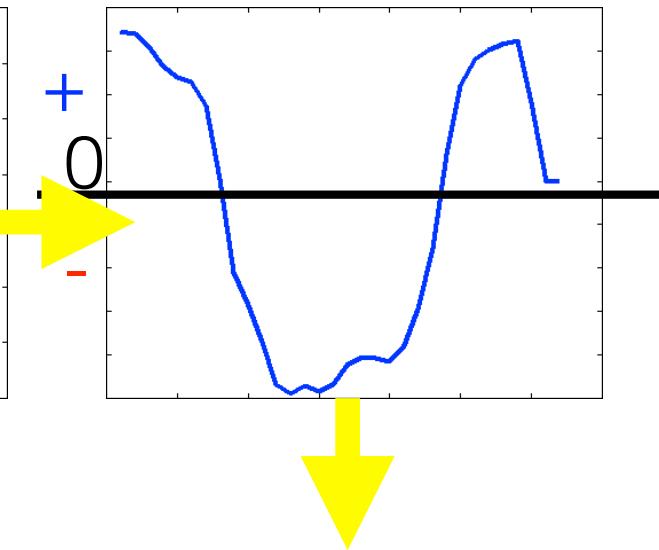
# Evaluation of Feature Discriminant Power



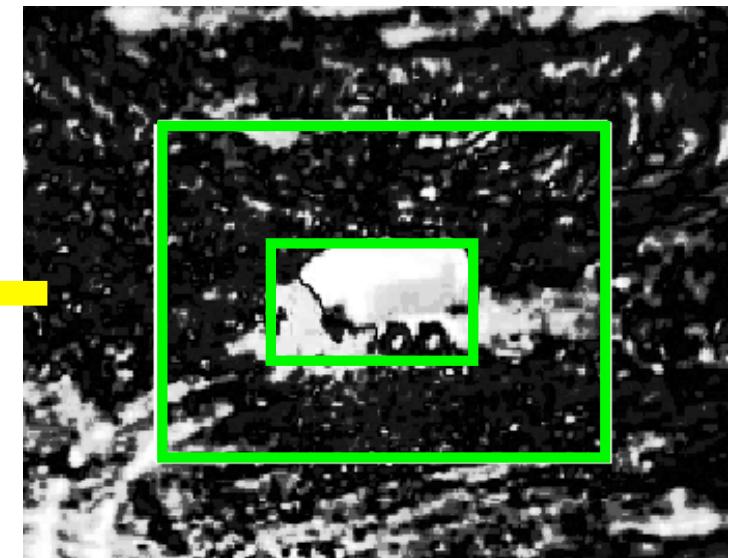
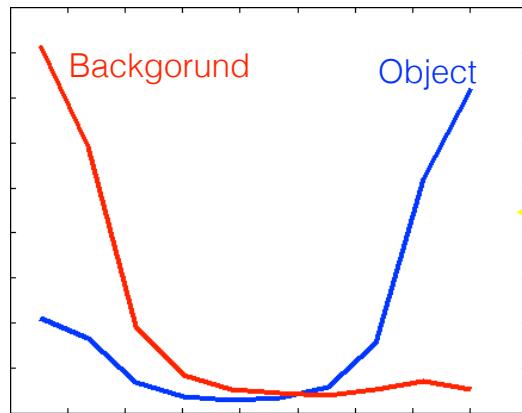
Feature Histograms



Log Likelihood Ratio

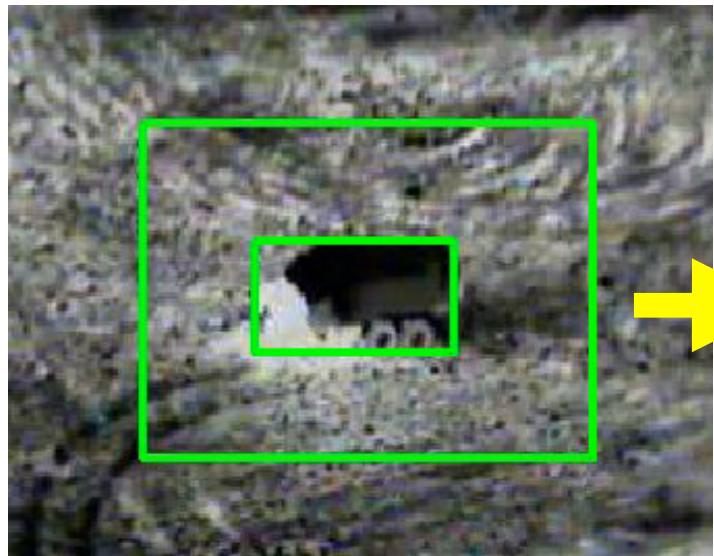


Likelihood Histograms

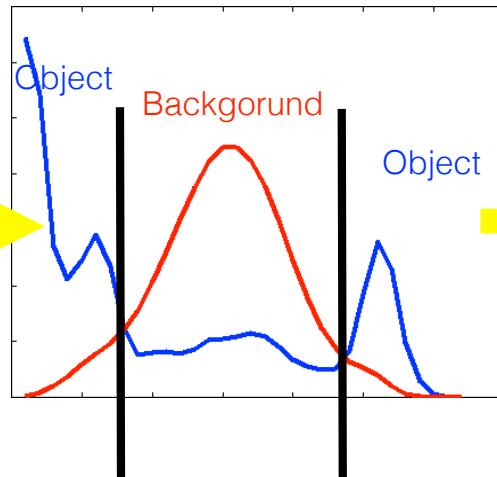


Likelihood Image

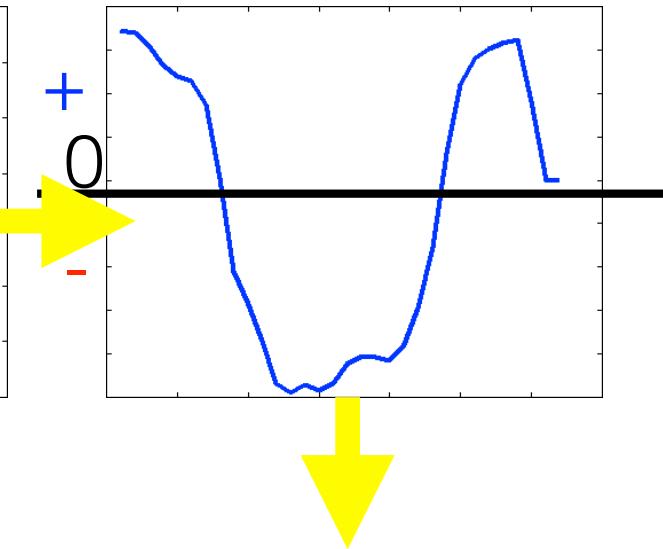
# Evaluation of Feature Discriminant Power



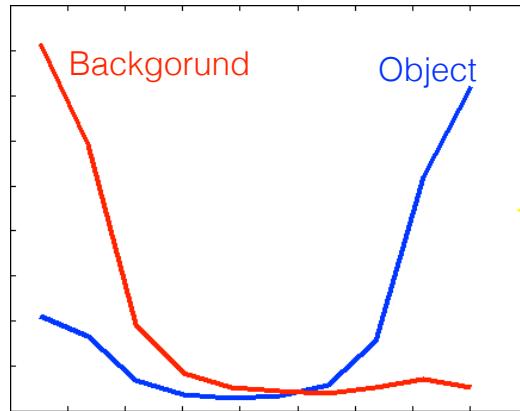
Feature Histograms



Log Likelihood Ratio



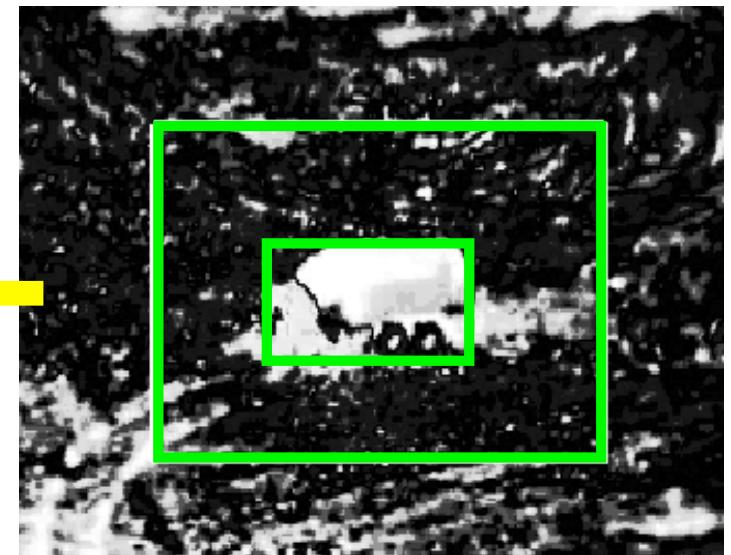
Likelihood Histograms



Feature Score

$$\frac{\text{Var between classes}}{\text{Var within classes}}$$

$$\text{VR}(L; p, q) \equiv \frac{\text{var}(L; (p + q)/2)}{[\text{var}(L; p) + \text{var}(L, q)]}$$



Likelihood Image

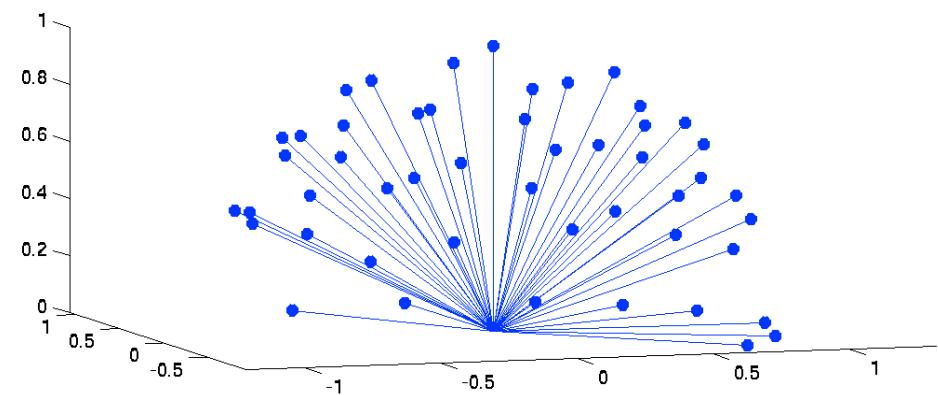
# Example: 1D Color Feature Spaces

Color features: integer linear combinations of R,G,B

$$\frac{aR + bG + cB}{|a| + |b| + |c|} + \text{offset} \quad a, b, c \in \{-2, -1, 0, 1, 2\}$$

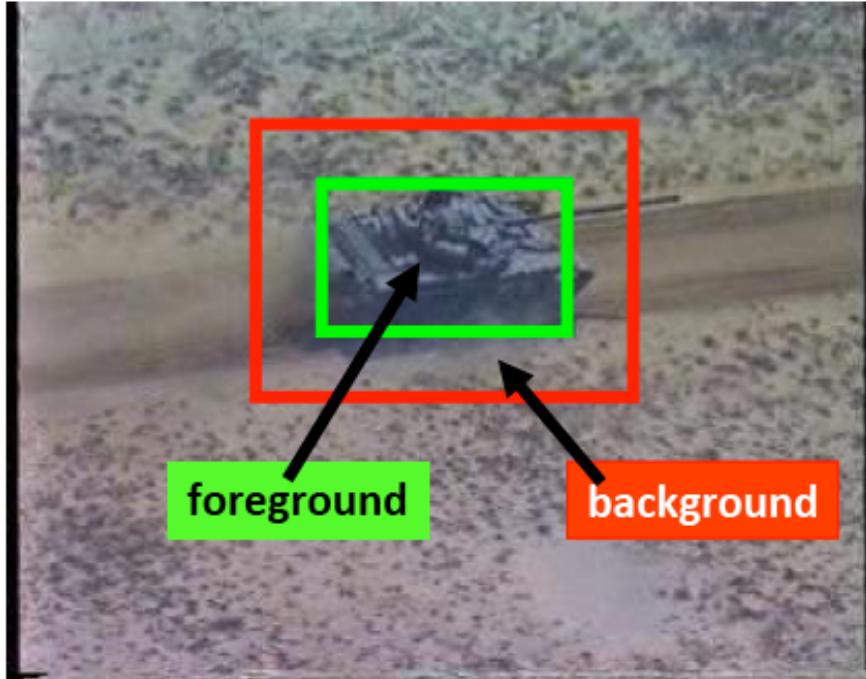
offset to bring feature  
back to 0-255 range

Once you remove duplicates and 0,  
there are 49 distinct features.  
They roughly uniformly sample the 1D  
marginal distributions of RGB.

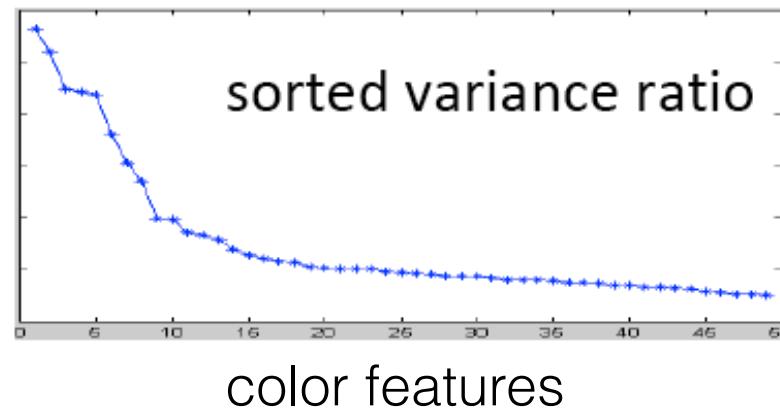


# Example

training frame

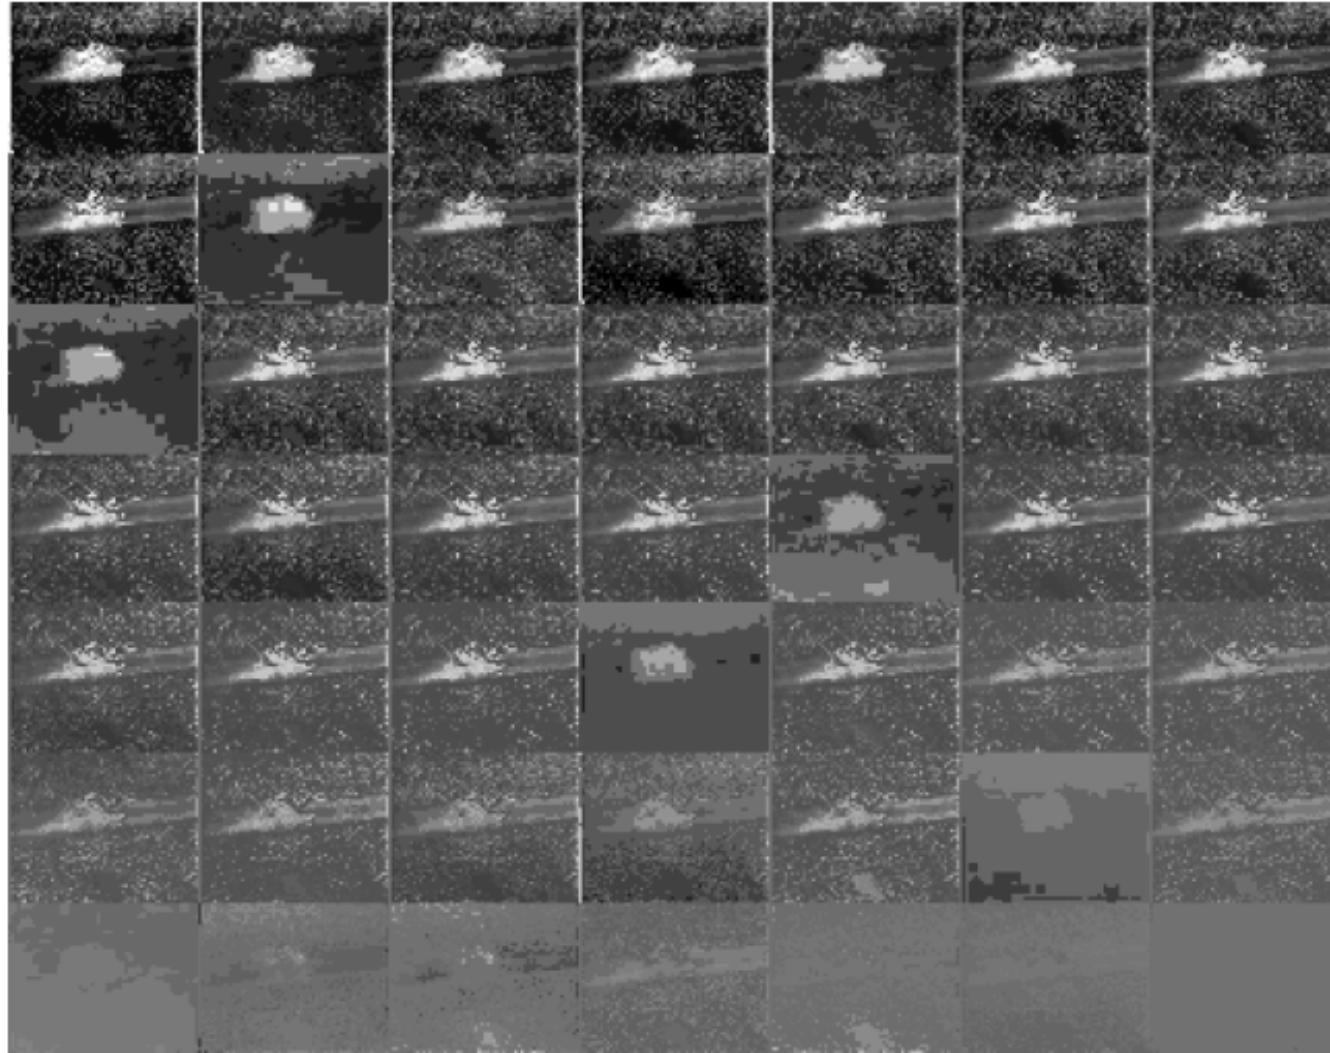


test frame



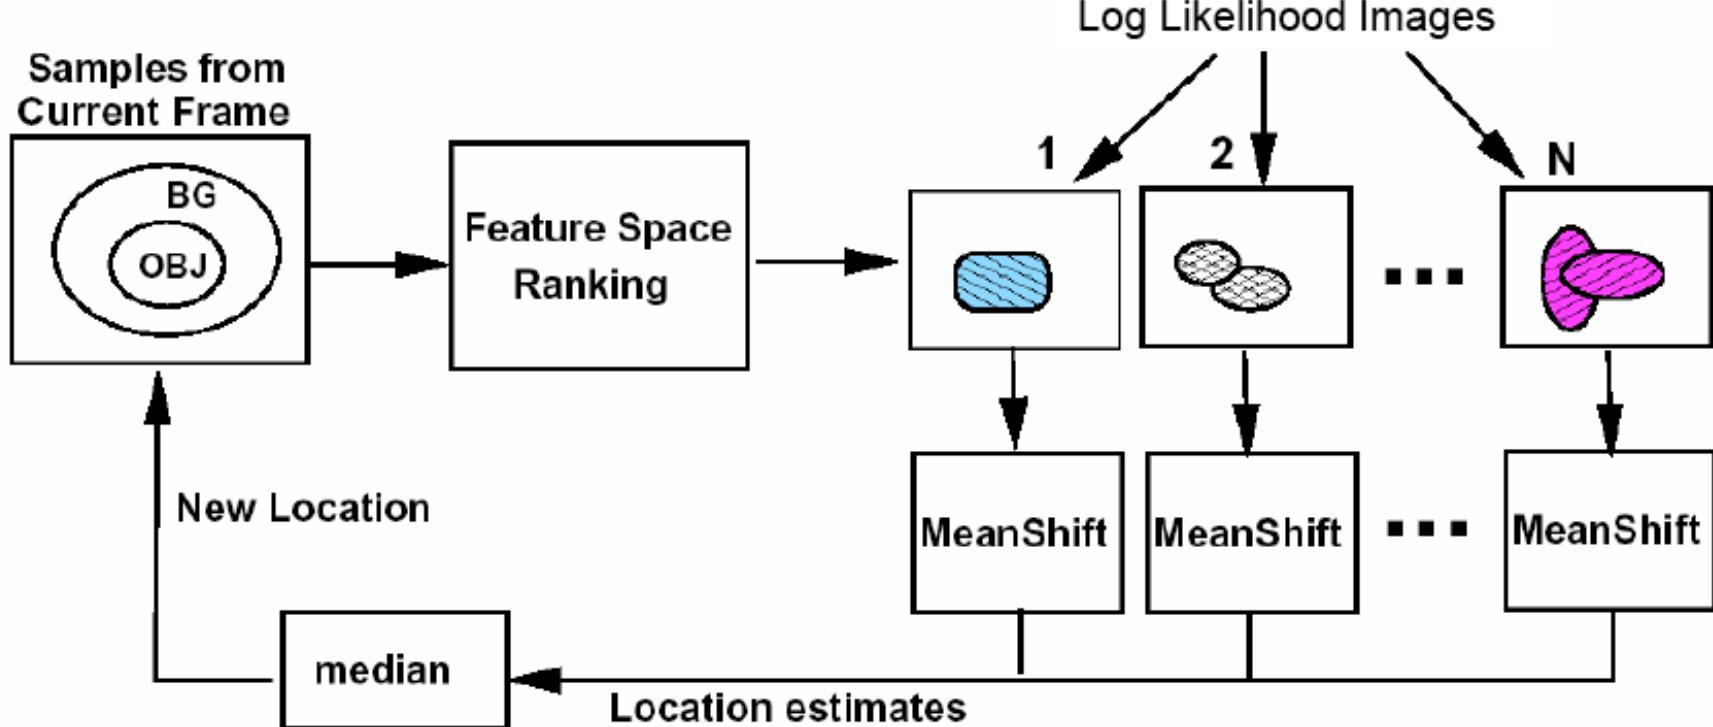
# Example: Feature Ranking

Best



Worst

# Overview Algorithm



Note: since log likelihood images contain negative values, must use modified mean-shift algorithm as described in Collins, CVPR'03

$$\Delta_x = \frac{\sum_a K(a - x) w(a)(a - x)}{\sum_a |K(a - x) w(a)|}$$

use absolute values

# Avoiding Model Drift

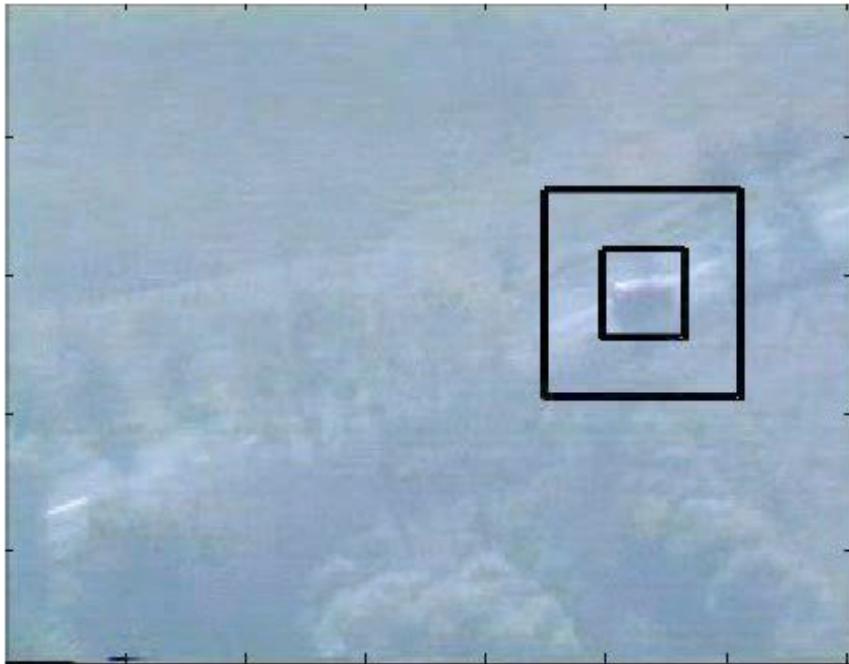
---

The boundingbox on the target is not tight:  
we incorporate background pixels into the model  
we will have drift

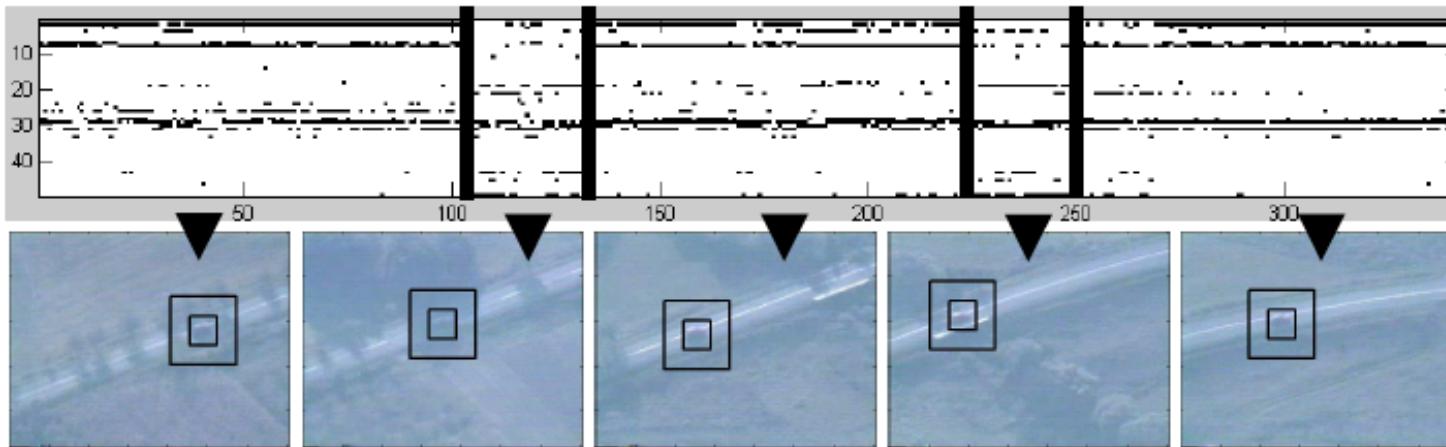
Use anchoring:  
combine current model with model from first frame  
solves drift, but limits the ability of adapting to new appearances

anchor distribution = object appearance histogram from first frame  
model distribution =  $(\text{current distribution} + \text{anchor distribution}) / 2$

# Example: Hard-to-see Objects



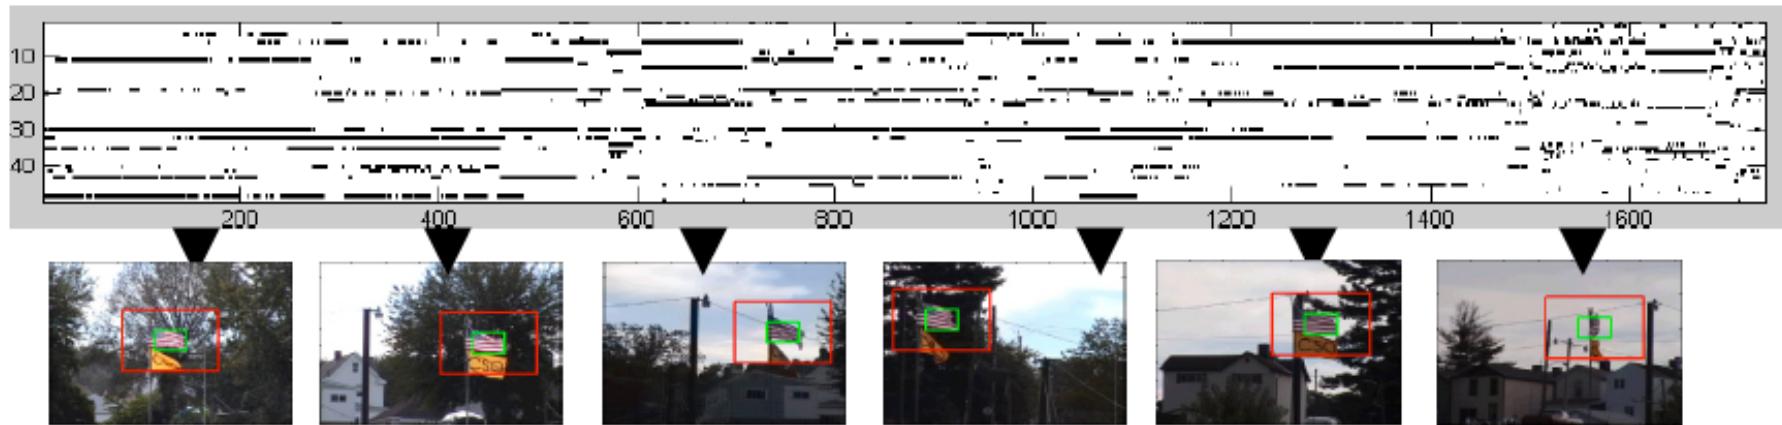
Trace of selected features



# Example: Changing Illumination/Background



Trace of selected features



# Tracking by Detection Using Dense Sampling

(Henriques, Caseiro, Martins & Batista, '12)



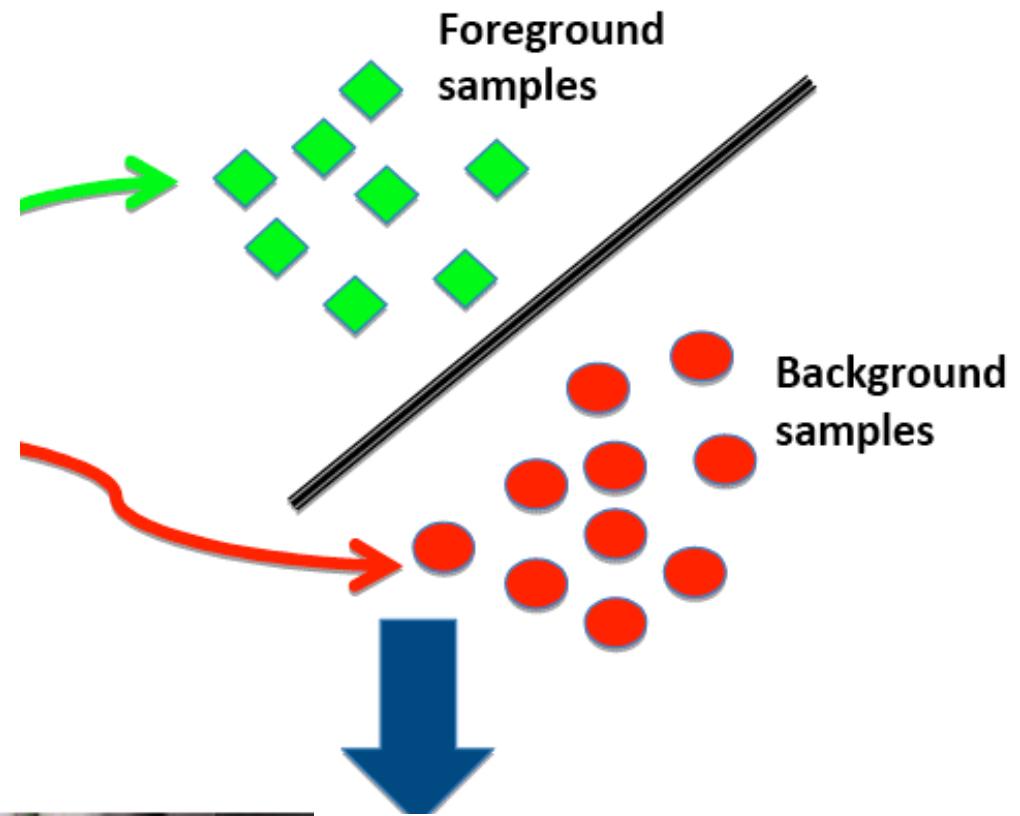
At each frame, collect a set of dense samples around the estimated location of the target.

Samples are continuously labeled (between 1 and 0) using a Gaussian centered at the target location with  $\text{stdev } s$

Uses a Kernel Regularized Least Square Classifier

# Overview

---



# Classifier Design Using Labeled Data

---

Given labeled training samples:

$$(\mathbf{x}_1, y_1), \dots, (\mathbf{x}_l, y_l)$$

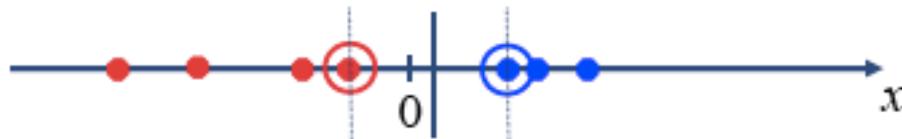
$$y_i \in \{-1, 1\}$$

- Would like to find a classifier function  $F(x; c)$  that given a sample  $x$ , would return its label.

# Classifier Using Labeled Data

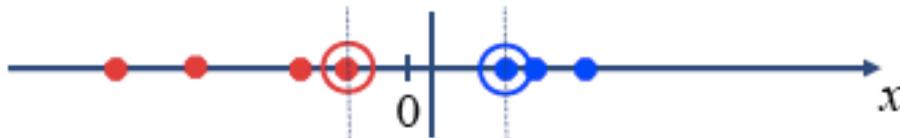
---

- Datasets that are linearly separable work out great:



# Classifier Using Labeled Data

- Datasets that are linearly separable work out great:

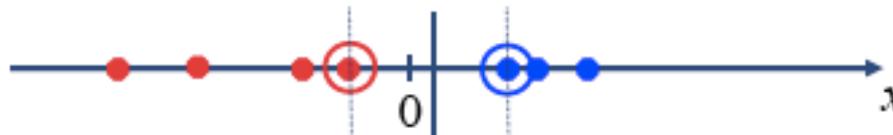


- But what if the dataset is just too hard?

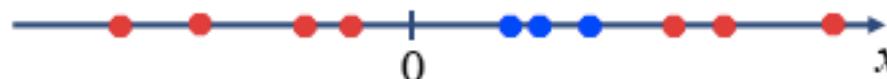


# Classifier Using Labeled Data

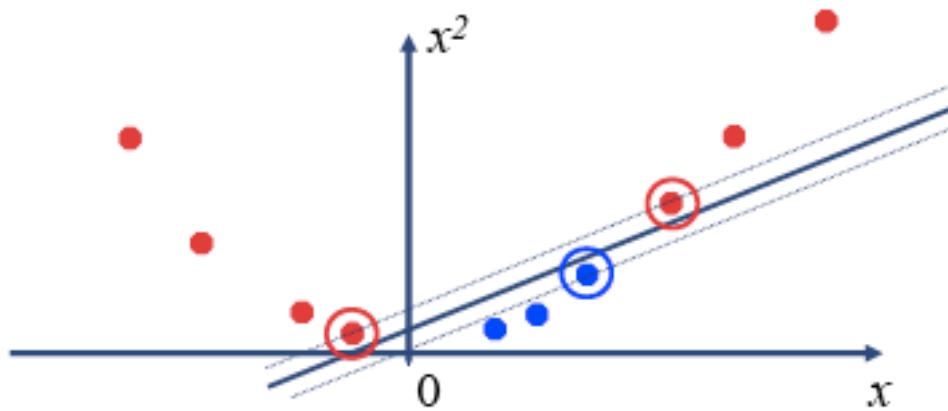
- Datasets that are linearly separable work out great:



- But what if the dataset is just too hard?



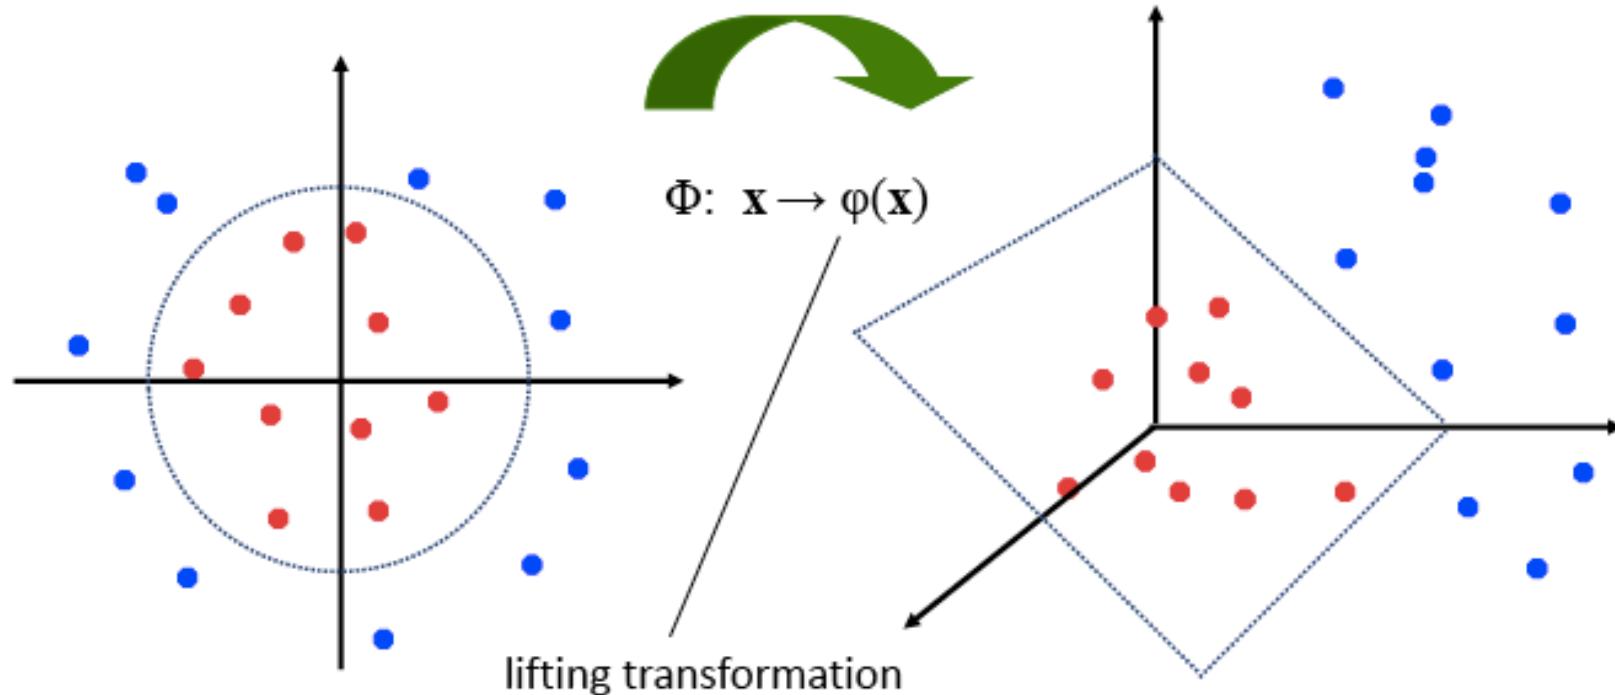
- We can map it to a higher-dimensional space:



Slide credit: Andrew Moore

# “Kernel Trick”

- General idea: the original input space can always be mapped to some higher-dimensional feature space where the training set is separable:



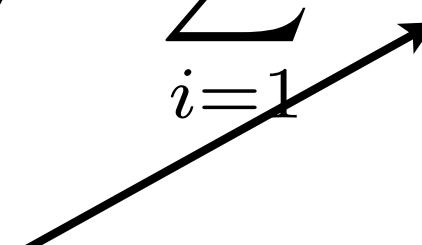
Slide credit: Andrew Moore

# A Classifier: Kernel Regularized Least Square Classifier

Given labeled training samples:

$$(\mathbf{x}_1, y_1), \dots, (\mathbf{x}_l, y_l)$$

- and a “kernel”  $\kappa(\mathbf{x}, \mathbf{x}') = \langle \varphi(\mathbf{x}), \varphi(\mathbf{x}') \rangle$   
  
(Non-linear) transformation → dot product
- Would like to find a classifier function  $f(z; \alpha)$  that given a sample  $z$ , would return its label:

$$y = f(\mathbf{z}; \alpha) = \sum_{i=1}^l K(\mathbf{z}, \mathbf{x}_i) \alpha_i$$


the “distances” between  $z$  and all the labeled samples

# A Classifier: Kernel Regularized Least Square Classifier

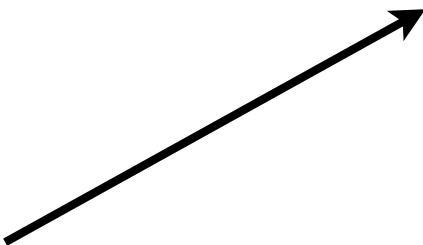
- Would like to find a classifier function  $f(\mathbf{x}; \alpha)$  that given a sample  $\mathbf{x}$ , would return its label:

$$y = f(\mathbf{z}; \alpha) = \sum_{i=1}^l K(\mathbf{z}, \mathbf{x}_i) \alpha_i$$

- Pose it as a LSE problem:

$$\min F(\alpha) = \min_{\alpha \in R^l} \frac{1}{2l} \sum_{i=1}^l (\mathbf{y} - K\alpha)^T (\mathbf{y} - K\alpha) + \frac{\lambda}{2} \alpha^T K \alpha$$

“good oracle” term



“regularization term”  
puts a bound on a

# A Classifier: Kernel Regularized Least Square Classifier

$$\min F(\alpha) = \min_{\alpha \in R^l} \frac{1}{2l} \sum_{i=1}^l (\mathbf{y} - K\alpha)^T (\mathbf{y} - K\alpha) + \frac{\lambda}{2} \alpha^T K \alpha$$

$$\nabla_{\alpha} F(\alpha) = 0 = \frac{1}{l} \sum_{i=1}^l (-K\mathbf{y} + K^2\alpha) + \lambda K\alpha = -K\mathbf{y} + K^2\alpha + \lambda K\alpha$$

$$\boxed{\mathbf{y} = (K + \lambda I)\alpha}$$

Can solve for a 😊

Need to invert a LARGE  $l \times l$  matrix ☹

## Learning:

Given a set of labeled samples, need to compute a

$$\mathbf{y} = (K + \lambda I)\alpha$$

Matrix Inversion

## Testing:

Given an unlabeled sample  $\mathbf{z}$ , need to compute  $y$ :

$$y = f(\mathbf{z}; \alpha) = \sum_{i=1}^l K(\mathbf{z}, \mathbf{x}_i) \alpha_i = K(\mathbf{z}, \mathbf{x}_i)^T \alpha$$

Distances computation, Cross-correlation

# Efficiency considerations

---

We want:

- to use a large number of samples (dense sampling)
- the samples themselves are high dimensional (hundreds, thousands of pixels)
- to learn a new classifier at each frame
- test many unlabeled samples to find a “good” positive example

Solution:

Use “circulant” matrices

Make computations in the FREQUENCY domain

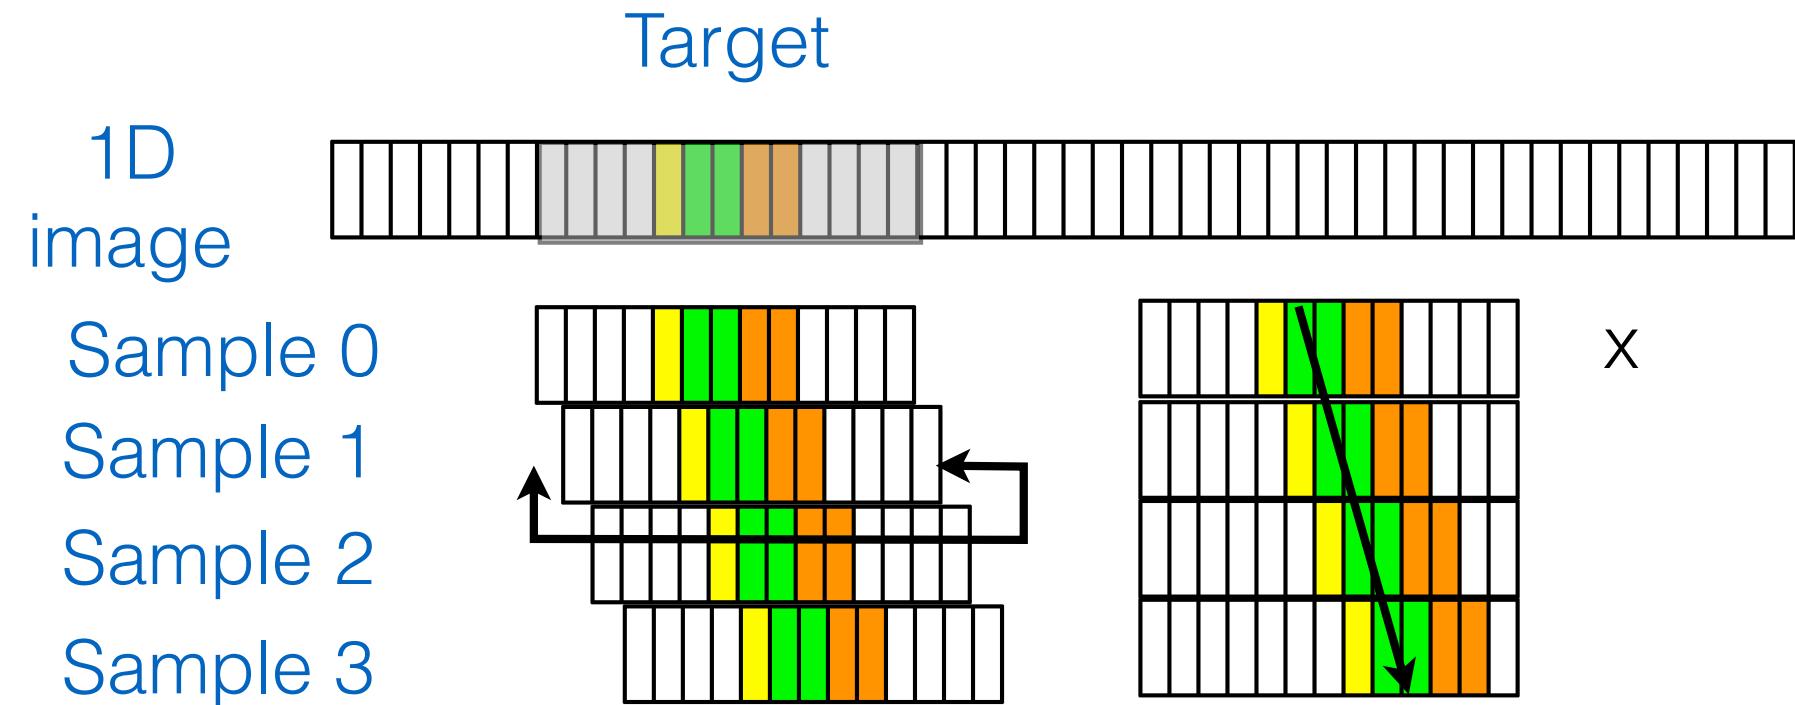
instead of multiplying matrices, we can multiply element by element

instead of inverting matrices, we can divide element by element

Price we pay:

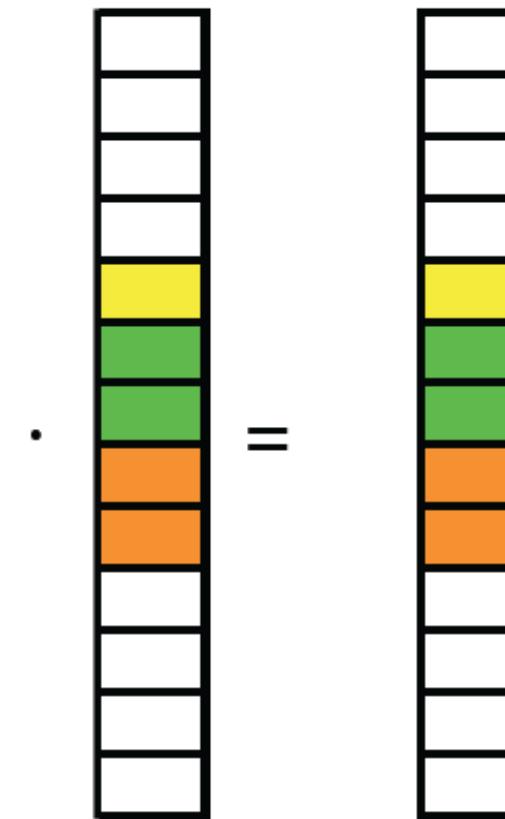
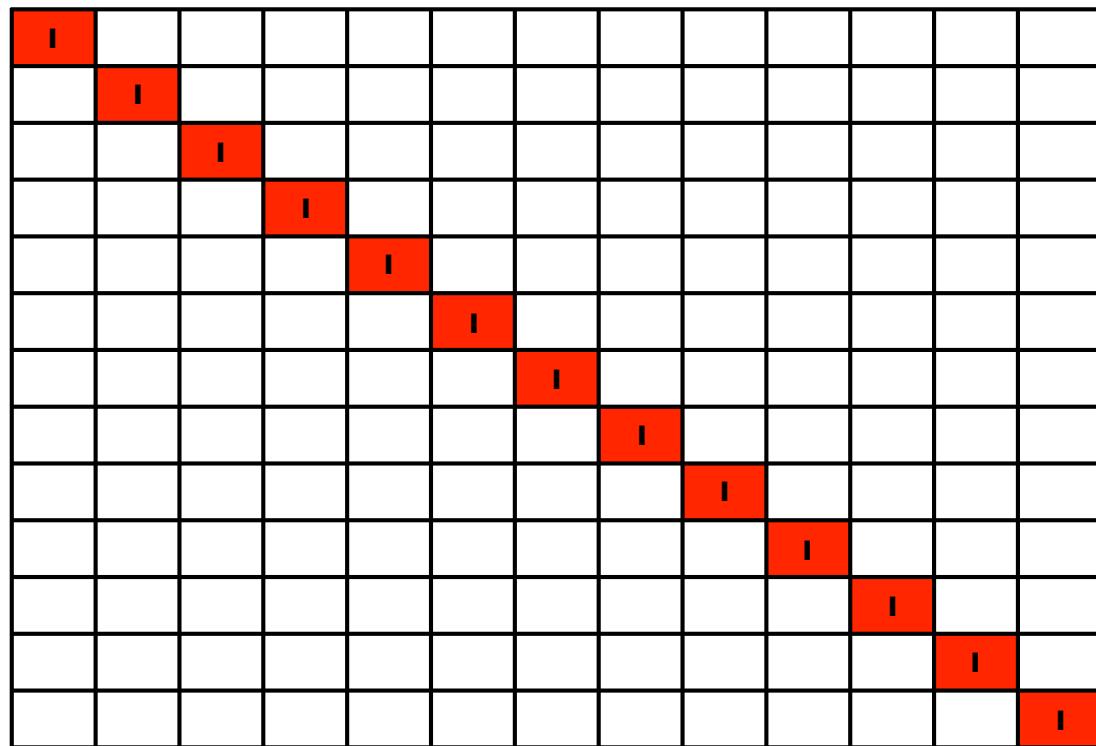
We need to go to the Frequency domain and come back

# “Dense Sampling” = Circulant Matrix



$C(x)$  Circulant Matrix with constant diagonals

$$\mathbf{x}_o = I\mathbf{x} = P^0 \mathbf{x}_o$$



# Permutation Matrix: P

$$\mathbf{x}_1 = P\mathbf{x}$$

A 10x10 grid of black lines on a white background. Red 'I' symbols are placed at various intersections. The symbols are positioned as follows: Row 1, Col 1; Row 2, Col 2; Row 3, Col 3; Row 4, Col 4; Row 5, Col 5; Row 6, Col 6; Row 7, Col 7; Row 8, Col 8; Row 9, Col 9; and Row 10, Col 10.

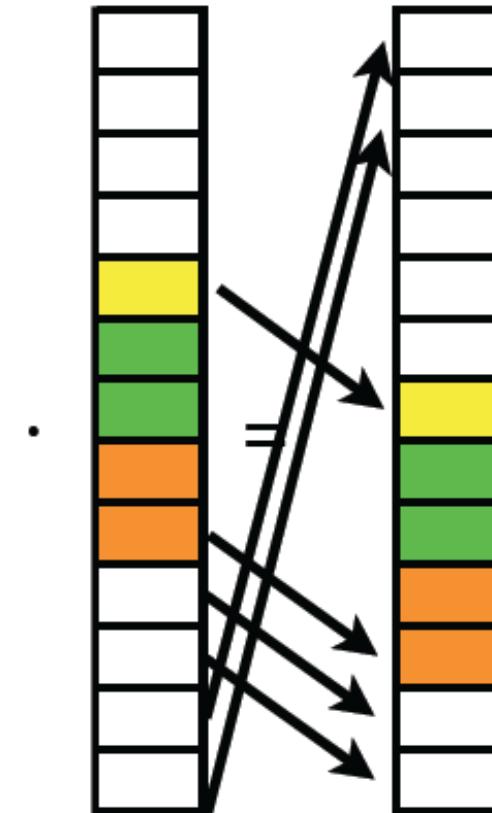


Only one 1 in each row and in each column

# Permutation Matrix: P

$$\mathbf{x}_2 = P\mathbf{x}_1 = PP\mathbf{x} = P^2\mathbf{x}$$

A 10x10 grid with black lines forming a 9x9 grid of cells. Some cells are filled with a solid red color and contain a black capital letter 'I' in their center. The red cells are located at the following coordinates: (1,1), (1,3), (2,2), (3,1), (3,3), (4,2), (5,1), (5,3), (6,2), (6,4), (7,3), (7,5), (8,4), (8,6), and (9,5). The rest of the grid is white.



Only one 1 in each row and in each column

# Kernels & Circulant Matrices

$$C(\mathbf{u}) = \begin{bmatrix} u_0 & u_1 & u_2 & \cdots & u_{n-1} \\ u_{n-1} & u_0 & u_1 & \cdots & u_{n-2} \\ u_{n-2} & u_{n-1} & u_0 & \cdots & u_{n-3} \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ u_1 & u_2 & u_3 & \cdots & u_0 \end{bmatrix} \quad \begin{aligned} c_{ij} &= u_{(j-i) \bmod n} \\ \mathbf{x}_i &= P^i \mathbf{x}, \quad \forall i = 0, \dots, n-1 \\ n &\text{ is the number of samples} \end{aligned}$$

**Theorem 1.** *The matrix  $K$  with elements  $K_{ij} = \kappa(P^i \mathbf{x}, P^j \mathbf{x})$  is circulant if  $\kappa$  is a unitarily invariant kernel.*

Proof:  $K_{ij} = \kappa(P^i \mathbf{x}, P^j \mathbf{x})$

$\kappa$  is unitarily invariant if  $\kappa(\mathbf{x}, \mathbf{x}') = \kappa(U\mathbf{x}, U\mathbf{x}')$  for any unitary matrix  $U$

Permutations are unitary matrices.

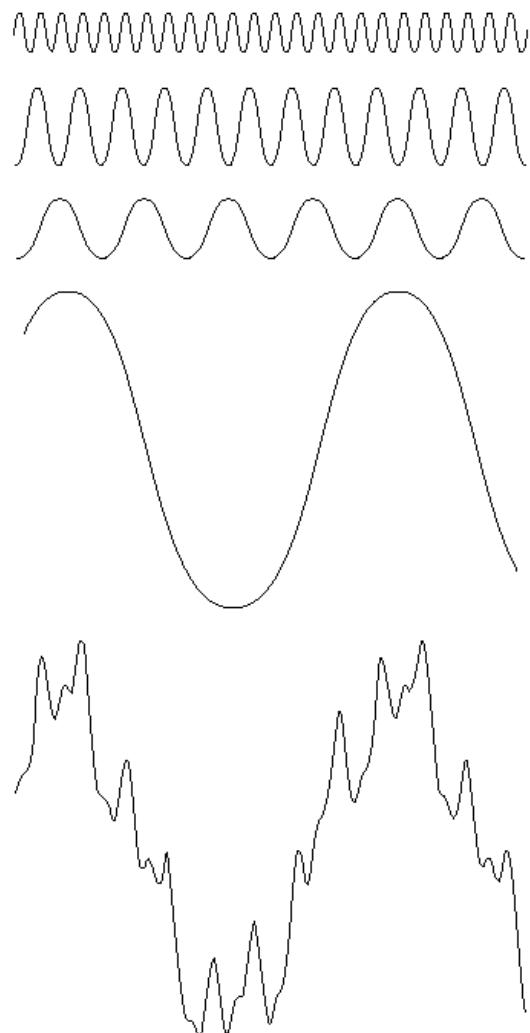
$$K_{ij} = \kappa(P^{-i} P^i \mathbf{x}, P^{-i} P^j \mathbf{x}) = \kappa(\mathbf{x}, P^{j-i} \mathbf{x})$$

depends only on  $(j-i) \bmod n$ .

QED

# The Frequency Domain

---



The sum of the four top functions (single frequencies) gives this “noisy” looking signal.

**FIGURE 4.1** The function at the bottom is the sum of the four functions above it. Fourier’s idea in 1807 that periodic functions could be represented as a weighted sum of sines and cosines was met with skepticism.

# Review: Complex Numbers

Complex Numbers:  $C = R + jI$

Conjugate:  $C^* = R - jI$

Polar  
Coordinates:  $C = |C|(\cos \theta + j \sin \theta)$

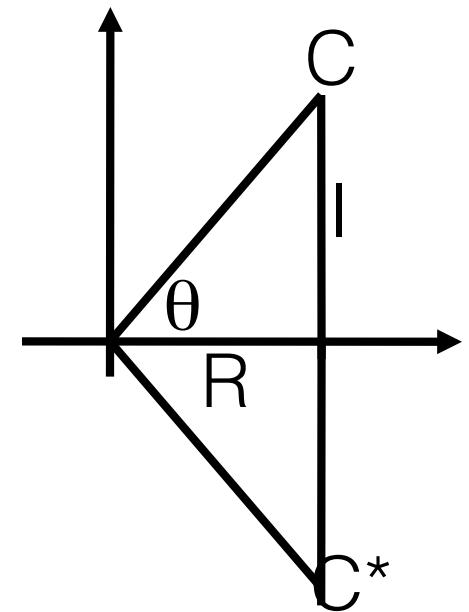
with  $|C| = \sqrt{R^2 + I^2}$

$$\tan \theta = \frac{I}{R}$$

Using Euler's formula:

$$C = |C|e^{j\theta}$$

$$e^{j\theta} = \cos \theta + j \sin \theta$$



# Sampling f(x)

---

Samples of  $f(x)$  and  $F(u)$  are not necessarily at integer values, but they should be at equally spaced points:

$$f(x) \doteq f(x_o + k\Delta x) \quad \text{With } k \text{ integer}$$

$$F(u) \doteq F(u\Delta u) \quad \text{With } u \text{ integer}$$

$\Delta x$  and  $\Delta u$  are related:

$$\Delta u = \frac{1}{M \Delta x}$$

# Discrete Fourier Transform (DFT)

---

One Dimension :

Let  $f(x)$  be a discrete function of  $x = 0, 1, 2, \dots, M-1$

$$F(u) = \sum_{x=0}^{x=M-1} f(x) e^{-j2\pi ux/M} \quad u = 0, 1, \dots, M-1$$

$$f(x) = \frac{1}{M} \sum_{u=0}^{u=M-1} F(u) e^{+j2\pi ux/M} \quad x = 0, 1, \dots, M-1$$

NOTE:  $f(x)$  is real BUT  $F(u)$  will be, in general, complex

# Euler's Formula and the Frequency Domain

---

$$F(u) = \sum_{x=0}^{x=M-1} f(x)e^{-j2\pi ux/M} \quad u = 0, 1, \dots, M-1$$

Euler's formula:  $e^{j\theta} = \cos \theta + j \sin \theta$

$$F(u) = \sum_{x=0}^{x=M-1} f(x) \left[ \cos \frac{2\pi ux}{M} - j \sin \frac{2\pi ux}{M} \right] \quad u = 0, 1, \dots, M-1$$

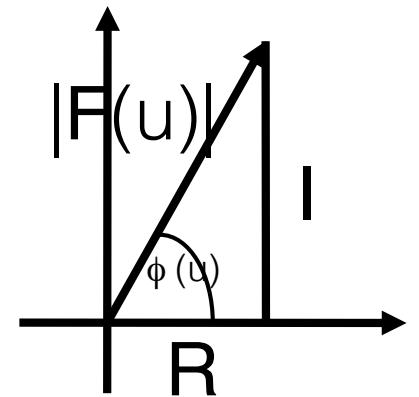
- For each value of  $u$ , we need to sum ALL values of  $f(x)$
- $u$  determines the frequency of the transform

# Fourier Transform

---

$$F(u) = R(u) + jI(u)$$

$$F(u) = |F(u)| e^{j\phi(u)}$$



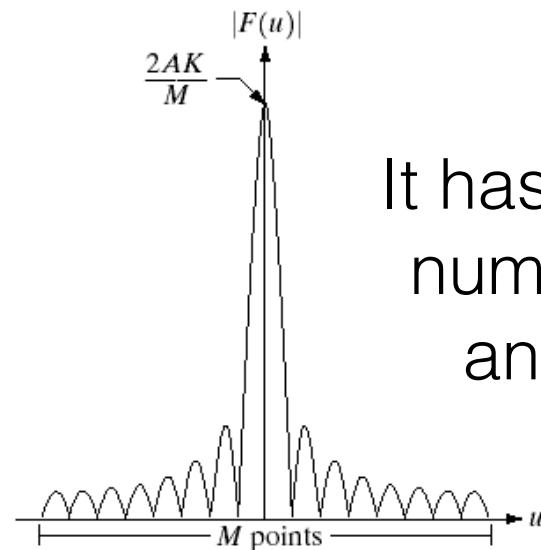
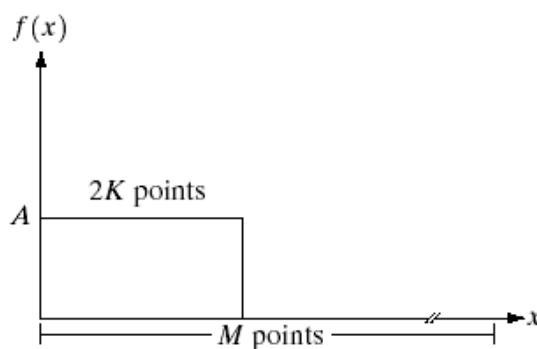
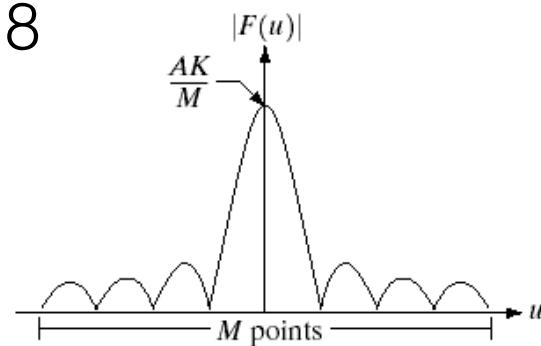
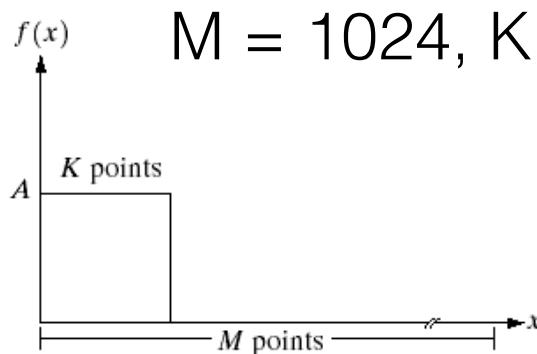
where  $|F(u)| = \sqrt{R^2(u) + I^2(u)}$  “spectrum of  $f(x)$ ”

$P(u) = |F(u)|^2$  “power spectrum or spectral density of  $f(x)$ ”

$\phi(u) = \tan^{-1} \frac{I(u)}{R(u)}$  “phase angle”

# 1D DFT Example

The functions are **discrete**, only drawn as continuous for visualization!



a	b
c	d

**FIGURE 4.2** (a) A discrete function of  $M$  points, and (b) its Fourier spectrum. (c) A discrete function with twice the number of nonzero points, and (d) its Fourier spectrum.

It has doubled the number of zeros and doubled height

# FT and Convolution

---

$$\mathcal{F}[f(t) * h(t)] = F(u)H(u)$$

$$\mathcal{F}[f(t)h(t)] = F(u) * H(u)$$

$$f(t) * h(t) = \mathcal{F}^{-1} [\mathcal{F}(u) \cdot \mathcal{H}(u)]$$

# Circulant Matrices, Kernels + FFT = Efficient!

## Learning:

Given a set of labeled samples, need to compute a

$$\mathbf{y} = (K + \lambda I)\boldsymbol{\alpha}$$

Matrix Inversion

$K = C(\mathbf{k})$  Circulant matrix with columns:

$$k_i = \kappa(\mathbf{x}, P^i \mathbf{x}), \quad \forall i = 0, \dots, n - 1$$

$$\boldsymbol{\alpha} = \mathcal{F}^{-1} \left( \frac{\mathcal{F}(\mathbf{y})}{\mathcal{F}(\mathbf{k}) + \lambda} \right), \text{ element-wise division!}$$

For  $n \times n$  images, the proposed algorithm has a complexity of only  $\mathcal{O}(n^2 \log n)$

# Circulant Matrices, Kernels + FFT = Efficient!

- Testing:

- Given an unlabeled sample  $\mathbf{z}$ , need to compute  $\mathbf{y}$ :

$$y = f(\mathbf{z}; \boldsymbol{\alpha}) = \sum_{i=1}^l K(\mathbf{z}, \mathbf{x}_i) \alpha_i = K(\mathbf{z}, \mathbf{x}_i)^T \boldsymbol{\alpha}$$

$K = C(\mathbf{k})$  Distances computation, Cross-correlation

$$k_i = \kappa(\mathbf{x}, P^i \mathbf{x}), \quad \forall i = 0, \dots, n-1$$

$$\bar{k}_i = \kappa(\mathbf{z}, P^i \mathbf{x})$$

$$\hat{\mathbf{y}} = \mathcal{F}^{-1} (\mathcal{F}(\bar{\mathbf{k}}) \odot \mathcal{F}(\boldsymbol{\alpha}))$$
 element-wise multiplication!

| For  $n \times n$  images, the proposed algorithm has a complexity of only  $\mathcal{O}(n^2 \log n)$

# Fast Kernel Computation

Linear case: no Kernel Trick, just dot product

$$\kappa(\mathbf{x}, \mathbf{x}') = \langle \mathbf{x}, \mathbf{x}' \rangle$$

$$\mathbf{w} = \mathcal{F}^{-1} \left( \frac{\mathcal{F}(\mathbf{x}) \odot \mathcal{F}^*(\mathbf{y})}{\mathcal{F}(\mathbf{x}) \odot \mathcal{F}^*(\mathbf{x}) + \lambda} \right)$$

# Fast Kernel Computation

## Gaussian Kernel

$$k^{\text{gauss}} = \exp\left(-\frac{1}{\sigma^2} \left( \|x\|^2 + \|x'\|^2 - 2\mathcal{F}^{-1}(\mathcal{F}(x) \odot \mathcal{F}^*(x')) \right)\right)$$

# Implementation Details

---

Images are not really periodic ... bound them using a sinusoidal window:

$$(x_{ij}^{\text{raw}} - 0.5) \sin(\pi i/n) \sin(\pi j/n), \quad \forall i, j = 0, \dots, n-1$$

- Use a “continuous” class label to yield smooth spatial responses:

$$y_{ij} = \exp\left(-\left((i - i')^2 + (j - j')^2\right) / s^2\right), \quad \forall i, j = 0, \dots, n-1$$

- “Incorporates temporal memory” by linearly interpolating the new parameters with the ones from the previous frame:

```
alphaf = (1 - interp_factor) * alphaf + interp_factor * new_alphaf;  
z = (1 - interp_factor) * z + interp_factor * new_z;
```

# CODE

---

---

**Algorithm 1 . MATLAB code for our tracker, using a Gaussian kernel**

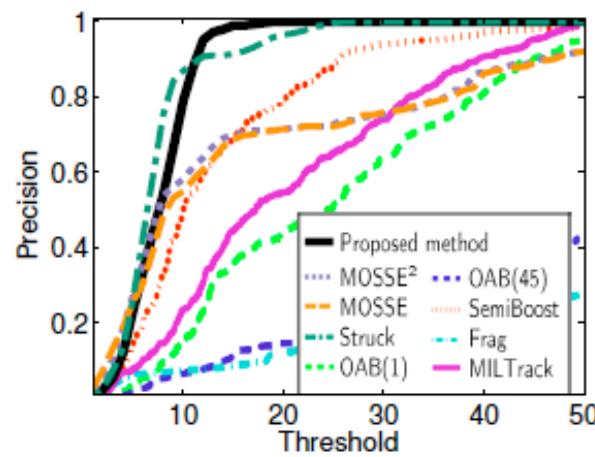
It is possible to reuse some values, reducing the number of FFT calls. An implementation with GUI is available at: <http://www.isr.uc.pt/~henriques/>

---

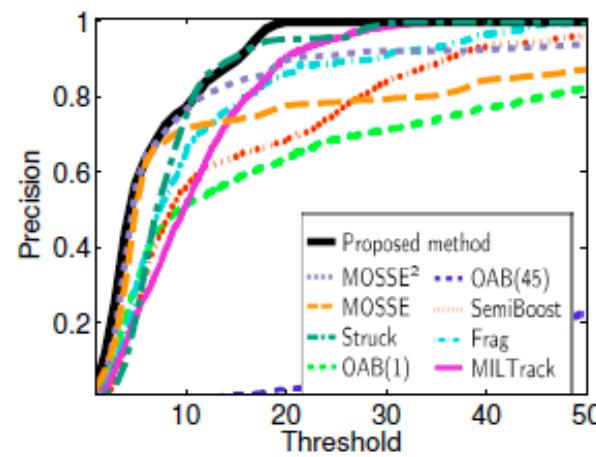
```
% Training image x(current frame) and test image z(next frame)  
% must be pre-processed with a cosine window. y has a Gaussian  
% shape centered on the target. x, y and z are M-by-N matrices.  
% All FFT operations are standard in MATLAB.  
  
function alphaf = training(x, y, sigma, lambda) % Eq. 7  
    k = dgk(x, x, sigma);  
    alphaf = fft2(y) ./ (fft2(k) + lambda);  
end  
  
function responses = detection(alphaf, x, z, sigma) % Eq. 9  
    k = dgk(z, x, sigma);  
    responses = real(ifft2(alphaf .* fft2(k)));  
end  
  
function k = dgk(x1, x2, sigma) % Eq. 16  
    c = fftshift(ifft2(fft2(x1) .* conj(fft2(x2))));  
    d = x1(:)' * x1(:) + x2(:)' * x2(:) - 2*c;  
    k = exp(-1 / sigma^2 * abs(d) / numel(x1));  
end
```

---

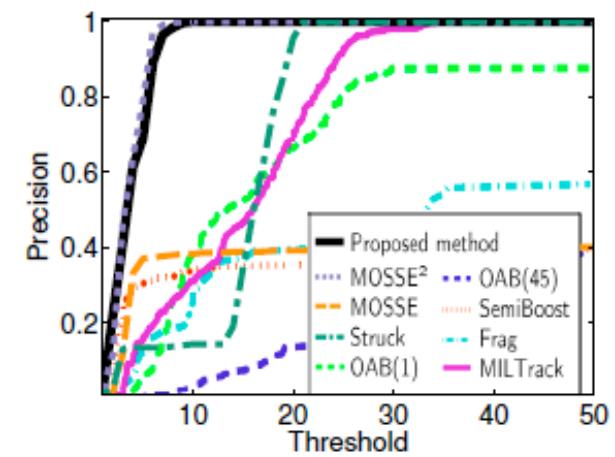
# Performance (300 fps)



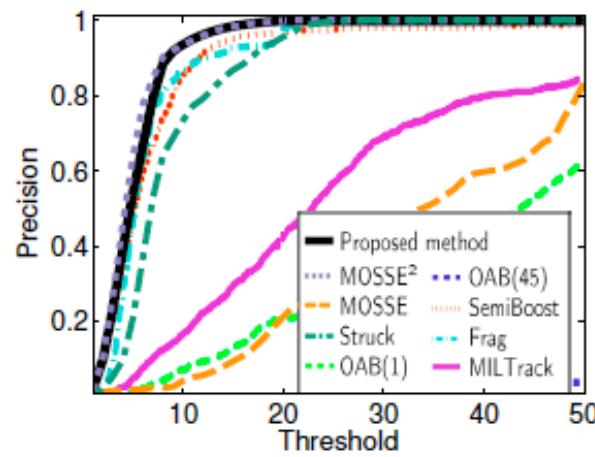
(a) coke11



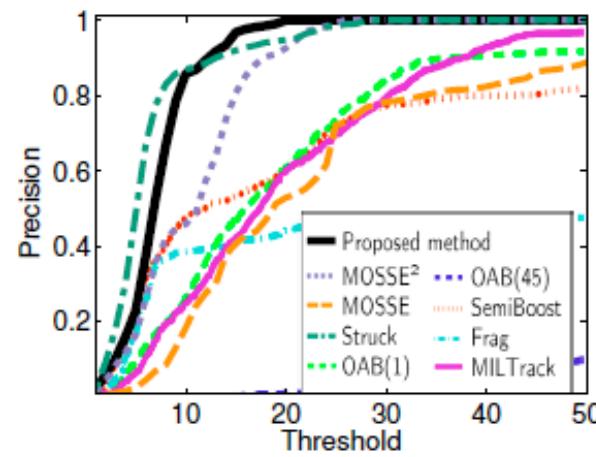
(b) sylvester



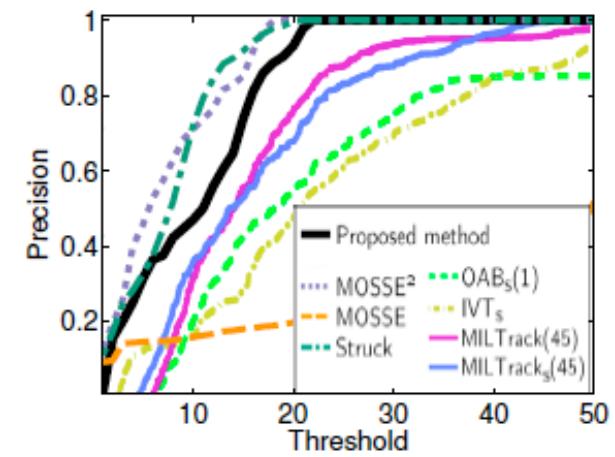
(c) dollar



(d) faceocc



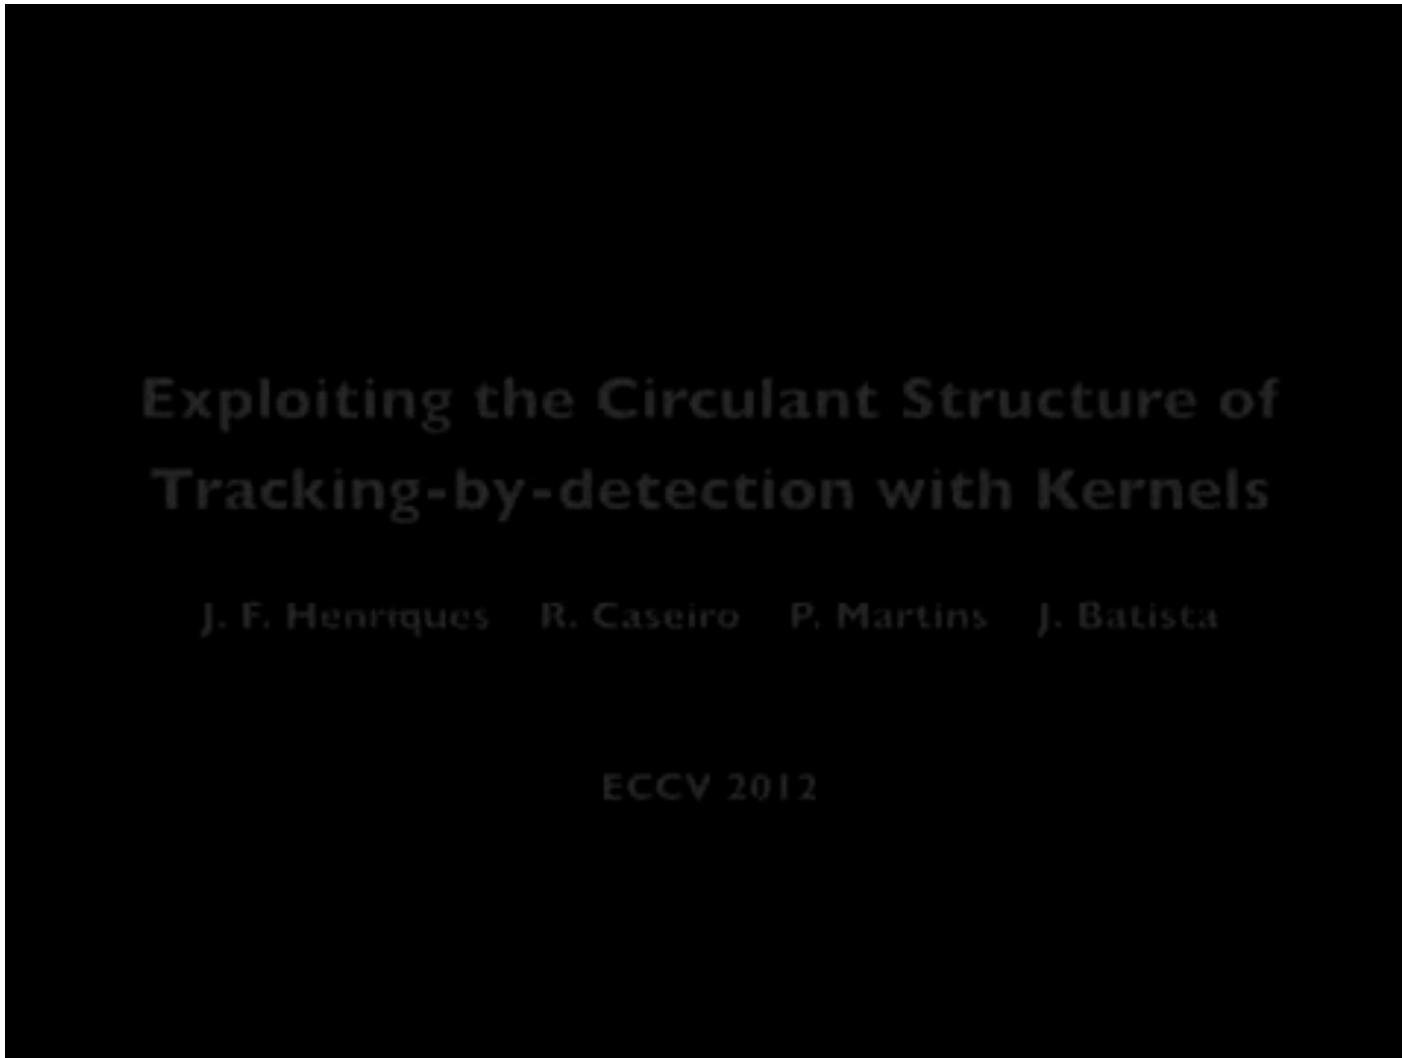
(e) faceocc2



(f) twinings

# Videos

---



<http://www.youtube.com/watch?v=sZXKnPUhAi8>